



Storage and Indexing

Discussion 9




Question 1

Consider a disk with sector size of 512 bytes, 2000 tracks per surface, 50 sectors per track, five double-sided platters and average seek time of 10 msec. Suppose a page size of 1024 bytes is chosen. Suppose a file containing 100,000 records of 100 bytes each is to be stored on such a disk and that no record is allowed to span two pages



Question 1

- How many pages fit in a track?
- What is the disk capacity in bytes?
- How many records fit onto a page?
- How many pages are required to store the entire file?





Given:

Page size = 1024 bytes

Sector size = 512 bytes

Number of records = 100,000

Size of a record = 100 bytes

Number of tracks/surface = 2000

bytes/track = bytes/sector \times sectors/track = $512 \times 50 = 25\text{K}$

bytes/surface = bytes/track \times tracks/surface = $25\text{K} \times 2000 = 50,000\text{K}$


bytes/disk = bytes/surface \times surfaces/disk = $50,000\text{K} \times 5 \times 2 = 500,000\text{K}$

1. Number of pages in a track = $25600 / 1024 = 25$
2. Disk capacity = $500,000\text{K}$
3. records/page = Size of page / size of record = $1024/100 > 10$. So 10 records can fit into a page
4. Number of pages required for the entire file = $100,000/10 = 10000$ pages



Question 2: B+ trees


Consider the following B+-tree, and assume the convention that the left pointer points to values that are strictly less than the key value. Each node in this B+-tree can hold up to four entries (i.e., the order of the tree is 2).





Question 2: B+ trees

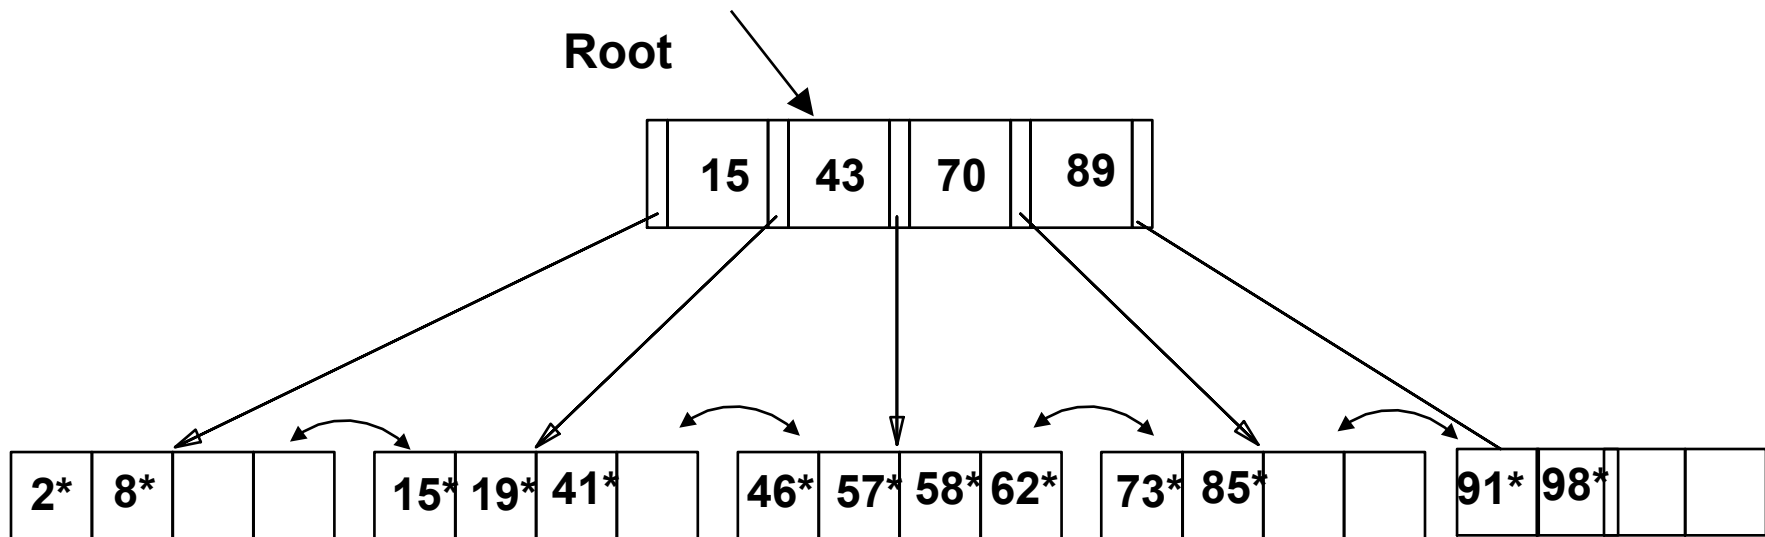
Part (a) Draw the B+-tree after inserting data entries with key values **15, 22, 51, and 102 (in that order)** into the B+-tree. Assume that the insert algorithm does not consider redistributing entries.





Question 2: B+ trees


Inserting 15





Question 2: B+ trees


Inserting 22





Question 2: B+ trees


Inserting 51





Question 2: B+ trees


Inserting 102





Question 2: B+ trees


Part (b) Consider the above B+-tree (prior to the insertions). Draw the B+-tree after deleting the data entries with key values **73, 46, 2 (in that order)**. **Assume that the delete algorithm** considers redistributing entries.





Question 2: B+ trees


Deleting 73





Question 2: B+ trees


Deleting 46





Question 2: B+ trees

Deleting 2





Question 3: Extendible Hashes

a) Consider an extendible hash index that uses the following hash function: **hashvalue = key mod 2^d** , where d is the global or local depth of the index. Assume that for this index the bucket capacity is 2 entries. d is initially 1.

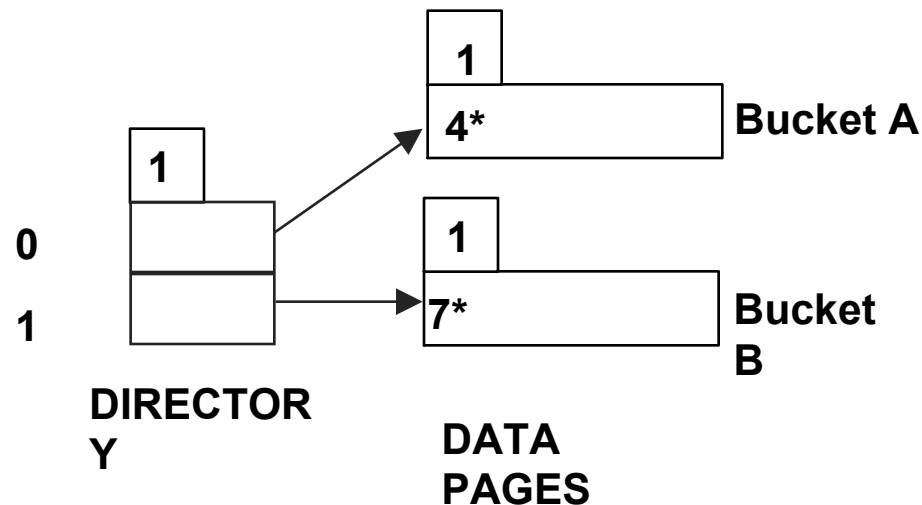
Assume that the index is initially empty. Then two entries with key values **4 and 7 are inserted.**

Draw the resulting index. *Clearly show the global and local depths and all bucket pointers.*



Question 3: Extendible Hashes


a) Assume that the index is initially empty. Then two entries with key values **4** and **7** are inserted. Draw the resulting index. *Clearly show the global and local depths and all bucket pointers.*





Question 3: Extendible Hashes

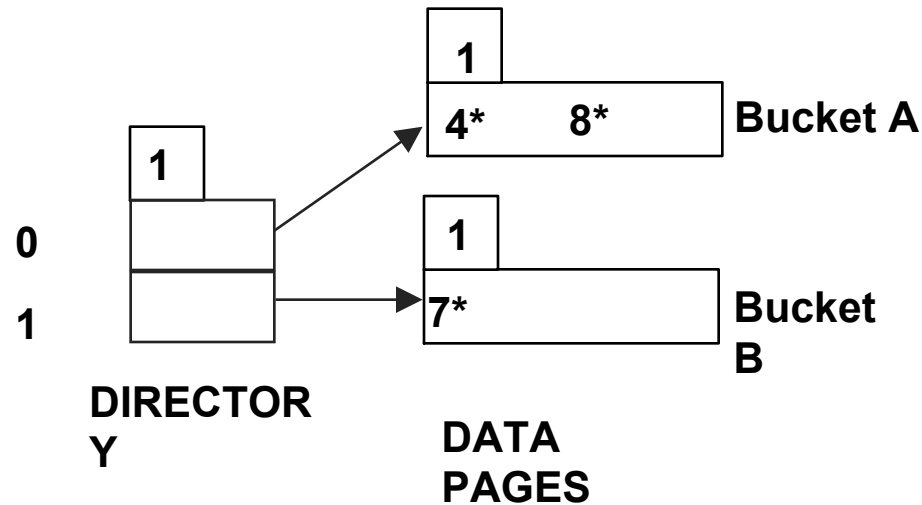
b) After inserting the keys 4 and 7 above, the following three additional keys are inserted into the index: **8, 2, 5**. Draw the final index structure with all the five keys. Clearly show the global and local depths, and all bucket pointers





Question 3: Extendible Hashes

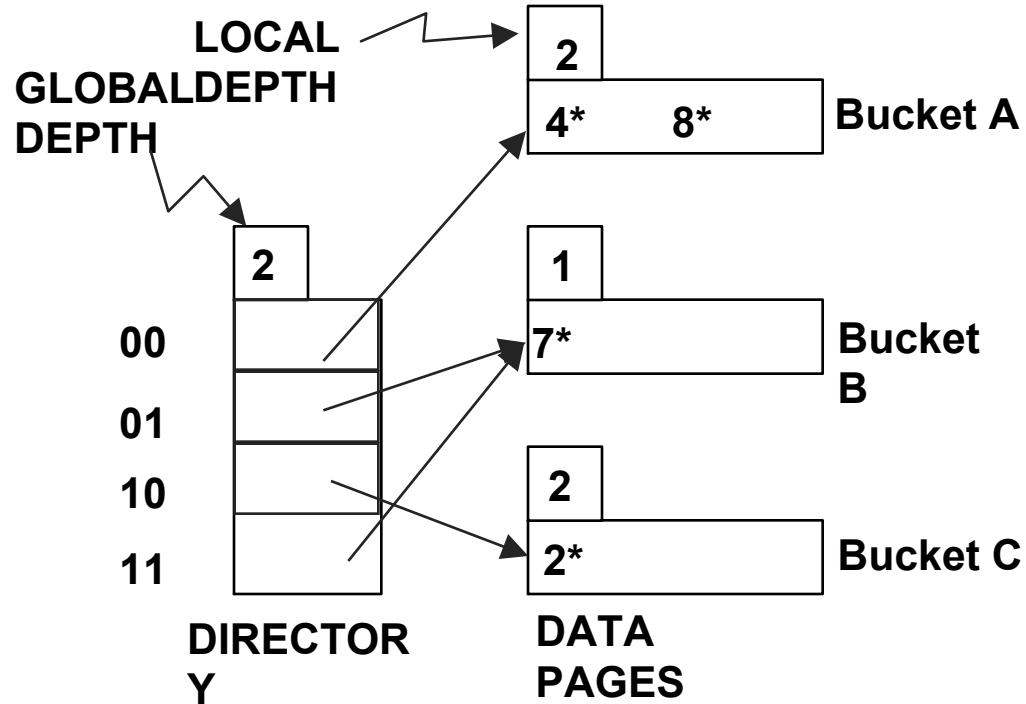
Inserting 8





Question 3: Extendible Hashes


Inserting 2





Question 3: Extendible Hashes

Inserting 5






Question 4: Linear Hashes

Consider a linear hash index. As was the case in the lecture example, let the initial function be $h_0(k) = f(k) \% m$, then any later hash function $h_i(k) = f(k) \% 2^i m$. This way, it is guaranteed that if h_i hashes a key to bucket j element of $[0 \dots 2^i m - 1]$, h_{i+1} will hash the same key to either bucket j or bucket $j + 2^i m$. At any time, two hash functions h_i and h_{i+1} are used.

Assume we are starting at round 0, with the "Next" pointer set to bucket 0. Each page can hold at most 4 entries. The current state of the hash table is as follows:





Question 4: Linear Hashes

a) Please show the state of the hash buckets after inserting **11**. Be sure to indicate any new buckets, the location of the "Next" pointer, any overflow pages, and the two hash functions currently in use.

