



# Discussion 10

---

Query Evaluation and Homework 3 solutions



# Query Evaluation

---

- B+ tree can be used for range and equality selection
- Hash based indexing can be used for only equality selection
- Composite keys - combination of keys
  - Hash: Requires all the keys on which hash is created
  - B+ tree : Prefix is sufficient



## Ex 12.5 (1)

---

- Consider the schema with the Sailors relation:
  - Sailors(*sid*: integer, *sname*: string, *rating*: integer, *age*: real)
  - Sailor tuple size = 50 bytes
  - A page can hold 80 tuples
  - Total 500 pages

What are the total number of tuples?

Total number of tuples = total number of pages \* tuples/page =  $80 * 500$   
= 40000



## Ex 12.5 (1)

- Consider the schema with the Sailors relation:
  - Sailors(*sid*: integer, *sname*: string, *rating*: integer, *age*: real)
  - Sailor tuple size = 50 bytes
  - A page can hold 80 tuples
  - Total 500 pages
- Assume that we have a B+-tree index T on the search key Sailors.sid , and assume that Height(T) = 4, INPages(T ) = 50, Low(T ) = 1, and High(T ) = 100,000. Assume that the root is in main memory
  - $\sigma_{\text{Sailors.sid} < 50,000}(\text{Sailors})$
  - $\sigma_{\text{Sailors.sid} = 50,000}(\text{Sailors})$
- Estimate the number of pages retrieved



# Ex 12.5 (1) Solution

- $\sigma_{\text{Sailors.sid} < 50,000}(\text{Sailors})$
- Assuming uniform distribution, 20,000 tuples will match the condition
- Clustered index:
  - Cost for finding leaf node: 4 I/Os
  - Cost for accessing leaf nodes: 24 I/Os
  - Cost for accessing matching tuples: 250 I/Os
  - Total cost = 278 I/Os
- Unclustered index:
  - Cost for accessing matching tuples: 20,000 I/Os
  - Total cost = 20,028 I/Os using index
  - Using file scan = 500 I/Os
- $\sigma_{\text{Sailors.sid} = 50,000}(\text{Sailors})$
- Sid is primary key so only one matching tuple
- Cost = Height of tree + cost to read qualifying tuple = 4+1=5 I/Os



## Ex 12.5 (2)

---

- Consider the schema with the Sailors relation:
  - Sailors(*sid*: integer, *sname*: string, *rating*: integer, *age*: real)
  - Sailor tuple size = 50 bytes
  - A page can hold 80 tuples
  - Total 500 pages
- Assume that we have a hash index T on the search key Sailors.sid , and assume that  $IHeight(T) = 2$ ,  $INPages(T) = 50$ ,  $Low(T) = 1$ , and  $High(T) = 100,000$ .
  - $\sigma_{Sailors.sid < 50,000}(Sailors)$
  - $\sigma_{Sailors.sid = 50,000}(Sailors)$
- Estimate the number of pages retrieved



## Ex 12.5 (2) Solution

---

- $\sigma_{\text{Sailors.sid} < 50,000}(\text{Sailors})$ 
  - Hash index on sid cannot be used for range queries
  - Need to do a file scan that require 500 I/Os
- $\sigma_{\text{Sailors.sid} = 50,000}(\text{Sailors})$ 
  - Sid is primary key so only one matching tuple
  - Cost = Depth (or levels) of hash table + cost to read qualifying tuple =  $2+1=3$  I/Os



# External Merge Sort

---

- **Run:** Sorted subfile
- B buffer pages
- Pass 0: Uses B buffer pages
  - Produces  $\lceil N/B \rceil$  sorted runs of B pages each
- Pass 1, 2,...: merge **B-1** runs, using (B-1)-way merge
- Number of passes:

$$1 + \lceil \log_{B-1} \lceil N/B \rceil \rceil$$

- Cost =  $2N * (\# \text{ of passes})$





## Ex 13.1

---

- Suppose you have a file with 20,000 pages and you have five buffer pages. Answer the following questions:
  - How many runs will you produce in the first pass?
  - How many passes will it take to sort the file completely?
  - What is the total I/O cost of sorting the file?



# Ex 13.1 Solution

- Number of runs after the first pass
- Pass 0 produces  $\lceil N/B \rceil$  runs of B pages each, where N is the number of file pages and B is the number of available buffer pages
- $\text{ceil}(20000/5) = 4000$  sorted runs.
- Passes to completely sort the file
- Number of passes =  $\lceil \log_{B-1} N_1 \rceil + 1$  where  $N_1 = \lceil N/B \rceil$   
 $\lceil \log_4(4000) \rceil + 1 = 7$  passes
- Total number of page I/Os =  $2 * N * (\text{\#passes})$   
 $= 2 * 20000 * 7 = 280000$



# Question 1

---

Consider a hard disk with 5400 RPM and 2 double-sided platters. Each surface has 300,000 tracks and 3000 sectors per track. (For simplicity, we assume that the number of sectors per track does not vary between the outer and inner area of the disk.) Each sector has 512 bytes. Suppose a page size of 4096 bytes is chosen. A file containing 100,000 records of 400 bytes each is to be stored on such a disk and that no record is allowed to span two blocks.



# Question 1

---

- What is the capacity(max number of bytes) of the disk?
- How many records fit in a page?
- How many I/Os are required to read the entire file sequentially?
- How many records of 100 bytes each can be stored using this disk?

Given:

Page size = 4096 bytes

Sector size = 512 bytes

Number of records = 100,000

Size of a record = 400 bytes

# of tracks/surface = 300,000

# of sectors/track = 3000

# of bytes/sector = 512 bytes

# of surfaces/disk = 4

bytes/track = bytes/sector  $\times$  sectors/track =  $512 \times 3000 = 1,536,000$

bytes/surface = bytes/track  $\times$  tracks/surface =  $1536K \times 300K = 460,800 \times 10^6$

bytes/disk = bytes/surface  $\times$  surfaces/disk =  $460,800 \times 10^6 \times 4 = 1843200 \times 10^6$

Number of blocks in a track =  $1536000 / 4096 = 375$

1) What is the capacity of the disk?

A)  $1843200 \times 10^6$  bytes

2) How many records fit into a page?

A Size of page/ size of record =  $4096/400 > 10$ . So 10 records can fit into a block



3) How many I/Os are required to read the entire file sequentially?

A  $100,000/10 = 10000$  I/Os

4) How many records of 100 bytes each can be stored using this disk?

Number of pages in a disk = Disk capacity / 4096 bytes =  $1843200 * 10^6 / 4096 = 450 * 10^6$  pages


Number of records that can fit in a page =  $4096/100 > 40$

Total number of records that can fit =  $40 * 450 * 10^6 = 18 * 10^9$



## Question 2


Consider the following B+-tree, and assume the convention that the left pointer points to values that are strictly less than the key value. Each node in this B+-tree can hold up to four entries (i.e., the order of the tree is 2).





# Question 2 - Solution

Inserting 12








# Question 2 - Solution


Inserting 33





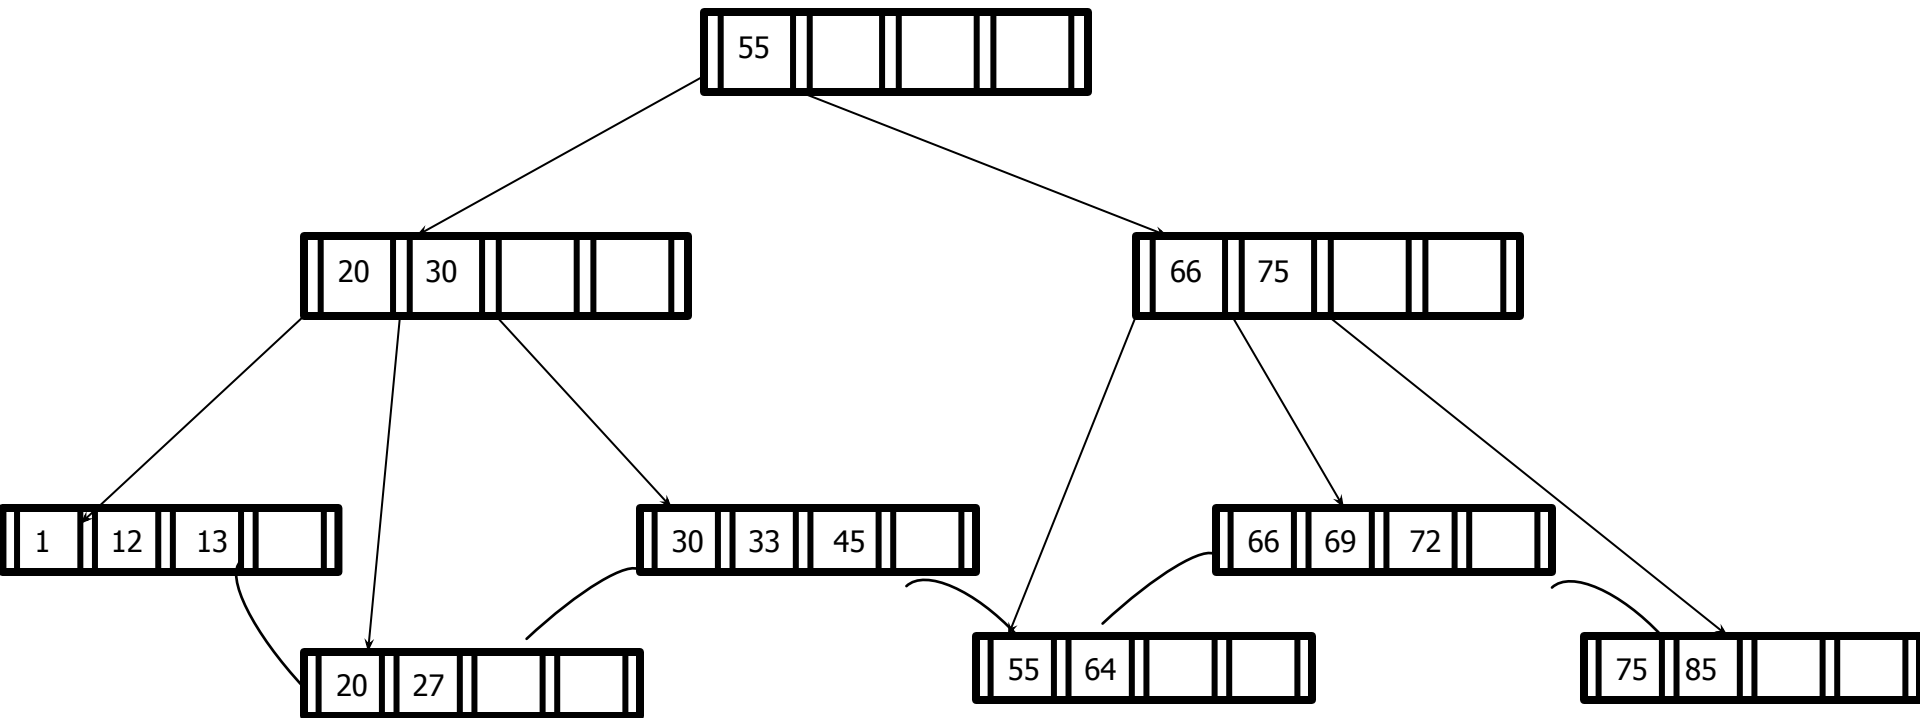
# Question 2 - Solution

Inserting 69



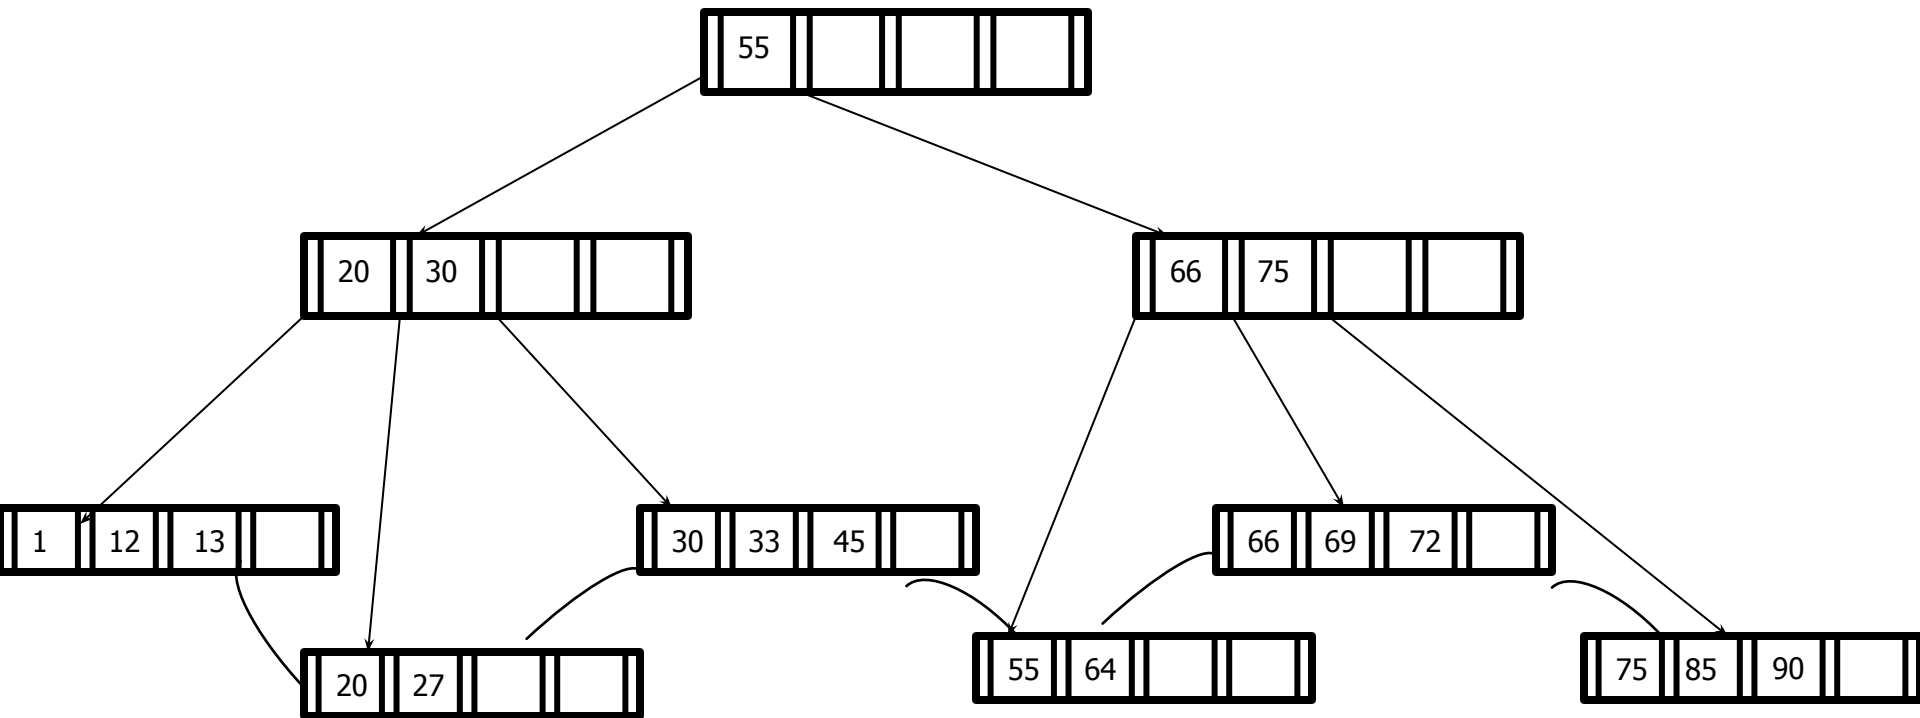
# Solution 2a

Inserting 72 Assume no redistribution of entries.



# Solution 2a


Inserting 90 Assume no redistribution of entries.





# Problem 2b

Delete the data entries with key values 1, 75, 20, 64 (in that order).  
Assuming redistribution of entries.






# Solution 2b

Delete the data entries with key values 1, 75, 20, 64 (in that order).  
Assuming redistribution of entries.

DELETING 1






# Solution 2b

Delete the data entries with key values 1, 75, 20, 64 (in that order).  
Assuming redistribution of entries.

DELETING 75






# Solution 2b

Delete the data entries with key values 1, 75, 20, 64 (in that order). Assuming redistribution of entries.

DELETING 20








# Solution 2b

Delete the data entries with key values 1, 75, 20, 64 (in that order).  
Assuming redistribution of entries.

DELETING 64






# Question 3(a)

- Consider an extendible hash index that uses the following hash function: **hashvalue = least significant bits**, where  $d$  is the global or local depth of the index. Assume that for this index the bucket capacity is 2 entries.

Assume that the index is initially empty. Then two entries with key values **16** and **23** are inserted. Draw the resulting index. *Clearly show the global and local depths and all bucket pointers.*




After adding 16



# Question 3(a)

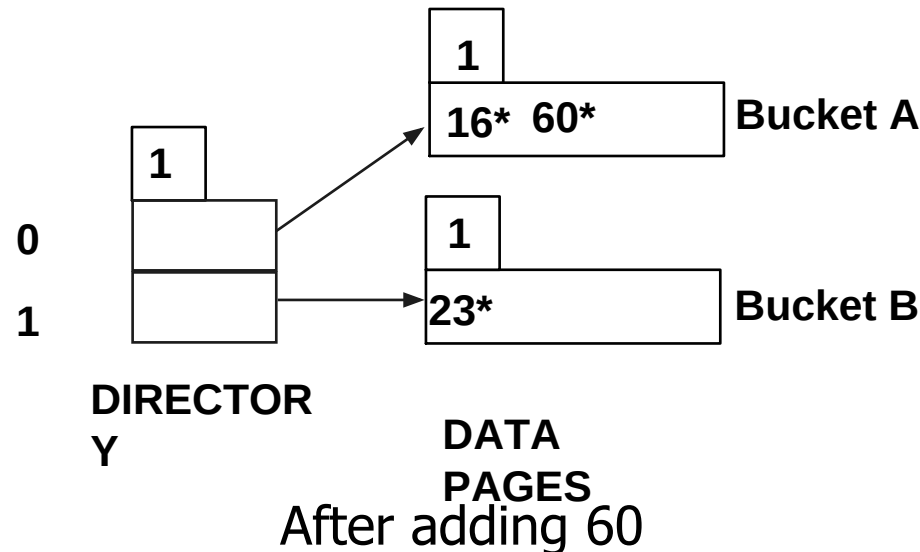
After adding 23





## Question 3(b)

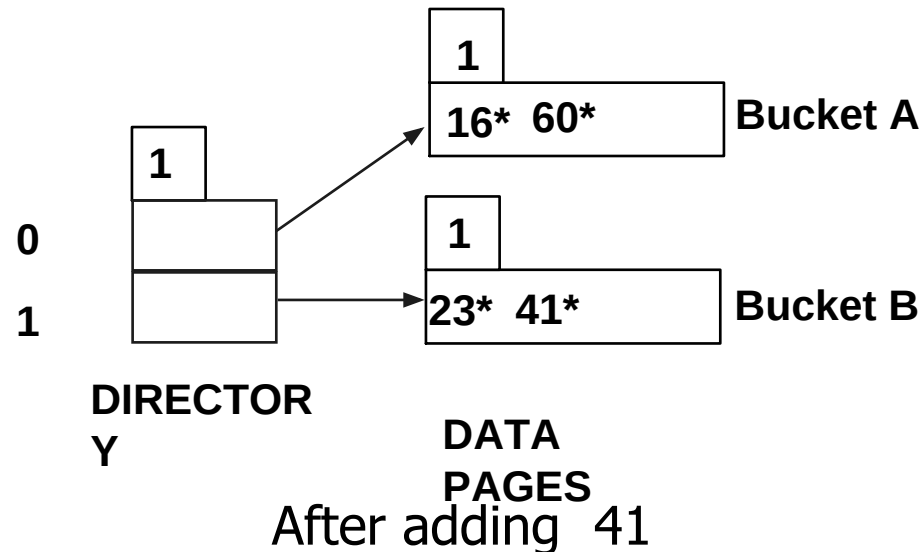
- After inserting the keys 16 and 23 above, the following three additional keys are inserted into the index: **60, 41, 26, 37, 28**. Draw the final index structure with all the five keys. Clearly show the global and local depths, and all bucket pointers.





## Question 3(b)

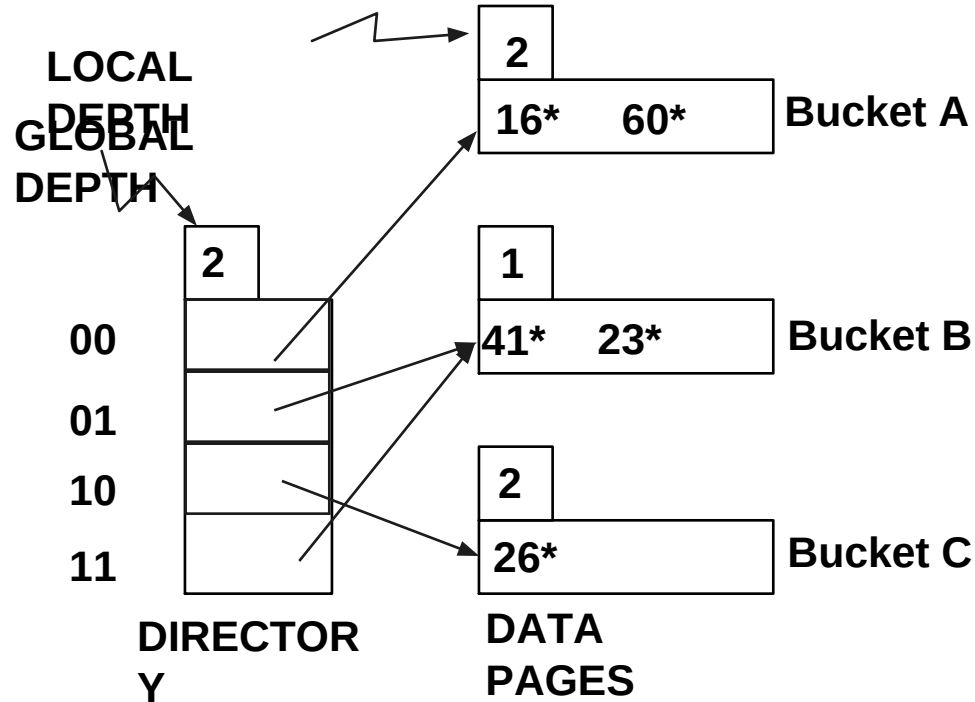
- After inserting the keys 16 and 23 above, the following three additional keys are inserted into the index: **60, 41, 26, 37, 28**. Draw the final index structure with all the five keys. Clearly show the global and local depths, and all bucket pointers.





## Question 3(b)


- After inserting the keys 16 and 23 above, the following three additional keys are inserted into the index: **60, 41, 26, 37, 28**. Draw the final index structure with all the five keys. Clearly show the global and local depths, and all bucket pointers.





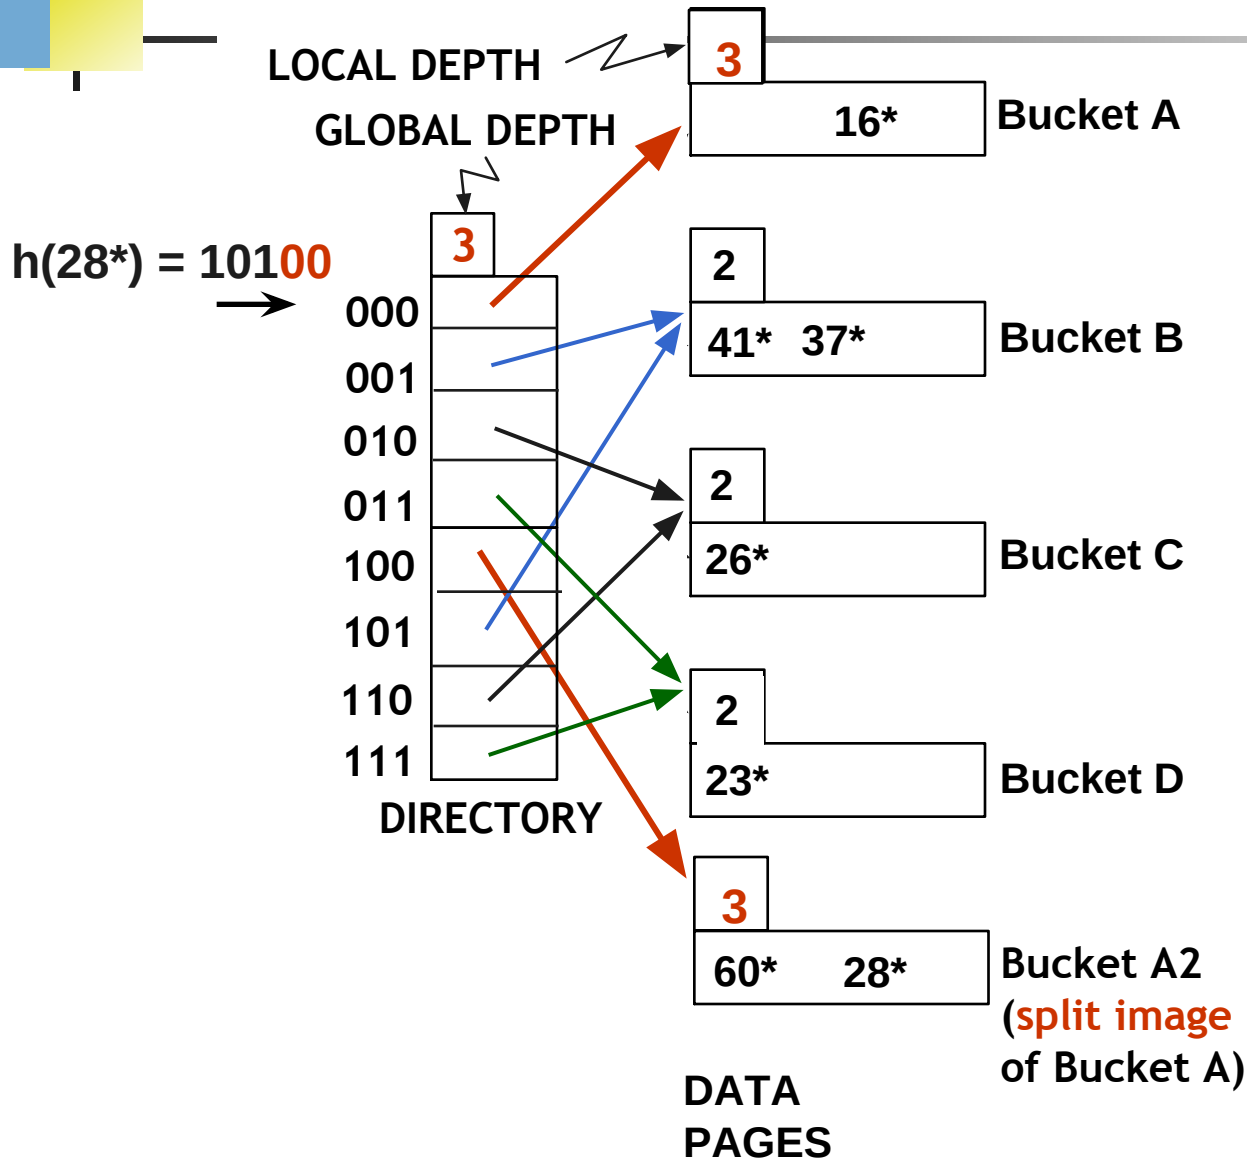
# Question 3(b)

- After inserting the keys 16 and 23 above, the following three additional keys are inserted into the index: **60, 41, 26, 37, 28**. Draw the final index structure with all the five keys. Clearly show the global and local depths, and all bucket pointers.





## After adding 26







c) What happens when the local depth becomes more than the global depth?

Possible answers:


- The index structure is split.
- This situation is impossible because when a split occurs the directory structure is doubled and then the global and local depth are increased simultaneously.




# Question 4

H Level 0,  $N = 2$

0 Next=0



0



1


Primary bucket Pages in the Hash File



# Question 4 - Solution

Inserting 43

H Level 0, N = 2




Primary bucket Pages in the Hash File



# Question 4 - Solution

Inserting 36

H Level 0, N = 2



Primary bucket Pages in the Hash File



# Question 4 - Solution

Inserting 21


$H_1$

H

Level 0, N = 2

0


00



0


Next=1

01



1

10



0

11

1



# Question 4 - Solution

Inserting 26


$H_1$

H

Level 0, N = 2

0


00



0


Next=1

01



1

10



0

11

1



# Question 4 - Solution

Inserting 54

