



Logical Database Design: Mapping ER to Relational

Chapter 3, Section 3.5



ER Model vs. Relational Model


- ER Model used for conceptual design
- Relational Model implemented by modern DBMS
- Important Step: Translate ER diagram to Relational schema



Recall ER Constructs

- Basic Constructs
 - Entity Sets
 - Relationship Sets
 - Attributes (of entities and relationships)
- Additional Constructs
 - ISA Hierarchies
 - Weak Entities
 - Aggregation
- Integrity Constraints
 - Key constraints
 - Participation constraints
 - Overlap / Covering constraints for ISA hierarchies


Entity Sets to Tables



```
CREATE TABLE Citizens
  (ssn CHAR(11),
   name CHAR(20),
   bday DATE,
   PRIMARY KEY (ssn))
```

Can ssn have
a null value?

Relationship Sets to Tables



Relationship set -> Table

Attributes:


- Participating entity set primary keys
 - Foreign key
- Descriptive attributes

```
CREATE TABLE Votes(
    ssn    CHAR(11),
    cid    INTEGER,
    when   DATE,
    PRIMARY KEY (ssn, cid),
    FOREIGN KEY (ssn) REFERENCES Citizens,
    FOREIGN KEY (cid) REFERENCES PR-Cands)
```


Generalizes to
n-ary
relationships
(we will see
example later)

Relationship Sets to Tables

```
CREATE TABLE Represents(  
    elected_ssn CHAR(11),  
    cons_ssn CHAR(11),  
    PRIMARY KEY (elected_ssn, cons_ssn),  
    FOREIGN KEY (elected_ssn) REFERENCES Citizens,  
    FOREIGN KEY (cons_ssn) REFERENCES Citizens)
```



Key Constraints




- Approach 1 – Three Tables

```
CREATE TABLE Votes
(  ssn    CHAR(11),
   cid    INTEGER NOT NULL,
   when   DATE,
   PRIMARY KEY (ssn),
   FOREIGN KEY (ssn) REFERENCES Citizens,
   FOREIGN KEY (cid) REFERENCES PR-Cands)
```

Same approach
as before –
Map each entity
set and relationship
set to a table

Key Constraints



- Approach 2 – Two Tables

```
CREATE TABLE Citizen_Votes (  
  ssn    CHAR(11), name CHAR(20),  
  bday   DATE,      when DATE,  
  cid    INTEGER,  
  PRIMARY KEY (ssn),  
  FOREIGN KEY (cid) REFERENCES PR-Cands)
```


Each citizen can only vote once, so OK to fold 'Votes' relationship into 'Citizens' entity

Q: Can cid be null?

Q: What if many citizens don't vote?

Q: Which approach is better?

Key Constraints



What about 1 Table?

No! This is bad design.

e.g., For each citizen that voted for a candidate, we would be storing candidate's information (cid, waddr, budget)

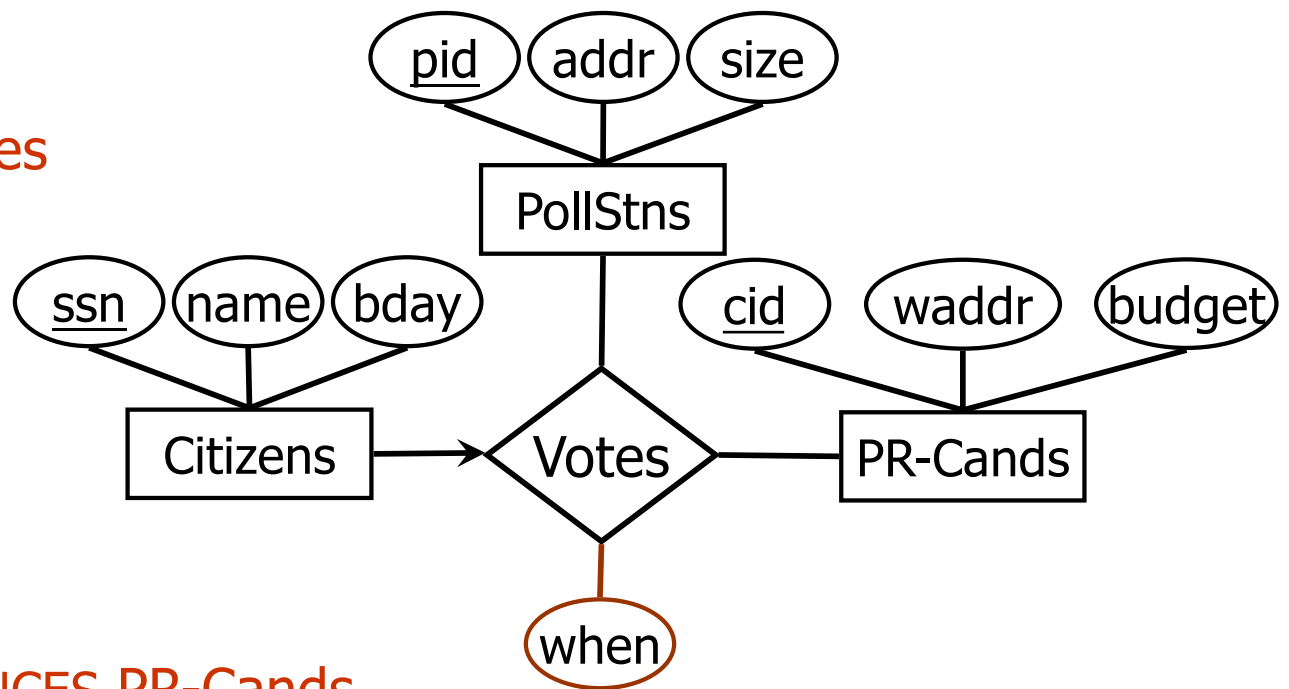
=> **REDUNDANCY!**

Key Constraints


Can generalize to n-ary relationships

CREATE TABLE Citizens_Votes

```
( ssn    CHAR(11),  
  name   CHAR(20),  
  bday   DATE,  
  when   DATE,  
  pid    INTEGER,  
  cid    INTEGER,  
  PRIMARY KEY (ssn),  
  FOREIGN KEY (cid) REFERENCES PR-Cands,  
  FOREIGN KEY (pid) REFERENCES PollStns)
```



Participation Constraints




Using Approach 2

```
CREATE TABLE Citizen_Votes(  
    ssn    CHAR(11),  
    name   CHAR(20),  
    bday   DATE,  
    when   DATE,  
    cid     INTEGER NOT NULL,  
    PRIMARY KEY (ssn),  
    FOREIGN KEY (cid) REFERENCES PR_Cands,  
    ON DELETE NO ACTION)
```

Can we enforce participation constraint using Approach 1 (three tables) ?


Mapping Participation Constraints



```
CREATE TABLE RAB(  
  r1 Integer,  
  a1 Integer,  
  a2 Integer,  
  a3 Integer,  
  b1 Integer,  
  b2 Integer,  
  b3 Integer ...)
```


Key constraints?

Mapping Participation Constraints



```
CREATE TABLE RAB(  
    r1 Integer,  
    a1 Integer,  
    a2 Integer,  
    a3 Integer,  
    b1 Integer NOT NULL,  
    b2 Integer,  
    b3 Integer,  
    UNIQUE (b1), PRIMARY KEY (a1))
```

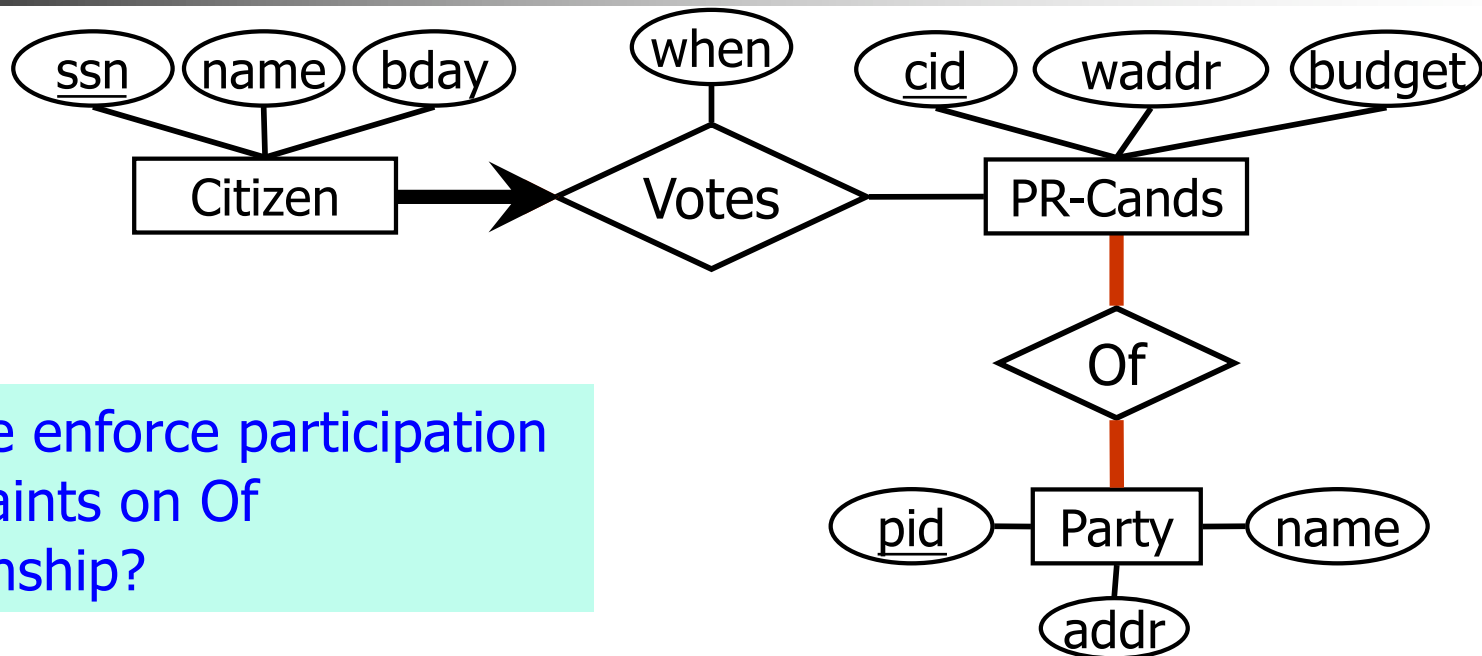
Weak Entities



- Use approach 2: Combine weak entity and owning relationship into one relation
 - Delete all weak entities when an owner entity is deleted.

```
CREATE TABLE Expert_Staff (  
  ename    CHAR(20),  
  jobDesc  CHAR(40),  
  salary   REAL,  
  cid      INTEGER,  
  PRIMARY KEY (ename, cid),  
  FOREIGN KEY (cid) REFERENCES PR-Cands ON DELETE CASCADE)
```

Participation Constraints



Can we enforce participation constraints on Of relationship?

Need table constraints & assertions – Later.


ISA Hierarchies – General Approach

- Three relations:

- Citizens (ssn, name, bday)
- PR-Cands (ssn, waddr, budget)
- Presidents (ssn, from, to)


- Queries

- Involving all citizens => Easy
- Involving just PR-Cands, need to join PR-cands with Citizens to get some attributes



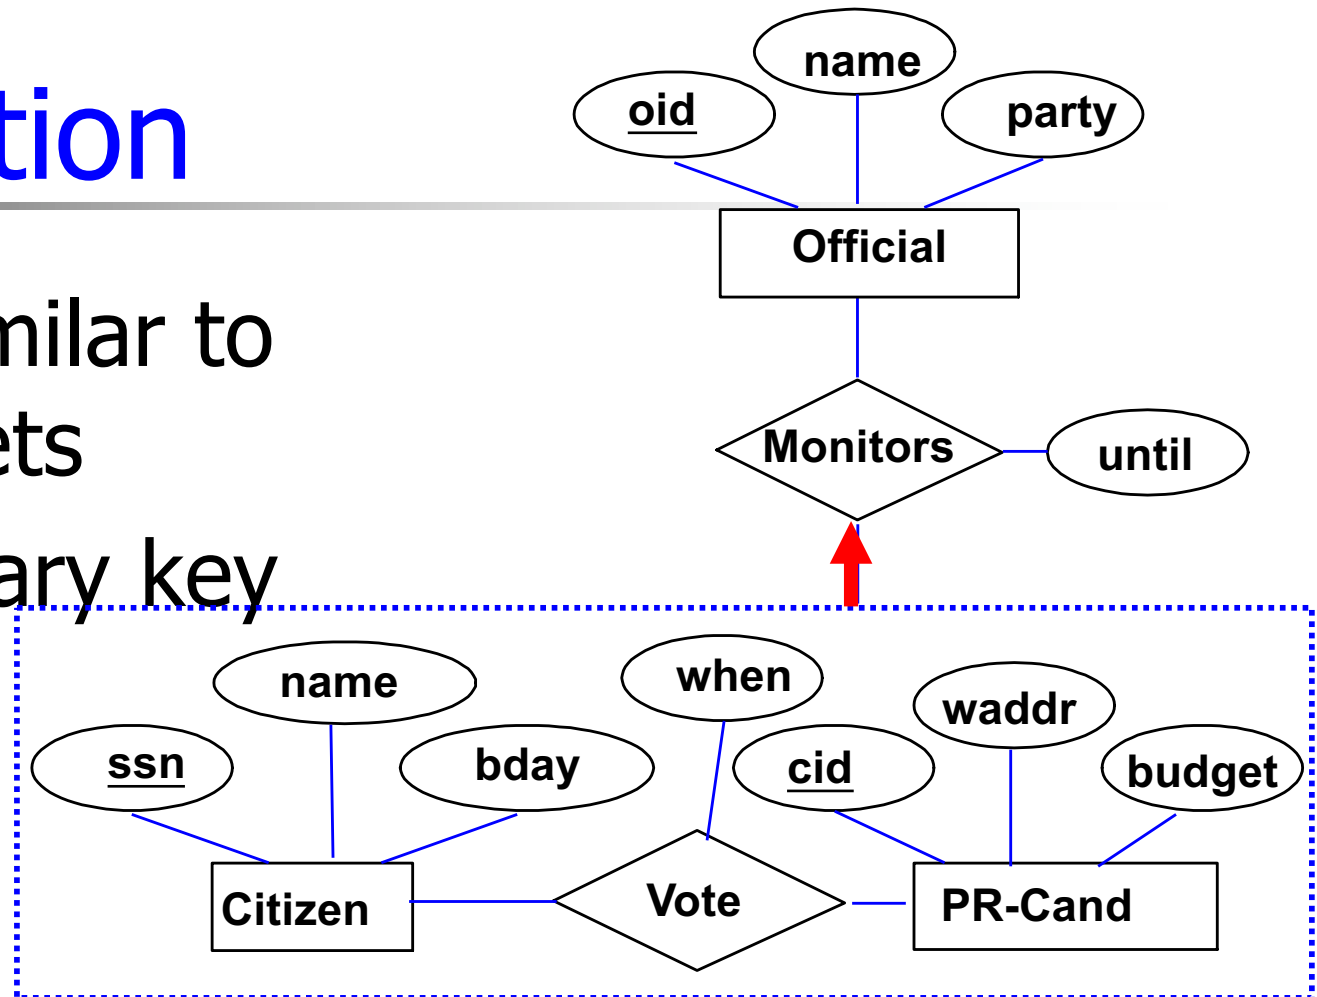
ISA Hierarchies – Alternative

- Two relations:
 - PR-Cands (ssn, name, bday, waddr, budget)
 - Presidents (ssn, name, bday, from, to)
- Problems
 - What if citizen is both?
 - Redundancy
 - What if citizen is neither?
 - Use General approach



Aggregation

- Translation similar to relationship sets
- Monitors primary key
 - (ssn, cid, oid)



- What if every Vote must have exactly one monitor?

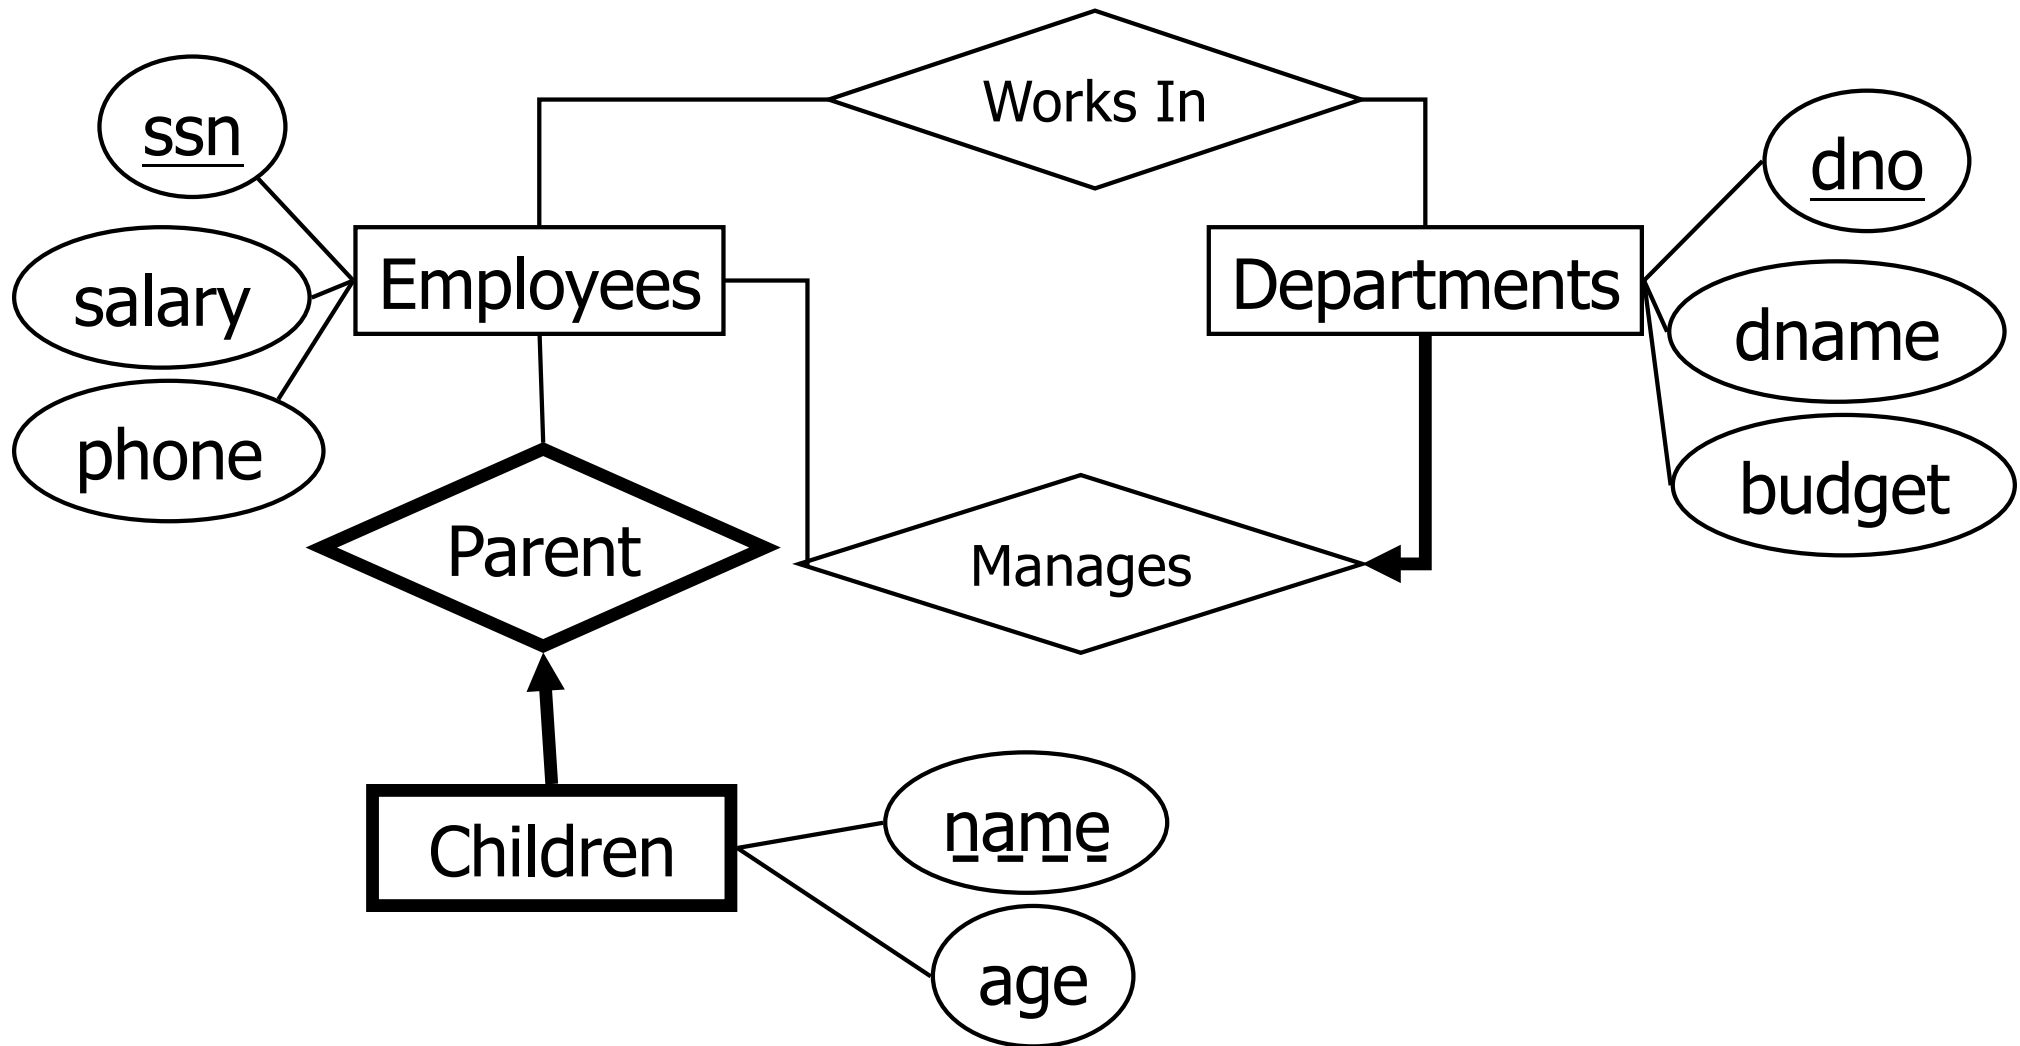
Fold Vote and
Monitors into 1 table



Exercise – Part 1

- A company database needs to store information about
 - employees (identified by ssn, with *salary* and *phone* attributes),
 - departments (identified by dno, with *dname* and *budget* attributes), and
 - children of employees (with *name* and *age* attributes).
- Employees work in (zero or more) departments;
- Each department is managed by exactly one employee
- A child must be identified uniquely by name when the parent (who is an employee; assume only one parent works for the company) is known.
- We are not interested in information about a child once the parent leaves the company.
- Draw an ER diagram that captures this information

ER Diagram (One Solution)





Exercise – Part 2

- Write SQL Statements to create the corresponding relations, and to capture as many of the constraints as possible.



SQL DDL (One Solution)

```
CREATE TABLE employees (  
  ssn INTEGER,  
  salary REAL,  
  phone CHAR(10),  
  PRIMARY KEY(ssn))
```

```
CREATE TABLE departments (  
  dno INTEGER,  
  dname CHAR(20),  
  budget real,  
  manager INTEGER NOT NULL,  
  PRIMARY KEY (dno),  
  FOREIGN KEY (manager)  
    REFERENCES employees)
```

```
CREATE TABLE works (  
  ssn INTEGER,  
  dno INTEGER,  
  PRIMARY KEY (ssn, dno),  
  FOREIGN KEY (ssn)  
    REFERENCES employees,  
  FOREIGN KEY (dno)  
    REFERENCES departments)
```

```
CREATE TABLE children (  
  name CHAR(20),  
  age REAL,  
  parent INTEGER NOT NULL,  
  PRIMARY KEY(name, parent),  
  FOREIGN KEY(parent)  
    REFERENCES employees  
    ON DELETE CASCADE)
```



Integrity Constraints

- Describes conditions that must be satisfied by every legal instance
- Types of integrity constraints
 - Domain constraints
 - Primary key constraints
 - Foreign key constraints
 - **General constraints**



Table Constraints (5.7 in book)

- More general than key constraints
- Can use a query to express constraint
 - Constraints checked each time table updated
 - CHECK constraint always true for empty relation

```
CREATE TABLE Sailors
( sid INTEGER,
  sname CHAR(10),
  rating INTEGER,
  age REAL,
  PRIMARY KEY (sid),
  CHECK ( rating >= 1
        AND rating <= 10)
);
```




Try it out in sqlplus or sqlite

```
CREATE TABLE Sailors
  ( sid INTEGER,
    sname CHAR(10),
    rating INTEGER,
    age REAL,
    PRIMARY KEY (sid),
    CHECK (rating >= 1 AND rating <= 10));
```

```
INSERT INTO Sailors VALUES (1, 's1', 11, 25);
```

Do you get a constraint violation error?



More general CHECK

- Interlake boats cannot be reserved. **Note: these are not supported in Oracle or SQLite**

```
CREATE TABLE Reserves
( sname CHAR(10),
  bid INTEGER,
  day DATE,
  PRIMARY KEY (bid,day),
  CONSTRAINT noInterlakeRes
  CHECK (`Interlake' NOT IN
        ( SELECT B.bname
          FROM Boats B
          WHERE B.bid=bid)))
```



Constraints Over Multiple Relations

- For general constraint over multiple tables, SQL standard provides **assertions**.

*Number of boats
plus number of
sailors is < 100*

```
CREATE ASSERTION smallClub  
CHECK  
((SELECT COUNT (S.sid) FROM Sailors S) +  
 (SELECT COUNT (B.bid) FROM Boats B) < 100)
```



Practical Considerations

- CHECK with subqueries and ASSERTIONS
 - Part of SQL standard (since 1992?)
 - But, major databases do not support them
 - Main concern: performance.
- Instead, most major database systems require you to use *triggers* to achieve the same goals. Triggers are procedural.



Active Databases & Triggers

- **Trigger:** procedure that starts automatically if specified changes occur to the DBMS
- Three parts:
 - Event (activates the trigger)
 - Condition (tests whether the trigger should run)
 - Action (what happens if the trigger runs)
 - Before and After Triggers
- Trigger Execution
 - Row-level Triggers: Once per modified row
 - Statement-level Triggers: Once per SQL statement



Example

- Student table with columns: sid, name, age.
- Event: Insert new row into Students
- Condition: New student's age ≥ 18
- Event: Update a counter variables that tracks university's total enrollment



Oracle Trigger Basic Syntax

```
CREATE [OR REPLACE] TRIGGER <trigger_name>
{BEFORE | AFTER} {INSERT | DELETE | UPDATE} ON <table_name>
[REFERENCING [NEW AS <new_row>] [OLD AS <old_row>]]
[FOR EACH ROW [WHEN (<condition>)]]
<trigger_body>
```

In Oracle,
<trigger_body>
is a PL/SQL
block

Oracle Triggers
Not quite the same
As SQL Standard



Triggers – Use for Auditing

```
CREATE TABLE Score(uniquename CHAR(10), grade NUMBER);  
CREATE TABLE ScoreLog(uniquename CHAR(10) , gradeprev NUMBER, gradenew NUMBER);
```

```
CREATE OR REPLACE TRIGGER trig1 AFTER UPDATE ON Score  
FOR EACH ROW  
  WHEN (:NEW.grade <> :OLD.grade)  
  BEGIN  
    INSERT INTO ScoreLog VALUES(:NEW.uniquename, :OLD.grade, :NEW.grade);  
  END;
```

```
.  
RUN
```

```
INSERT INTO Score VALUES('s1', 23);  
INSERT INTO Score VALUES('s2', 50);  
UPDATE Score Set grade = 100 WHERE uniquename = 's1';  
SELECT * FROM Score;  
SELECT * FROM ScoreLog;
```

Try it Out!



To see compilation errors

- Use the following:

`SHOW ERRORS TRIGGER <triggername>`



Triggers – Use for error-checks

```
CREATE OR REPLACE TRIGGER trig2 BEFORE INSERT OR UPDATE ON Score
FOR EACH ROW
  WHEN (NEW.grade < 0 OR NEW.grade > 100)

  BEGIN
    RAISE_APPLICATION_ERROR(-20001, 'Bad grade');
  END;

.
```

RUN

- Aborts an INSERT or UPDATE command if grade is not in range [0,100]. Try doing:

```
INSERT INTO Score VALUES('s3', -10);
```



Triggers – Fix inserted data

```
CREATE OR REPLACE TRIGGER trig3 BEFORE INSERT OR UPDATE ON Score
FOR EACH ROW
  WHEN (NEW.grade > 100)
  BEGIN
    :NEW.grade := 100;
  END;
.
```

RUN

```
INSERT INTO Score VALUES('s4', 200);
```

What value is inserted into Score?



Simulate MySQL's AUTO_INCREMENT

```
CREATE TABLE Person  
(  
  id INT NOT NULL AUTO_INCREMENT,  
  last VARCHAR(24) NOT NULL,  
  first VARCHAR(24)  
)
```

```
INSERT INTO Person VALUES('john', 'doe');
```

AUTO_INCREMENT is **NOT** available in Oracle.
See Discussion Slides on how to use triggers in Oracle
to accomplish equivalent this.



Oracle Trigger Example Contd.

- First trigger executed before the activating statement, second executes after the activating statement.
- In combination with:
 - “FOR EACH ROW” - execute once per modified record
 - (default) - execute once per activating statement.
- Activating statements can be:
 - “INSERT”
 - “DELETE”
 - “UPDATE”



Recall CASCADE constraints

```
CREATE TABLE Athlete  
(aid INTEGER PRIMARY KEY,  
name CHAR(30),  
country CHAR(20),  
sport CHAR(20));
```

```
CREATE TABLE Olympics  
(oid INTEGER PRIMARY KEY,  
year INTEGER,  
city CHAR(20));
```

```
CREATE TABLE Compete  
(aid INTEGER, oid INTEGER,  
PRIMARY KEY (aid, oid),  
FOREIGN KEY (aid)  
REFERENCES Athlete  
ON DELETE CASCADE  
);
```



CASCADE Using Triggers

```
CREATE TABLE Compete
(aid INTEGER, oid INTEGER,
PRIMARY KEY (aid, oid),
FOREIGN KEY (aid)
REFERENCES Athlete );
```

```
CREATE OR REPLACE TRIGGER cascade_on_delete
AFTER DELETE ON Athlete
FOR EACH ROW
BEGIN
DELETE FROM Compete
WHERE Compete.aid = :OLD.aid;
END;
/
```



Trying out triggers

- Play with triggers in SQLite and Oracle.
- To drop the trigger:
 - DROP TRIGGER *triggername*;



Triggers: Pitfalls and Pain

- Triggers can trigger other triggers
 - Update to table A causes update to table B
 - Update to table B causes update to table C
- Even cycles can occur
- Debugging difficult
- You may see an error with “mutating” table
 - Typically, due to a SELECT on a table in the middle of being modified inside a trigger



Triggers vs. Constraints

- Both used to maintain consistency
- Constraints are easier to understand than triggers
 - Always prefer them when you have a choice
- Triggers are more powerful but harder to debug