



Relational Algebra

Chapter 4



“Formal” Query Languages

- Foundation for commercial query languages like SQL
- Two types
 - Declarative: **Relational Calculus**
 - Describe what a user wants, rather than how to compute it.
 - Procedural : **Relational Algebra**
 - Operational, very useful for representing execution plans.
- Query Languages **!=** programming languages!
 - QLs not expected to be “Turing complete”.
 - QLs not intended to be used for complex calculations.
 - QLs support easy, efficient access to large data sets.

Understanding Algebra & Calculus is key to understanding SQL, query processing!



Relational Algebra Preliminaries

- Query:

- Input: Relational instances
- Output: Relational instances!

Relational Algebra is
“closed”

- Specified using the schemas.
 - May produce different results for different instances.
 - But schema of the result is fixed.
- The algebra assumes “set semantics” for relations. OK – this is a formal model!



Relational Algebra

- Basic operations on relations:
 - *Selection* (σ) Selects a subset of rows from relation.
 - *Projection* (π) Deletes unwanted columns from relation.
 - *Cross-product* (\times) Allows us to combine two relations.
 - *Set-difference* ($-$) Tuples in reln. 1, but not in reln. 2.
 - *Union* (\cup) Tuples in reln. 1 and in reln. 2.
- Additional operations (constructed from basic ops):
 - Intersection, *Join*, Division, Renaming
 - Not essential, but (very!) useful.
- Because algebra is closed, we can compose operators

Selection

- Retrieve rows that satisfy a logical condition

$$\sigma_{\text{predicate}}(\text{relation})$$

Example:

$$\sigma_{\text{sport}='gymnastics' \wedge \text{country}='USA'}(\text{Athlete})$$

Athlete

aid	name	sport	country
1	Mary Lou Retton	gymnastics	USA
2	Jackie Joyner-Kersey	track	USA
3	Michael Phelps	swimming	USA
4	Johann Koss	skating	Norway
5	Natalie Coughlin	swimming	USA
6	Gabby Douglas	gymnastics	USA

aid	name	sport	country
1	Mary Lou Retton	gymnastics	USA
6	Gabby Douglas	gymnastics	USA

Projection

$\pi_{projectionlist}(Relation)$

Delete attributes that are not in projection list

- *Projectionlist*: a list of columns
- The result is a set (relational algebra uses set semantics)
- Remove duplicates!

Athlete

aid	name	sport	country
1	Mary Lou Retton	gymnastics	USA
2	Jackie Joyner-Kersey	track	USA
3	Michael Phelps	swimming	USA
4	Johann Koss	skating	Norway
5	Natalie Coughlin	swimming	USA
6	Paul Hamm	gymnastics	USA

$\pi_{sport, country}(Athlete)$

sport	country
gymnastics	USA
track	USA
swimming	USA
skating	Norway

Set Operations: Union (\cup), Intersection(\cap), Set-Difference ($-$)

- Input: Two union-compatible relations
 - Same number and type of attributes, in same order
- Field names of result: uses the name from the FIRST input

name	Mary Lou Retton	Jackie Joyner-Kersee	Michael Phelps	Johann Koss	Natalie Coughlin	Gabby Douglas
------	-----------------	----------------------	----------------	-------------	------------------	---------------

AthleteName	Mary Lou Retton	Jackie Joyner-Kersee	Apollo Ono	Picabo Street	Natalie Coughlin	Bode Miller
-------------	-----------------	----------------------	------------	---------------	------------------	-------------

name	Michael Phelps	Johann Koss	Gabby Douglas
------	----------------	-------------	---------------

Duplicates? Relational algebra uses set semantics.
So no duplicates. Difference from SQL

Cross-Product (Cartesian Product) X

- Result Schema
 - One field from both relations (Names inherited)
- If both input relations have a field with the same name, can use the rename operator ρ . (See 4.2.2 and 4.2.3 in textbook)

Athlete

aid	name	sport	country
1	Mary Lou Retton	gymnastics	USA
2	Jackie Joyner-Kersey	track	USA

X

Venue

vid	venue
1	Los Angeles
2	Barcelona

=

aid	name	sport	country	vid	venue
1	Mary Lou Retton	gymnastics	USA	1	Los Angeles
2	Jackie Joyner-Kersey	track	USA	1	Los Angeles
1	Mary Lou Retton	gymnastics	USA	2	Barcelona
2	Jacki Joyner-Kersey	track	USA	2	Barcelona



Derived Operators: Joins

- Most common way of combining information from two tables
- Conditional join (sometimes called a Θ -join)
 - Definition: $R \bowtie_c S = \sigma_c(R \times S)$, where c is a condition
- Equijoin
 - Join condition consists only of equalities
- Natural Join
 - Equijoin in which equalities are specified on all fields with the same name in R and S
- Despite equivalence, usually faster ways to evaluate joins than to compute cross-product!



Examples: Writing Queries in RA

Example Schema:

Sailors (sid, sname, rating, age)

Reserves (sid, bid, day)

Boats (bid, bname, color)

Sailors

sid	sname	rating	age
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
32	Andy	8	25.5
58	Rusty	10	35.0
64	Horatio	7	35.0
71	Zorba	10	16.0
74	Horatio	9	35.0
85	Art	3	25.5
95	Bob	3	63.5

Reserves

sid	bid	day
22	101	10/10/98
22	102	10/10/98
22	103	10/8/98
22	104	10/7/98
31	102	11/10/98
31	103	11/6/98
31	104	11/12/98
64	101	9/5/98
64	102	9/8/98
74	103	9/8/98

Boats

bid	bname	color
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red

Find names of sailors who've reserved boat #103

Solution 1: (1) Extract reservations for boat ID 103
(2) Join with sailors and project on sname

Solution 2: Same as 1, but give temp names to intermediate results

Solution 3: (1) Join Reserves and Sailors
(2) Select on bid = 103
(3) project on sname

Find names of sailors
who've reserved a red
boat

$$\pi_{sname}((\sigma_{color='red'}Boats) \bowtie Reserves \bowtie Sailors)$$

An equivalent solution:

$$\pi_{sname}(\pi_{sid}((\pi_{bid}\sigma_{color='red'}Boats) \bowtie Res) \bowtie Sailors)$$

Query optimizer chooses from the (equivalent) expressions and chooses one for efficiency of evaluation.

Find the names of sailors
who've reserved at least
one boat

$\pi_{sname}(\text{Reserves} \bowtie \text{Sailors})$

Sailor appears in this
intermediate relation
only if there is at least
one Reserves tuple with
same sid.



Derived Operators: Division

- Useful for queries like:

*Find customers with accounts at **all** branches in Brooklyn.*

- Let A have 2 fields, x and y ; B have only field y :
 - $A/B = \{ \langle x \rangle \mid \forall \langle y \rangle \in B, \langle x, y \rangle \in A \}$
 - **A/B contains all x tuples (customers) such that for every y tuple (branches in Brooklyn) in B , there is an $\langle x, y \rangle$ tuple in A**
- In general, x and y can be lists of fields; y is the list of fields in B , and $x \cup y$ is the list of fields of A .

Examples of Division A/B

cid	bname
c1	b1
c1	b2
c1	b3
c1	b4
c2	b1
c2	b2
c3	b2
c4	b2
c4	b4

A

bname
b2

B1

cid
c1
c2
c3
c4

A/B1

bname
b2
b4

B2

cid
c1
c4

A/B2

bname
b1
b2
b4

B3

cid
c1

A/B3

Expressing A/B Using Basic Operators

- Can be equivalently expressed using basic operators
- **Idea:** For A/B, compute all x values that are not disqualified by some y value in B.
 - x value is **disqualified** if by attaching y value from B, we obtain an $\langle x, y \rangle$ tuple that is not in A

Can you express this operator using basic operators?

$$A/B: \pi_x(A) - \underbrace{\pi_x((\pi_x(A) \times B) - A)}_{\text{Disqualified } x \text{ values}}$$

Examples of Division A/B

cid	bname
c1	b1
c1	b2
c1	b3
c1	b4
c2	b1
c2	b2
c3	b2
c4	b2
c4	b4

A

bname
b2
b4

B

cid	bname
c1	b2
c1	b4
c2	b2
c2	b4
c3	b2
c3	b4
c4	b2
c4	b4

$\pi_x(A) \times B$

~~*A*~~

cid
c1
c4

A/B

Find names of sailors who've reserved a red OR green boat

- Identify all red or green boats, then
- find sailors who've reserved one of these boats:

$$\rho(Tempboats, (\sigma_{color='red' \vee color='green'} Boats))$$

$$\pi_{sname}(Tempboats \bowtie Reserves \bowtie Sailors)$$

Equivalent:

$$\rho(Tempboats, (\sigma_{color='red'}(Boats) \cup \sigma_{color='green'}(Boats)))$$

$$\pi_{sname}(Tempboats \bowtie Reserves \bowtie Sailors)$$

Find names of sailors who've reserved a red AND green boat

- Does this work: Change \vee to \wedge on previous slide?
- How about this?
 - Identify Sailor ids who've reserved red boats
 $\rho(TempRed, \pi_{sid}((\sigma_{color='Red'}Boats) \bowtie Reserves))$
 - Identify Sailor ids who've reserved green boats
 $\rho(TempGreen, \pi_{sid}((\sigma_{color='Green'}Boats) \bowtie Reserves))$
 - Then use the intersection to get the names
 $\pi_{sname}((TempRed \cap TempGreen) \bowtie Sailors)$

Find the sids of sailors over
age 20 who have not
reserved a red boat

$$\pi_{sid}(\sigma_{age>20}Sailors) - \pi_{sid}((\sigma_{color='Red'}Boats) \bowtie Reserves)$$

Find the names of sailors who've reserved all boats

- Uses division; schemas of the input relations must be carefully chosen:

$$\rho \text{ (} Tempids, (\pi_{sid, bid} Reserves) / (\pi_{bid} Boats) \text{)}$$
$$\pi_{sname} (Tempids \bowtie Sailors)$$



Suggested Review

- Exercises 4.1, 4.3, 4.5
 - Only RA required for the last two.

Find names of sailors
who've reserved boat #103

Solution 1:

$$\pi_{sname}((\sigma_{bid=103} Reserves) \bowtie Sailors)$$

Solution 2:

$$\rho(Temp1, \sigma_{bid=103} Reserves)$$

$$\rho(Temp2, Temp1 \bowtie Sailors)$$

$$\pi_{sname}(Temp2)$$

Solution 3:

$$\pi_{sname}(\sigma_{bid=103}(Reserves \bowtie Sailors))$$