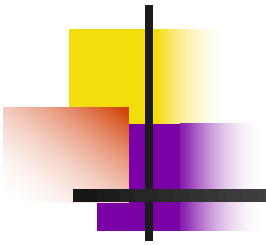

Hash-Based Indexes

Chapter 11

Index Design Space

Organization Structure for k*

- **Hash-based**
  - + Equality search
    - Static hashing
    - Extensible hashing
    - Linear hashing
- **Tree-based**
  - + Range, equality search

Data Entry (k*) Contents

1. Actual Data record
   - index = file
2. <k, rid>
   - actual records in a different file
3. <k, list of rids>

Static Hashing

- # primary bucket pages fixed, allocated sequentially, never de-allocated; overflow pages if needed.
- **h**(*key*) mod N = bucket to which data entry with *key* belongs. (N = # of buckets)

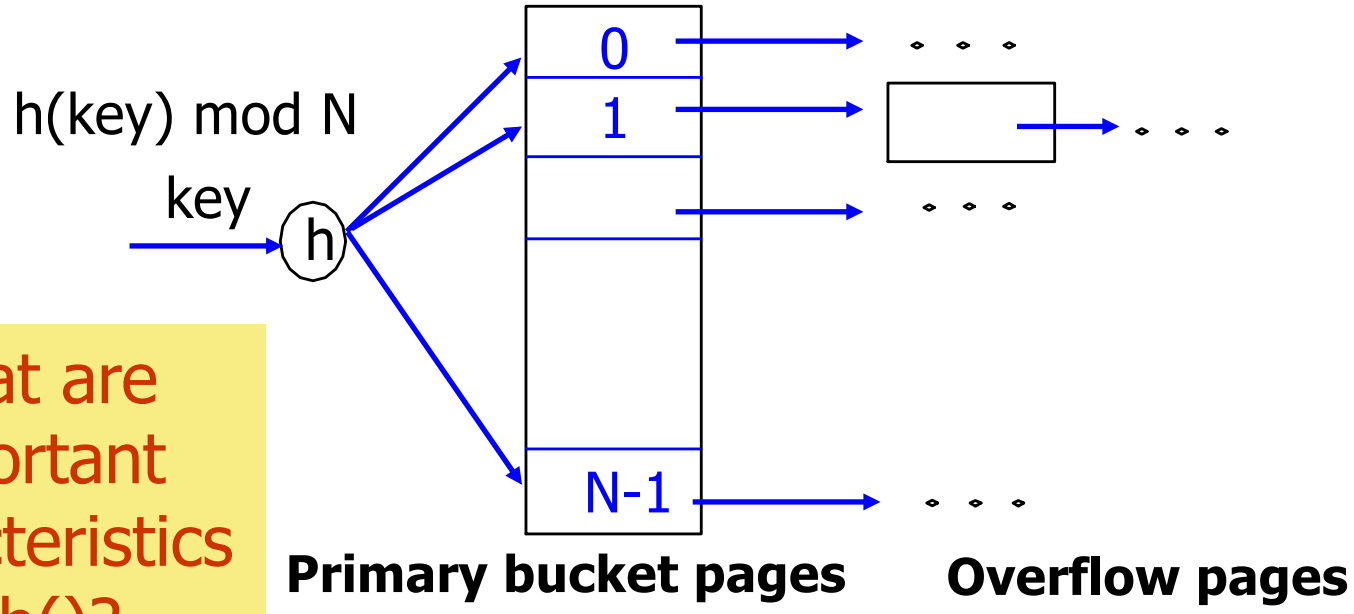

What are important characteristics of h()?

Static Hashing (Contd.)

- Buckets contain *data entries*
- Number of buckets (N) is fixed ahead of time
- Static structure can be problematic
  - Consider many insertions
  - Long overflow chains can develop (and degrade performance!)
- Might consider periodically doubling N and "rehashing" file
  - Entire file has to be read and written
  - Index unavailable while rehashing
  - *Extensible* and *Linear Hashing*: Dynamic techniques to fix this problem.

Extensible Hashing

- **Main Idea:** Use a directory of pointers to buckets
- On overflow, double the directory (not # number of buckets)
- Why does this help?
  - Directory much smaller than file
  - Only one page of data entries is split at a time
  - No overflow pages

Example

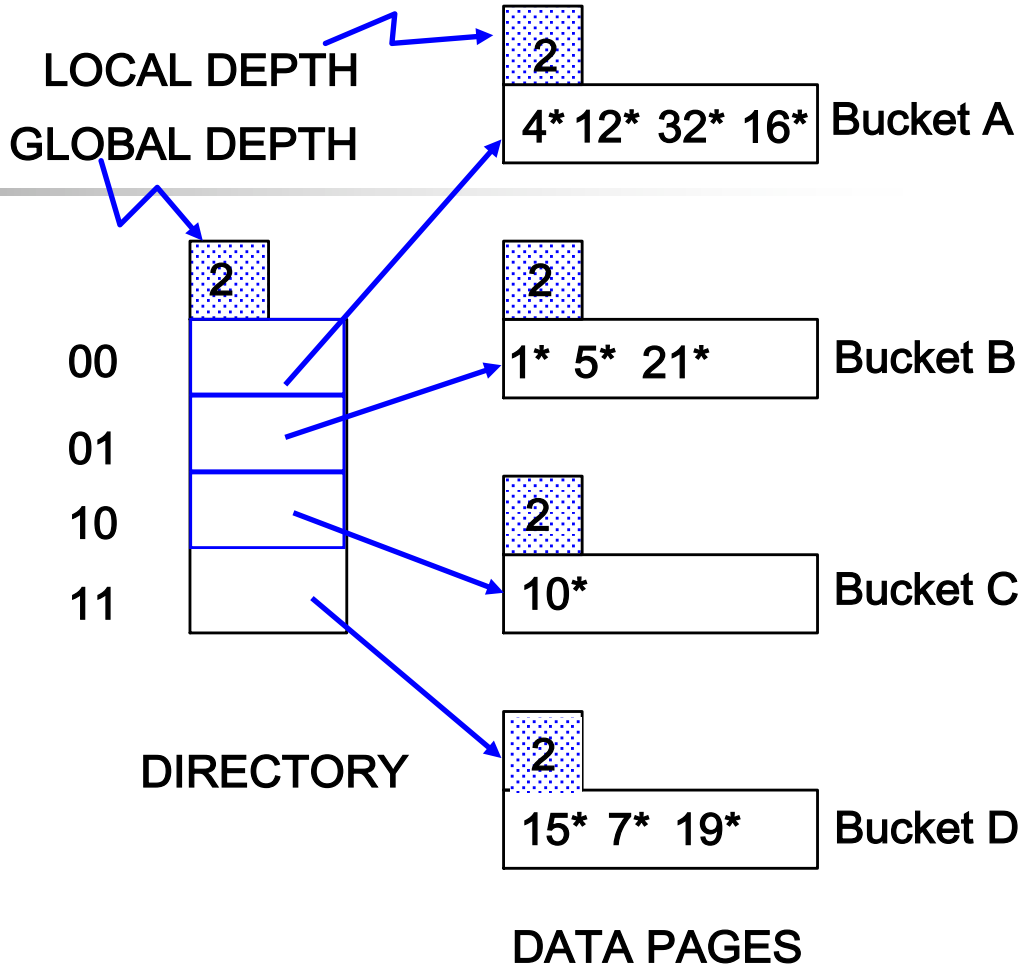

- Directory an array
- Search for k:
  - Apply hash function h(k)
  - Take last *global depth* # bits of **h**(*k*)
- Insert:
  - If bucket has space, insert, done
  - If bucket if full, split it, re-distribute
  - If necessary, double the directory

Example

- Insert 13*
- Suppose h(13*) = 1101

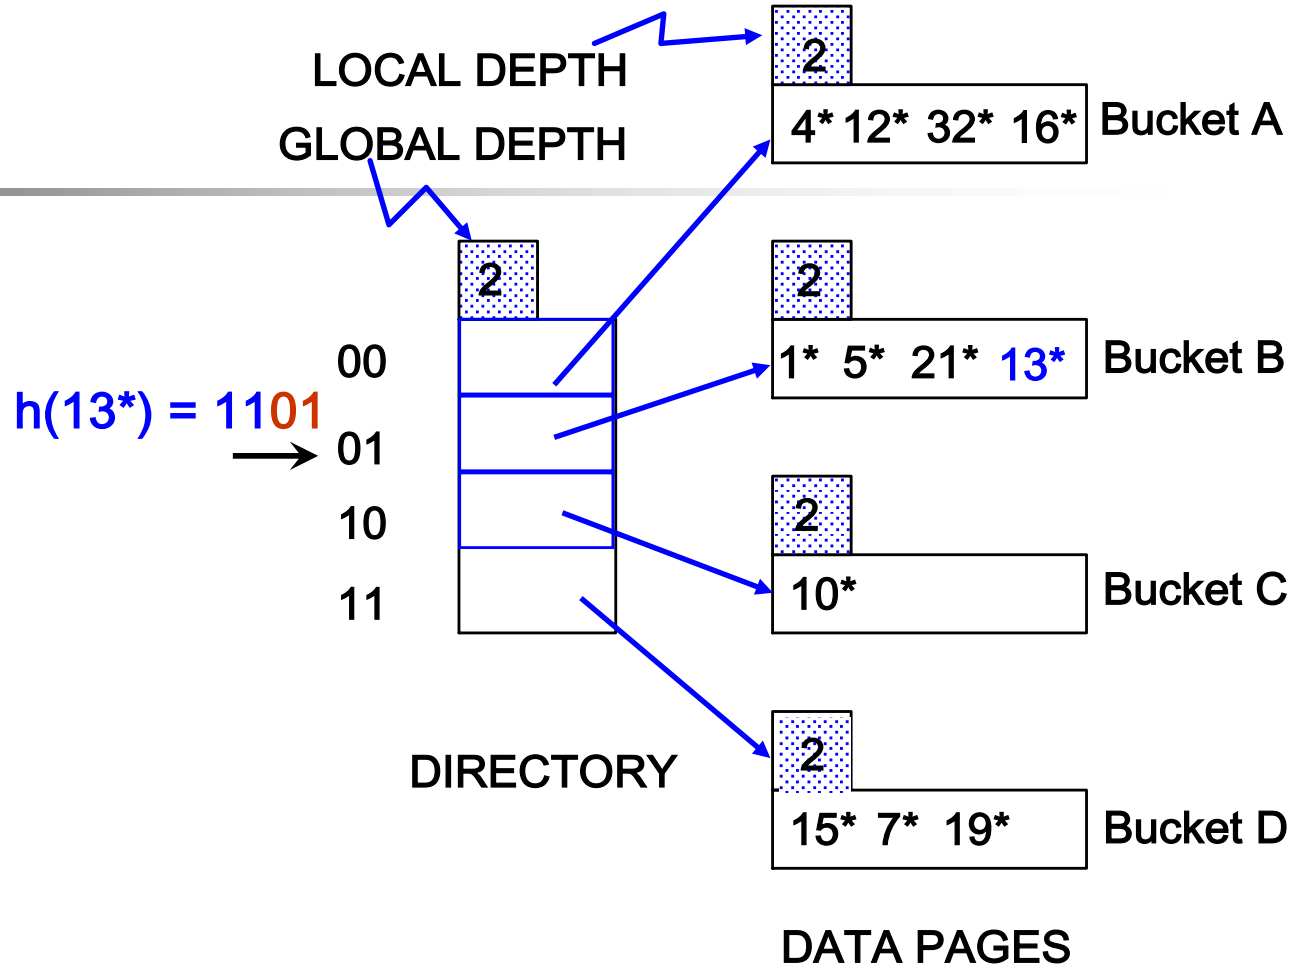

Insert 20

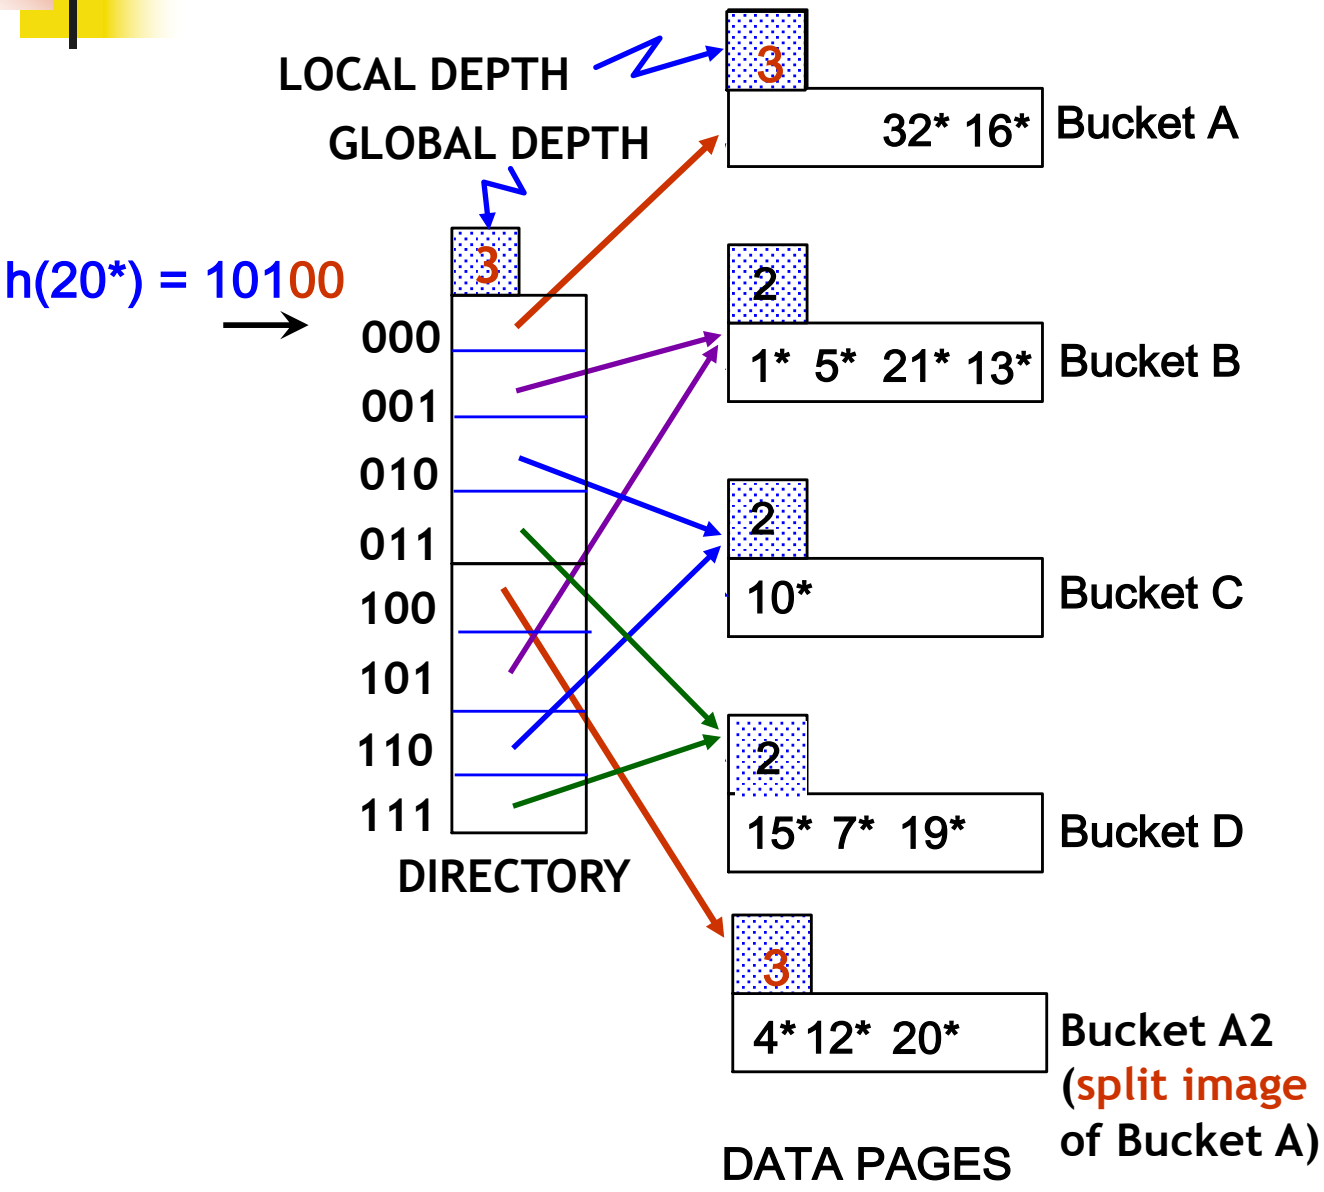

Notice that splitting a bucket only requires doubling the directory when LD = GD

Directory doubled by *copying it over* and fixing pointer to split image page

Directory Doubling

Why use least significant bits in directory?

Allows for doubling via copying!

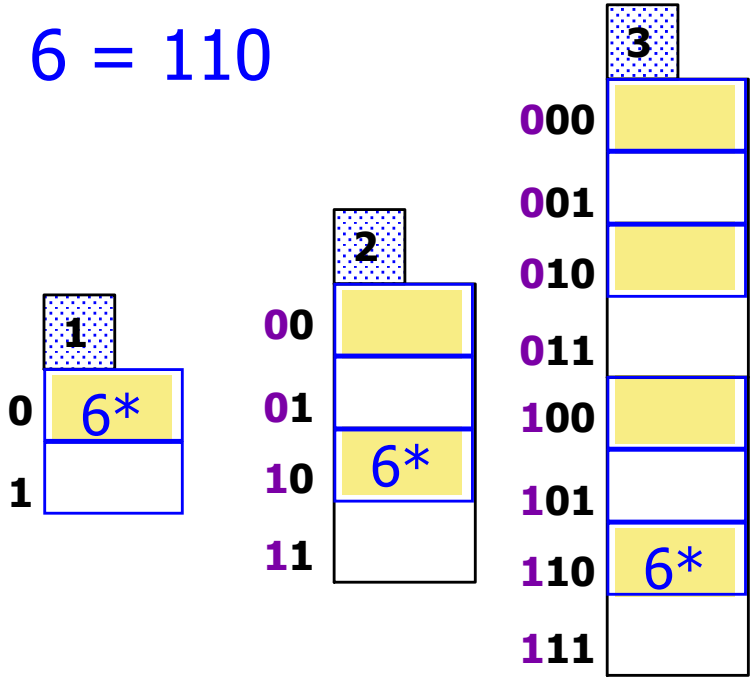

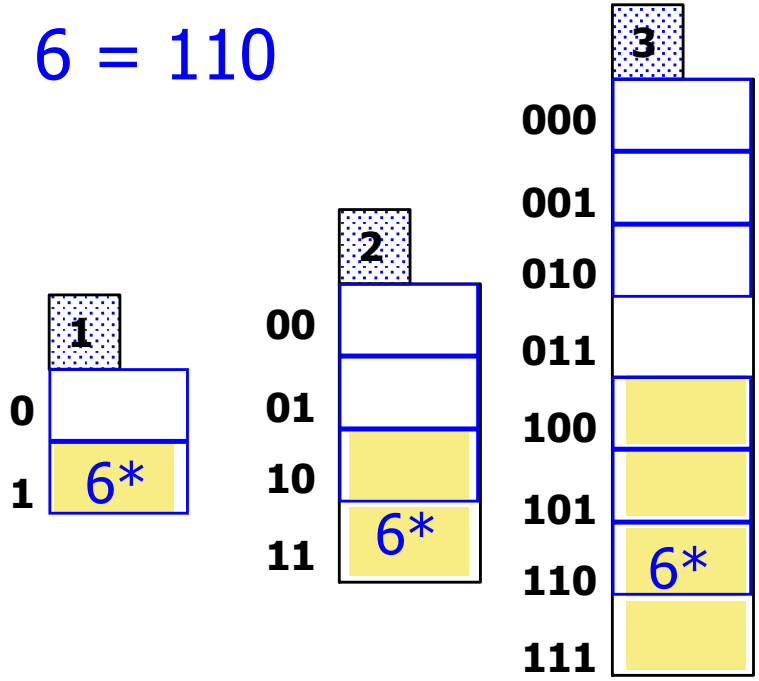

Least Significant vs. Most Significant

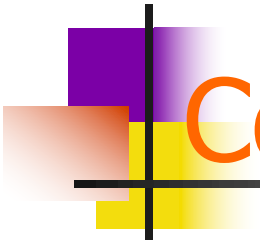

Comments on Extensible Hashing

- How many disk accesses for equality search?
  - One if directory fits in memory, else two.
- Directory grows in spurts, and, if the distribution *of hash values* is skewed, directory can grow large.
- Do we ever need overflow pages?
  - Multiple entries with same hash value cause problems!
- **Delete**: Reverse of inserts – see textbook.

Quiz: Insert 42, 9

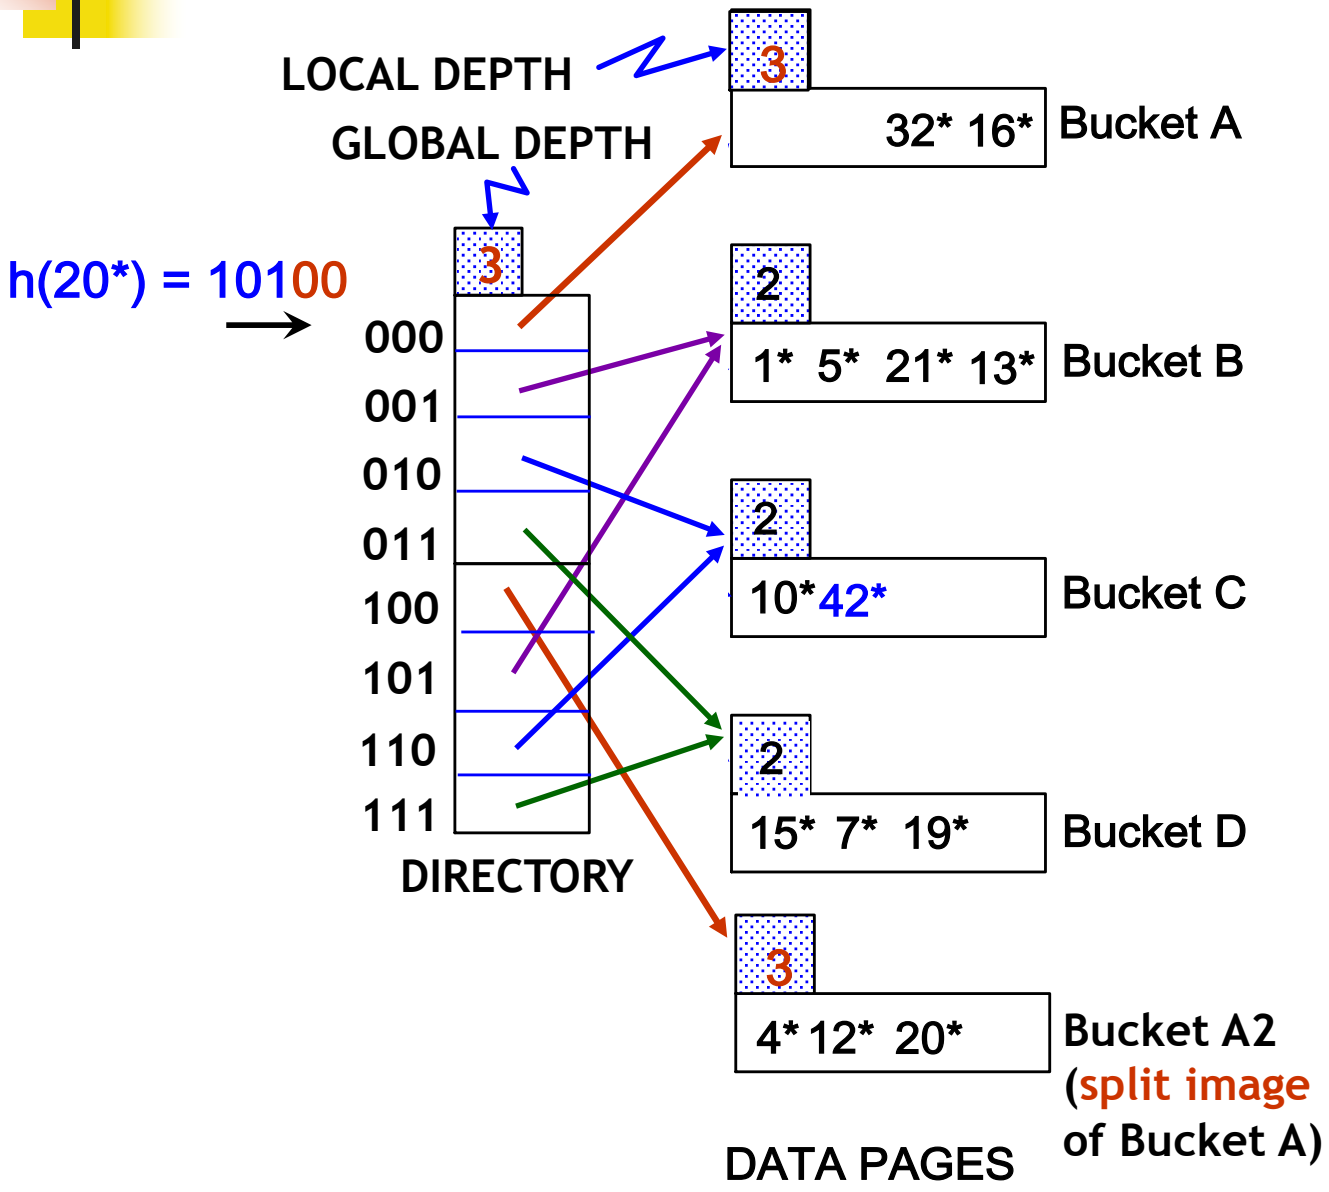

$42=(101010)_2$

$9=(1001)_2$

Answer: Insert 42, 9

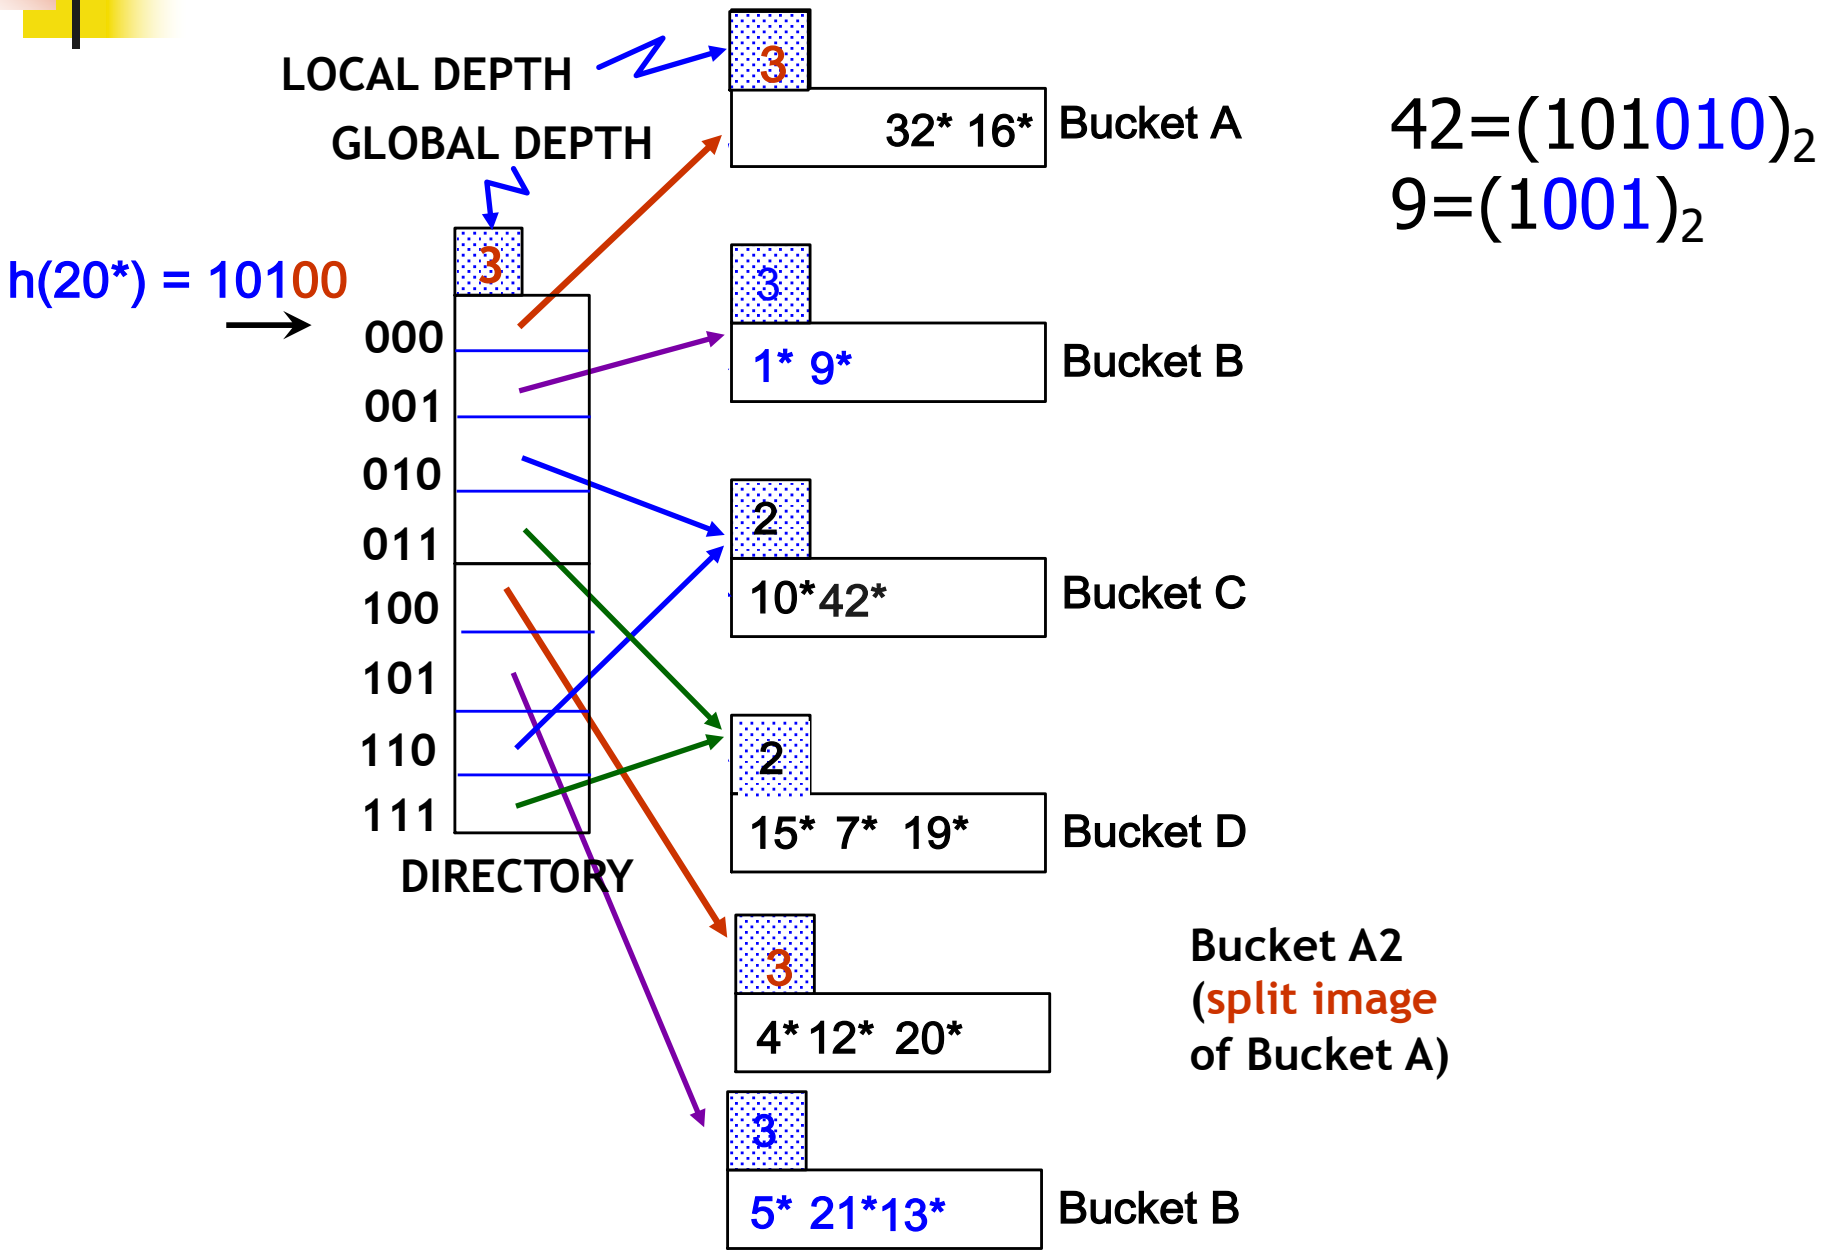

Linear Hashing

- Another dynamic hashing scheme
- Eliminates long overflow chains without using a directory.
- **Main idea:** Use a family of hash functions $\mathbf{h}_0$, $\mathbf{h}_1$, $\mathbf{h}_2$, ...
  - $\mathbf{h}_{i+1}$ doubles the range of $\mathbf{h}_i$ (similar to directory doubling)
  - Hash family typically obtained by choosing hash function *h()* and initial number of buckets *N*
  - Define $h_i(value) = h(value) mod (2^i N)$
  - If N is a power of 2, apply hash function *h(),* and look at last d_i bits
    - d_0 number of bits needed to represent *N*
    - $d_i = d_0 + i$

Linear Hashing

- Splitting proceeds in rounds
  - During round *level*, only $\mathbf{h}_{level}$ and $\mathbf{h}_{level+1}$ are in use
- Variables
  - Level: Initialized to 0
  - Next: Pointer to the bucket being split
- At the beginning of round # Level, the # buckets in the file = N * 2^{Level}
  - N is initial number of buckets

Linear Hashing Example

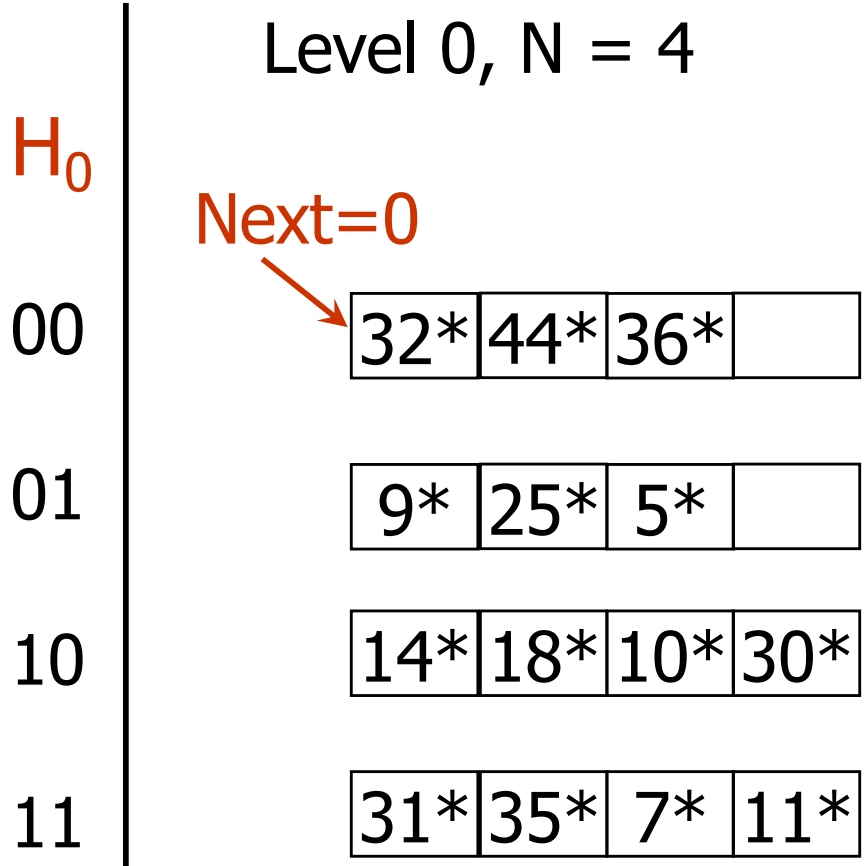

What happens if we add 37*?
$37=(100101)_2$

Linear Hashing Example

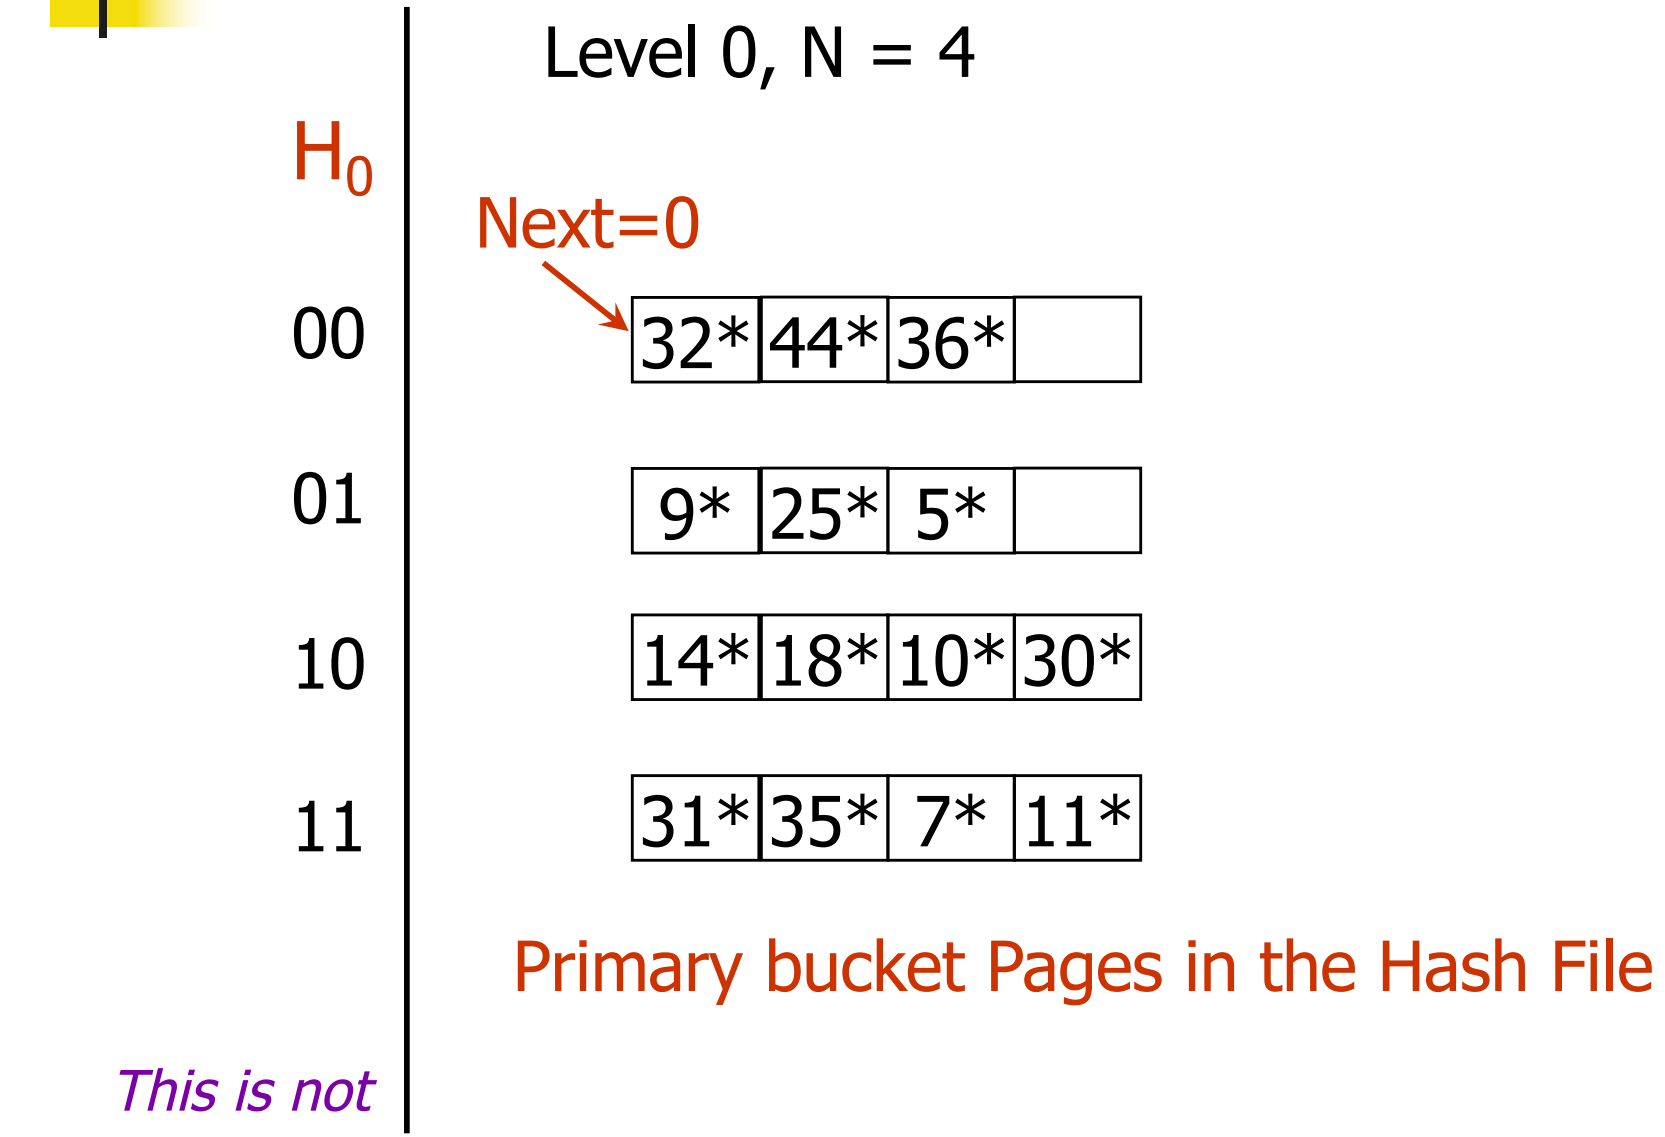

Linear Hashing Example

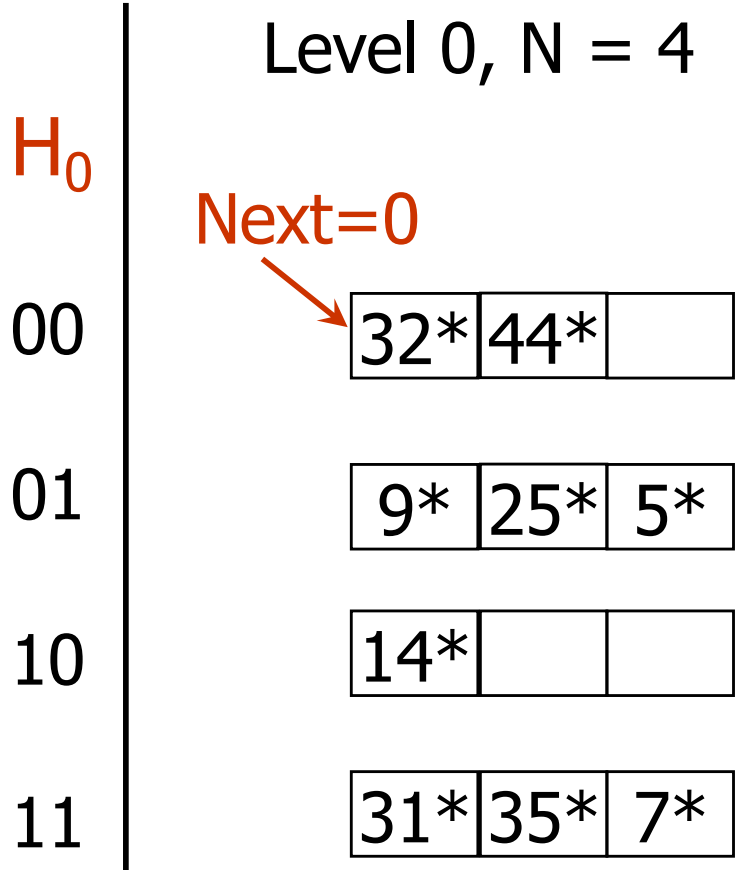

What happens if we add 37*?
$37=(100101)_2$

This is not actually stored

Primary bucket Pages in the Hash File.
Next points to the next bucket to be split on an overflow

Linear Hashing Example

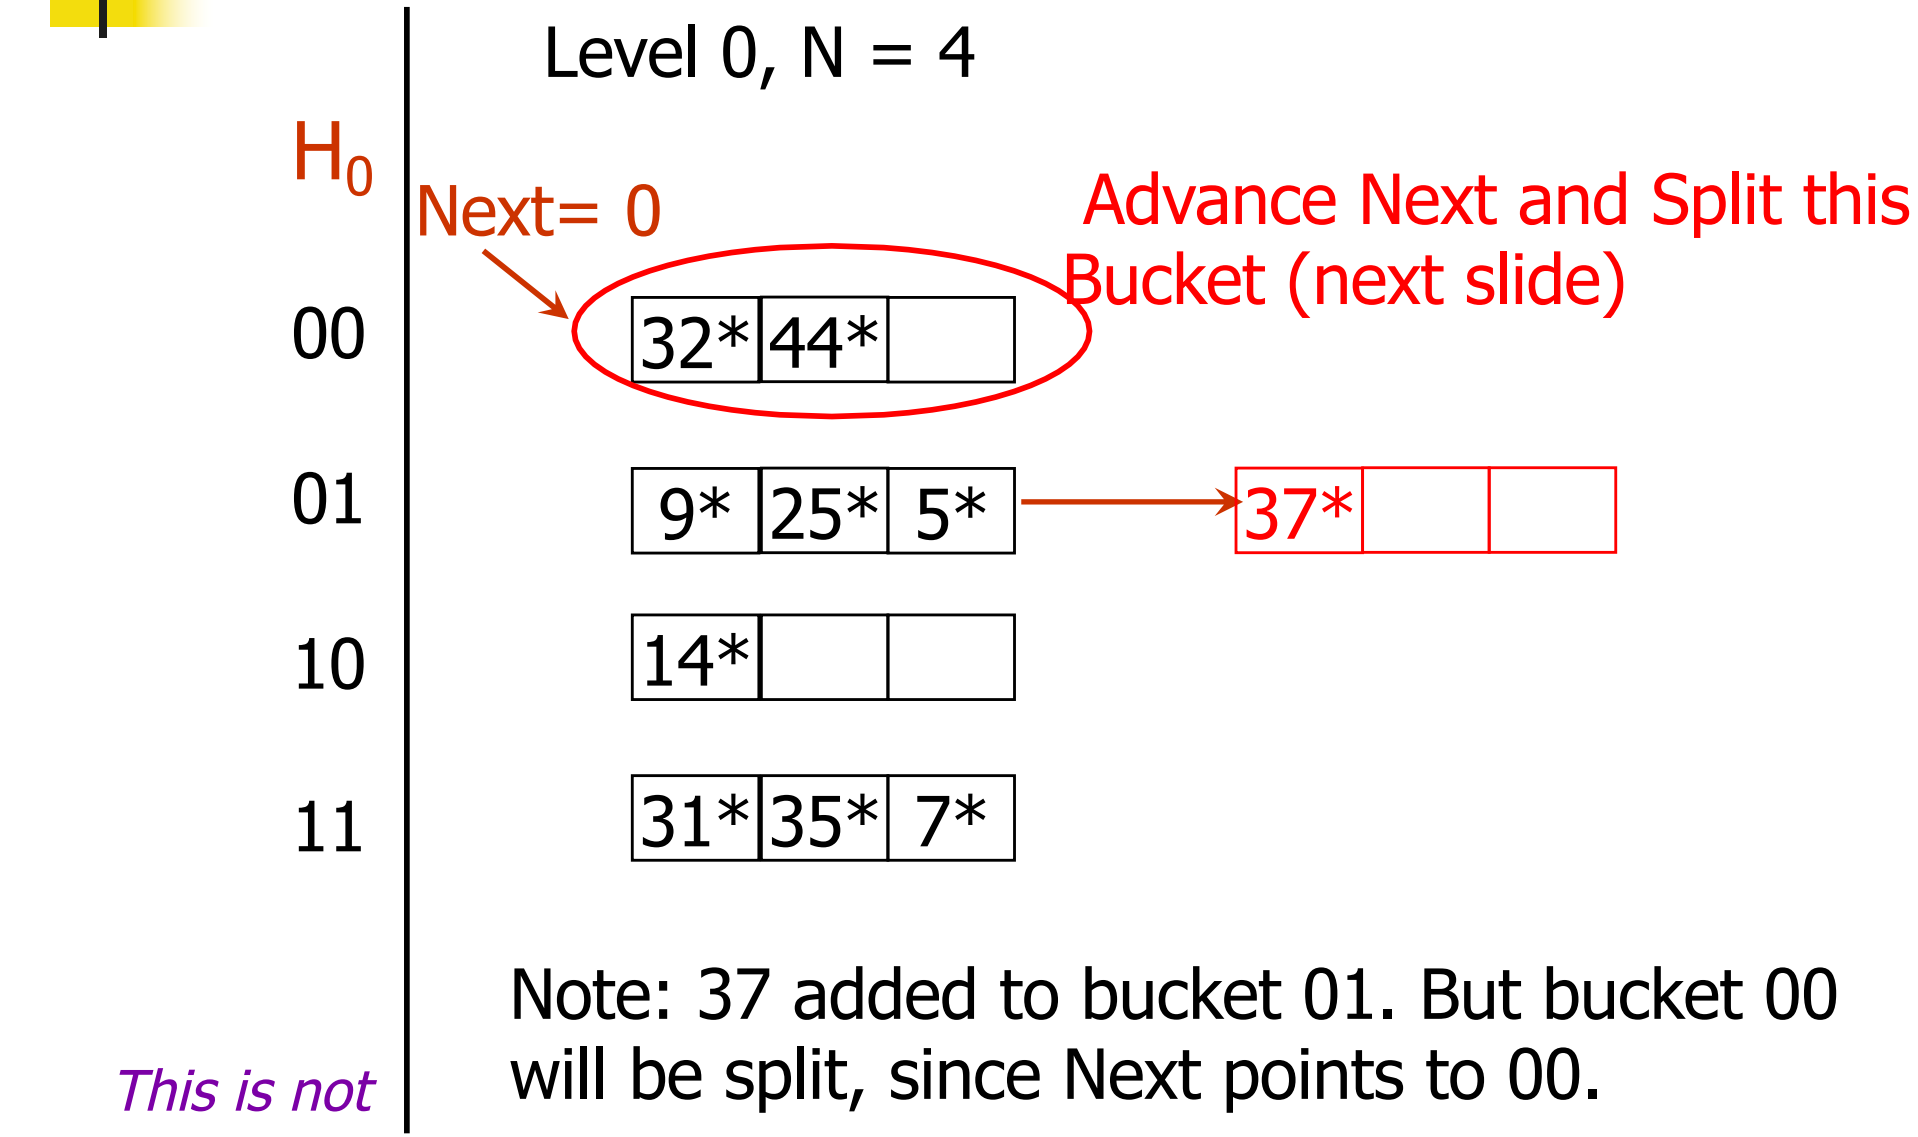

Linear Hashing Example

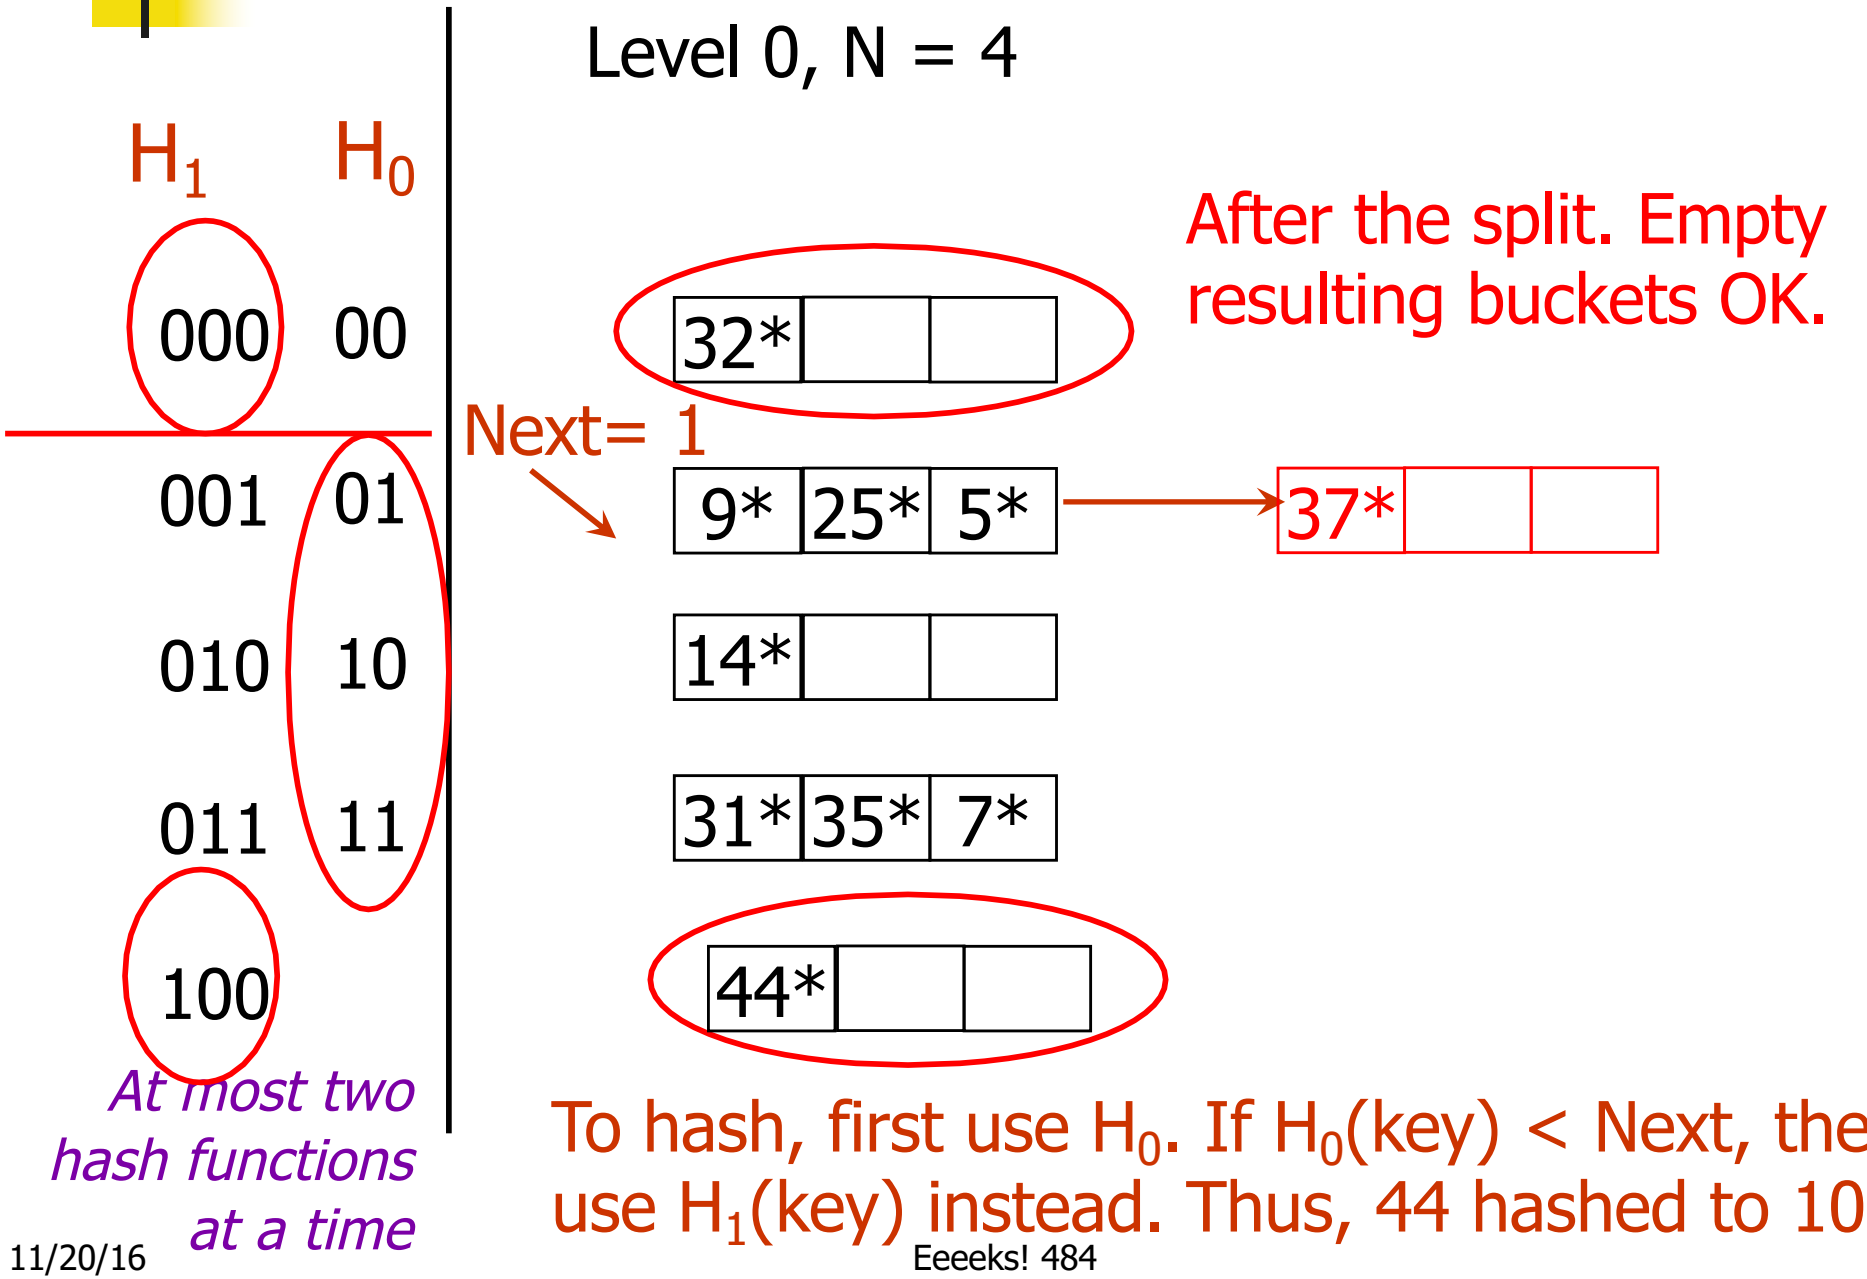

Linear Hashing Example

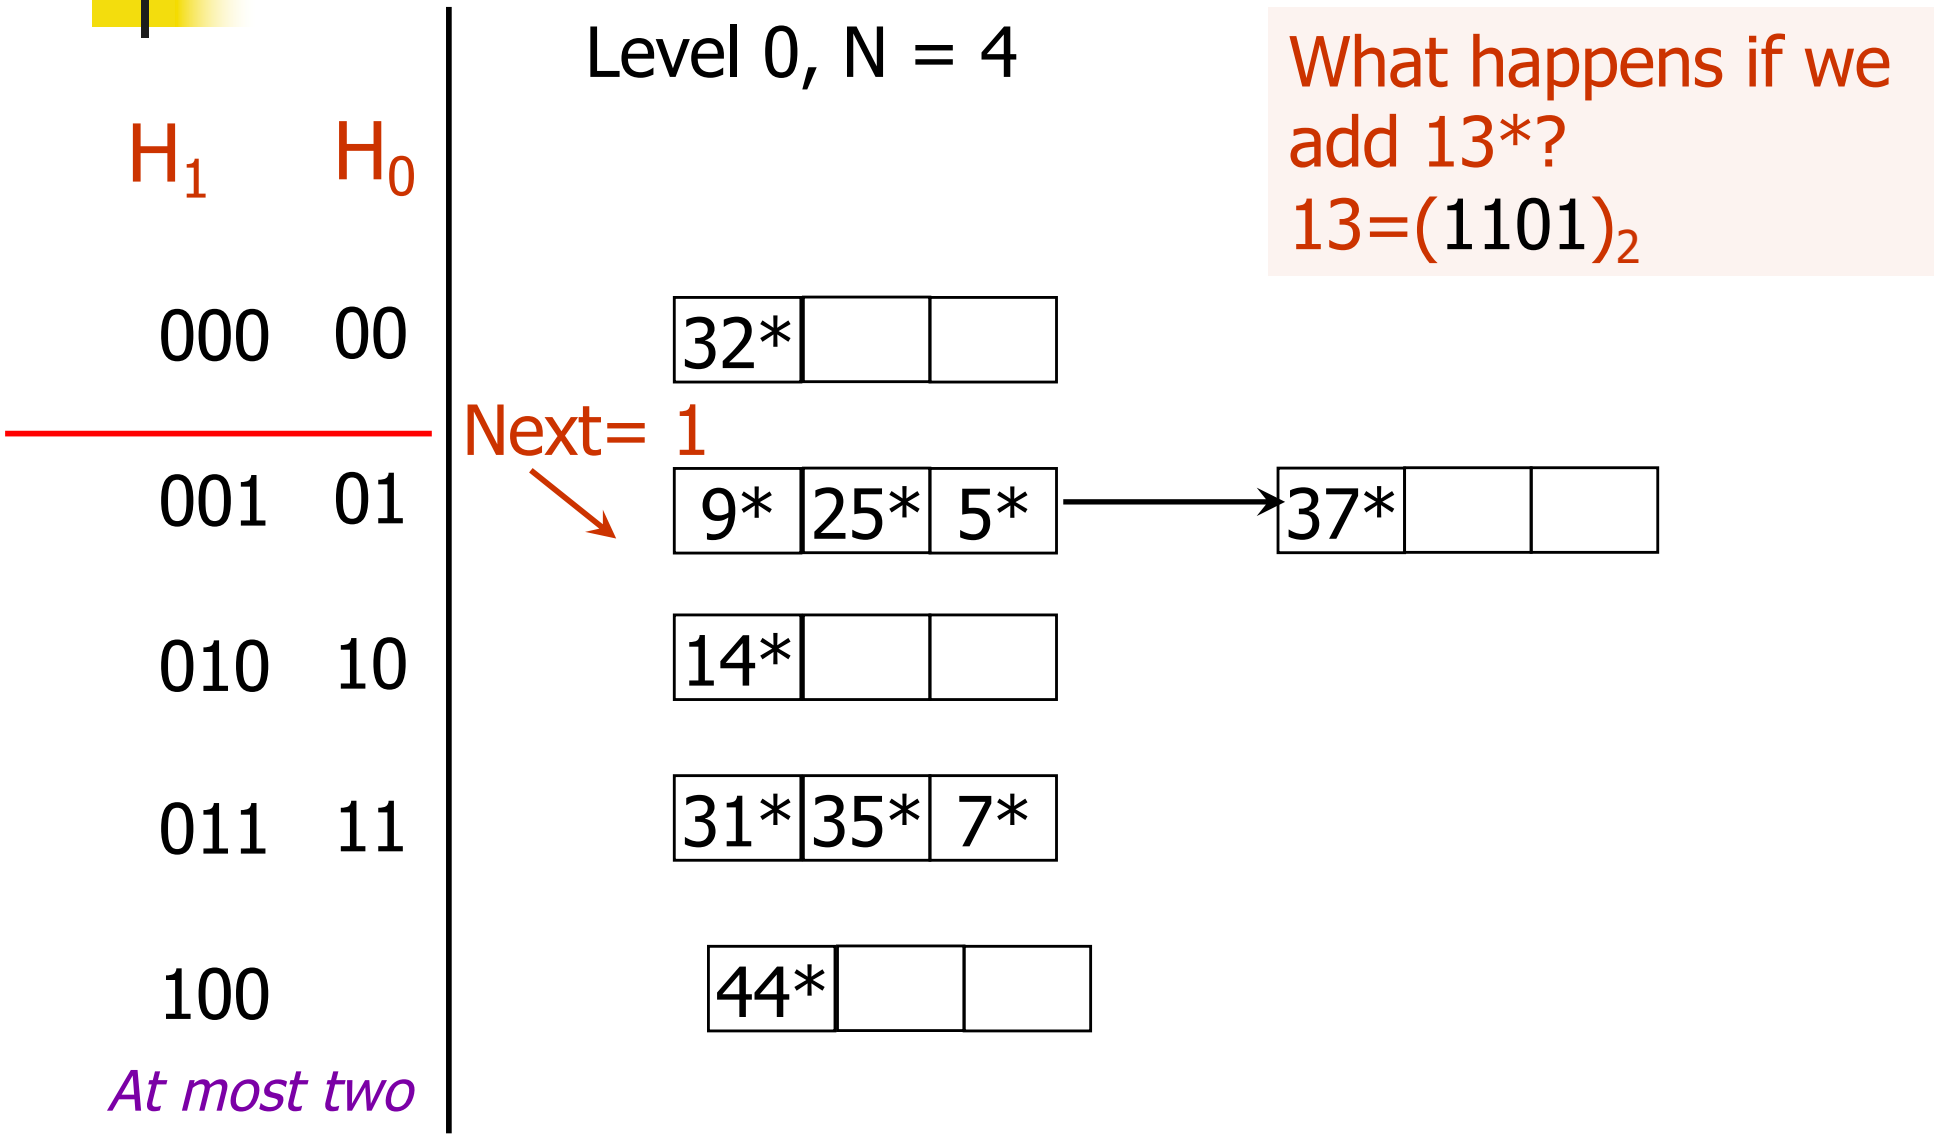

Linear Hashing Example

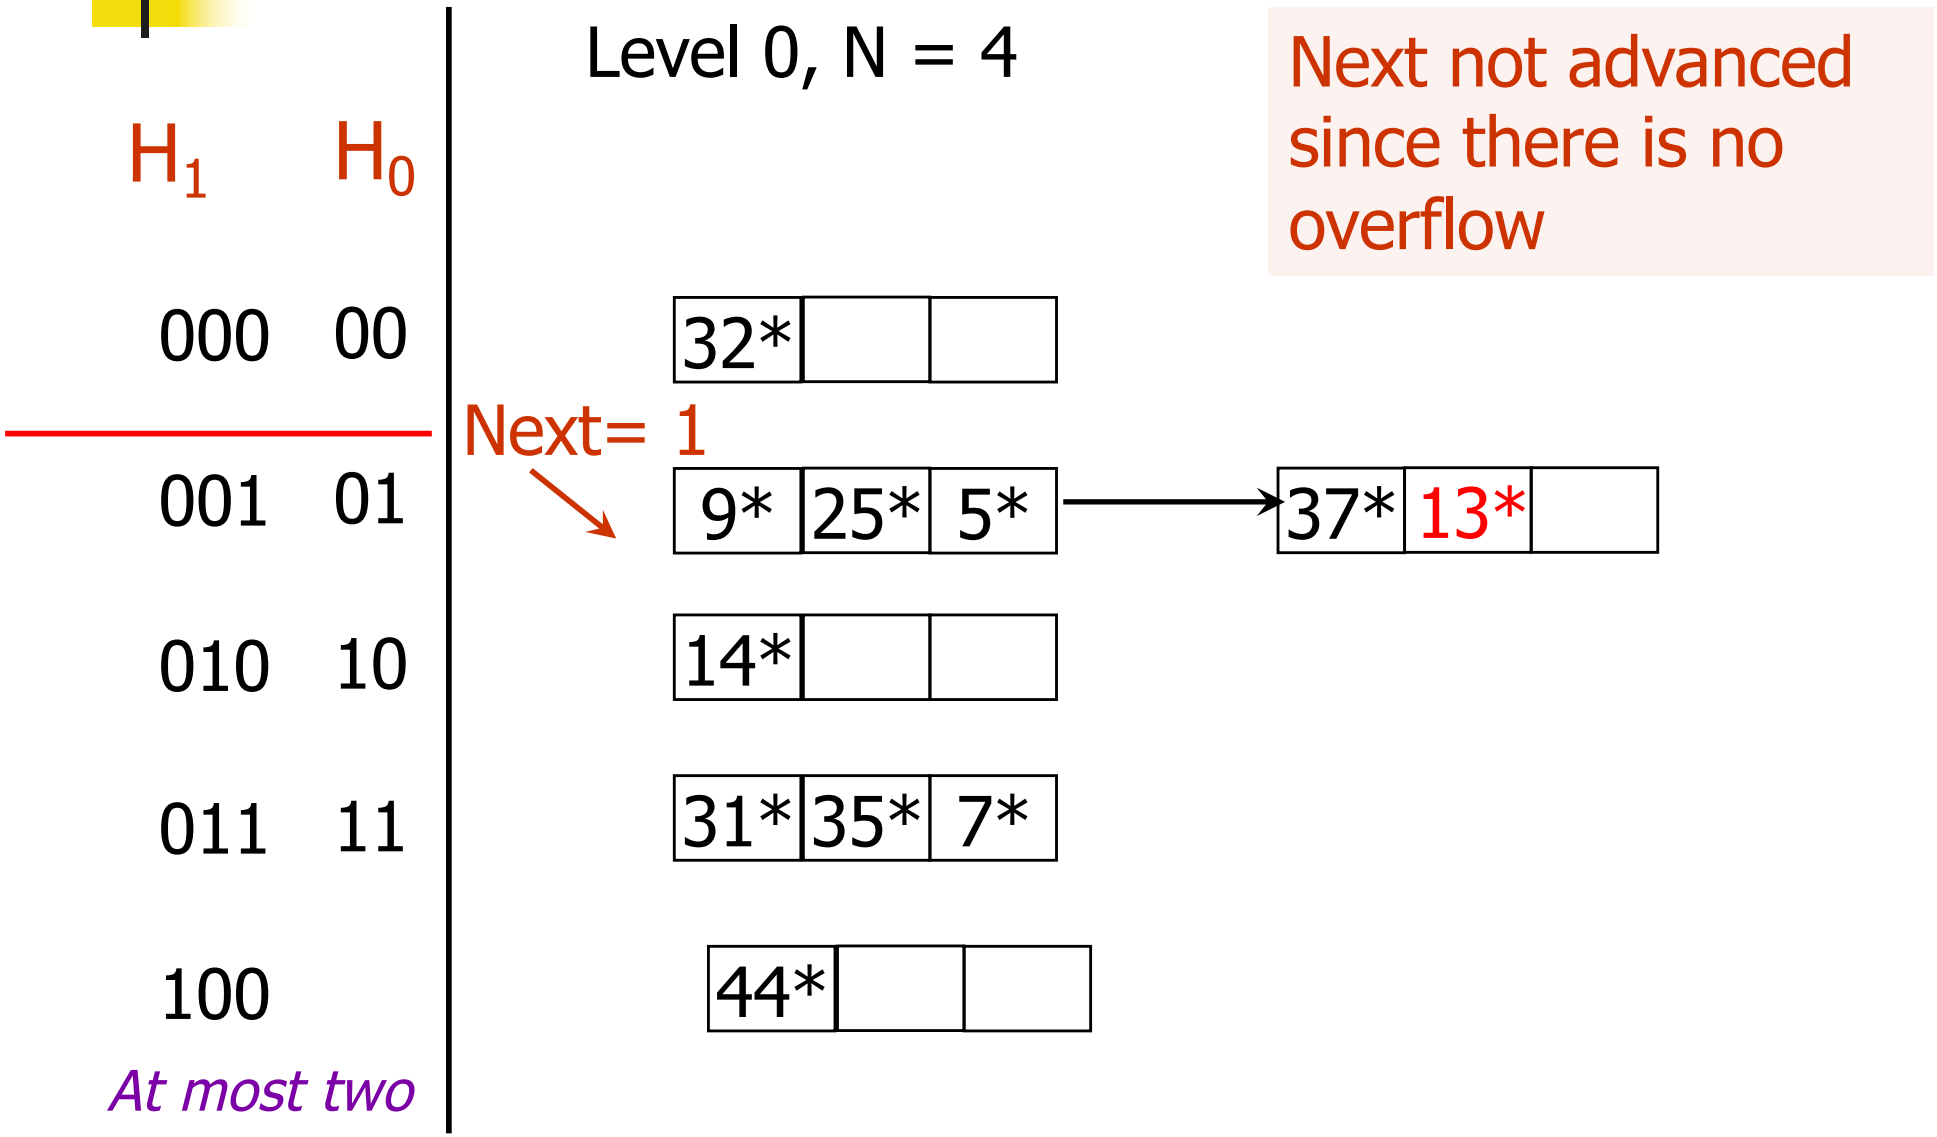

Linear Hashing Example

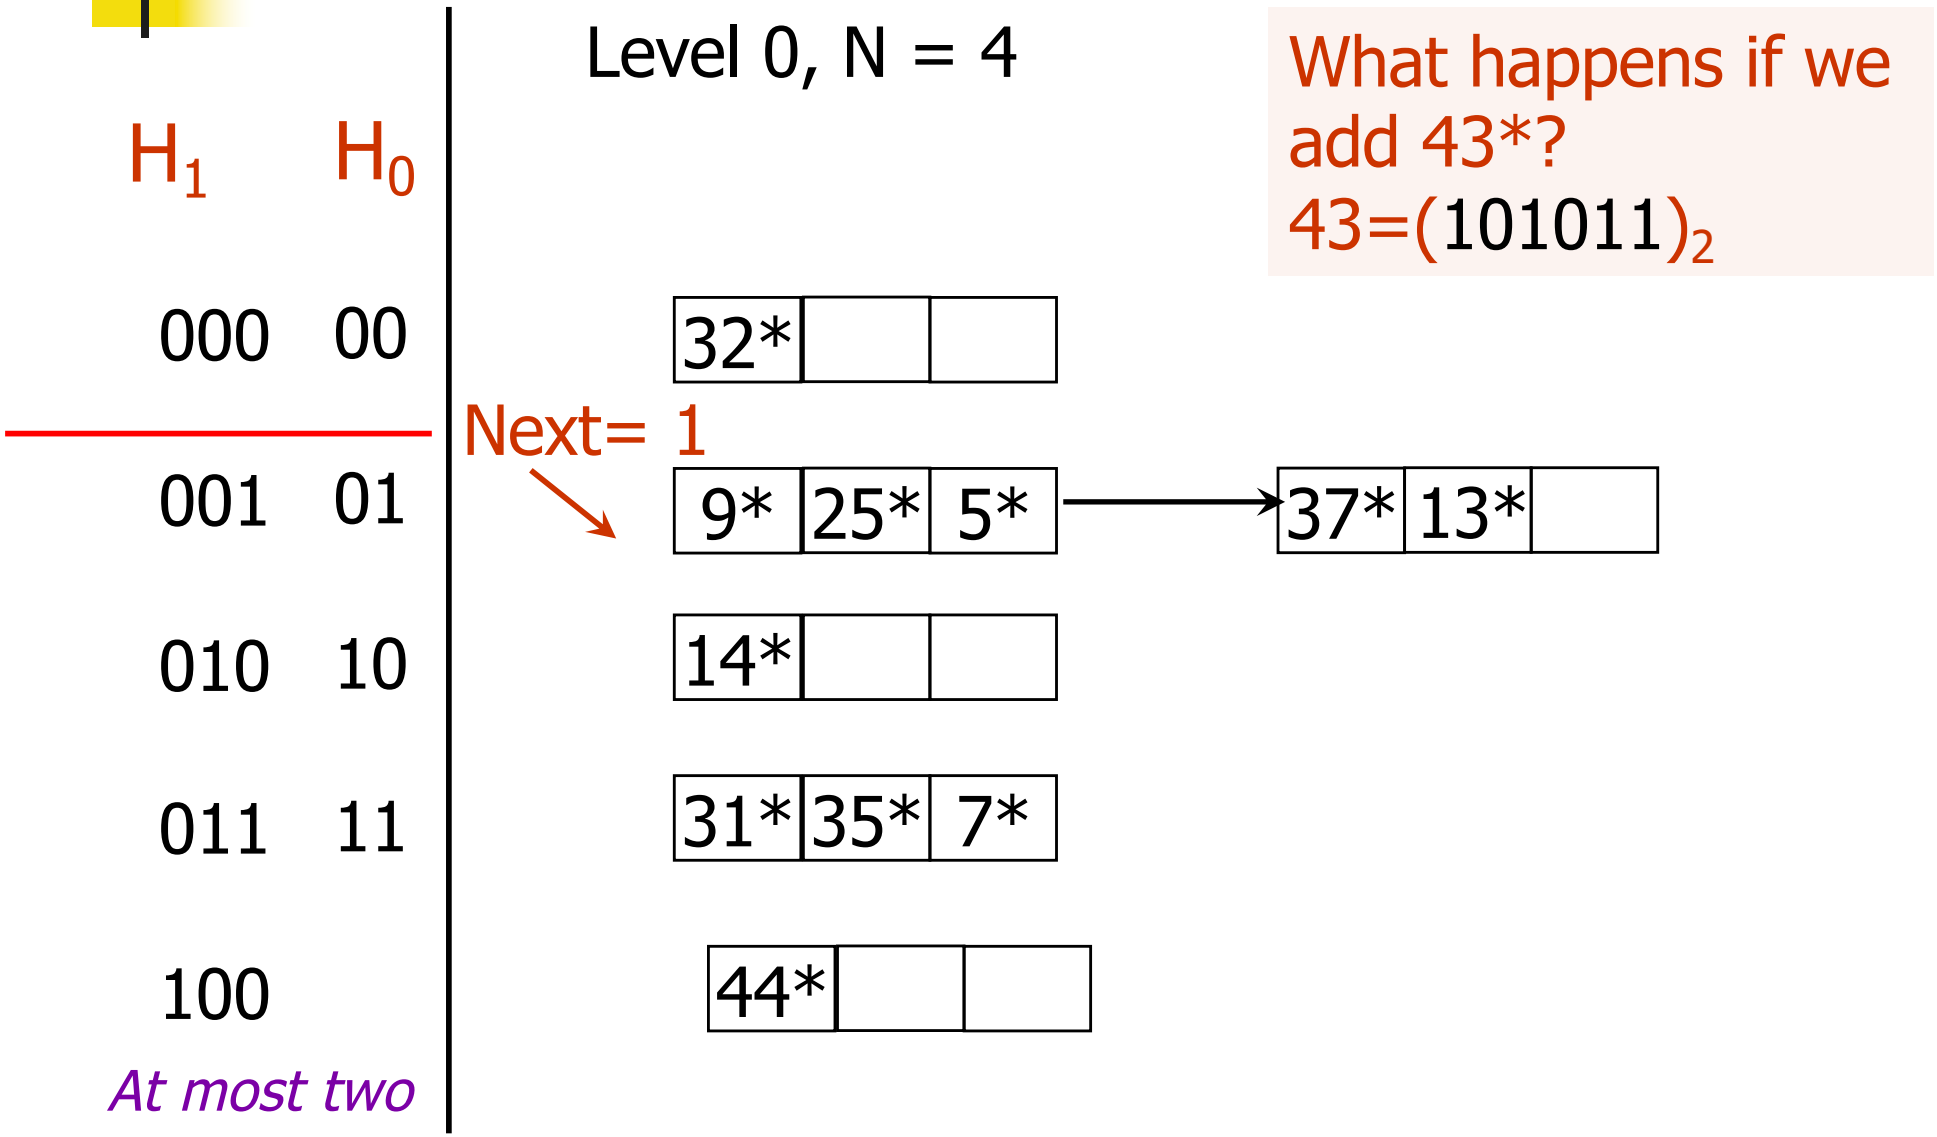

Linear Hashing Example

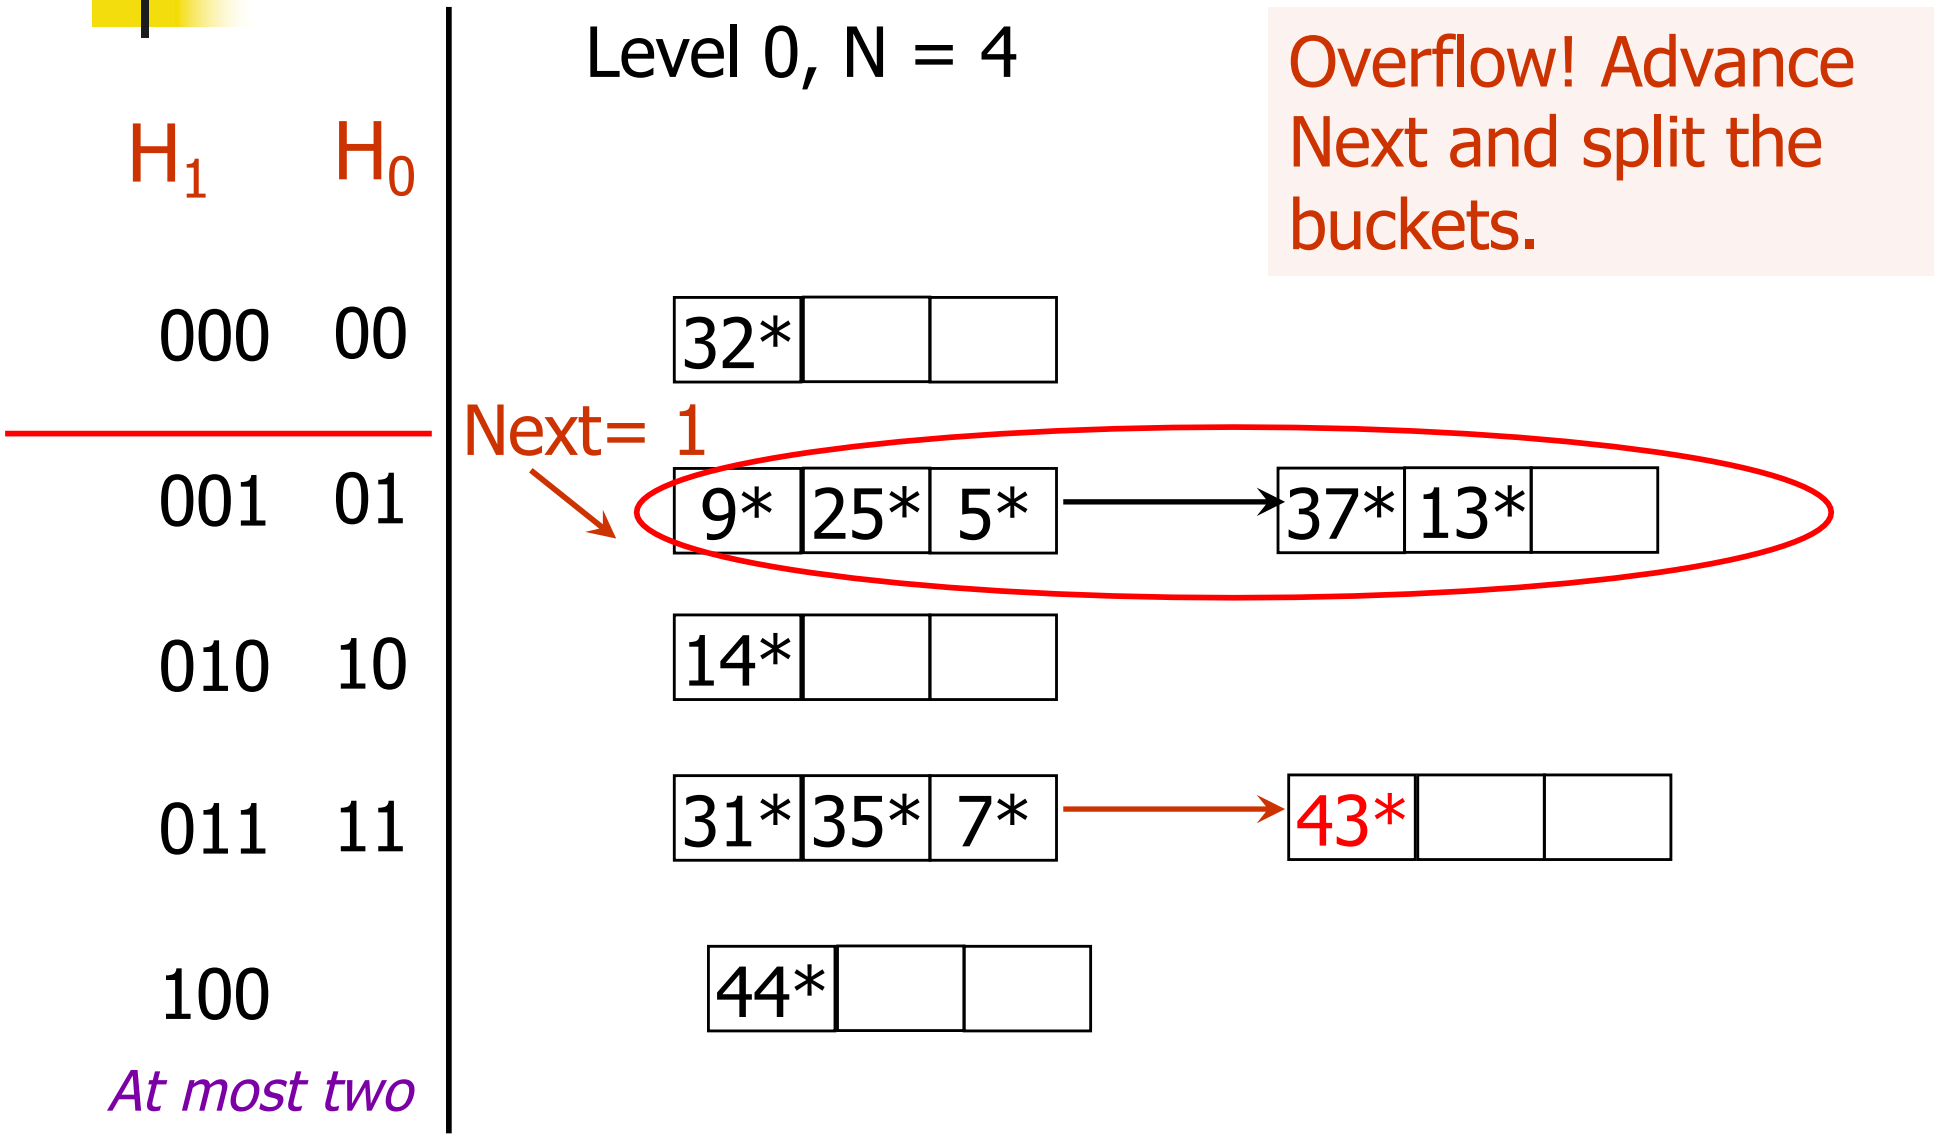

Linear Hashing Example

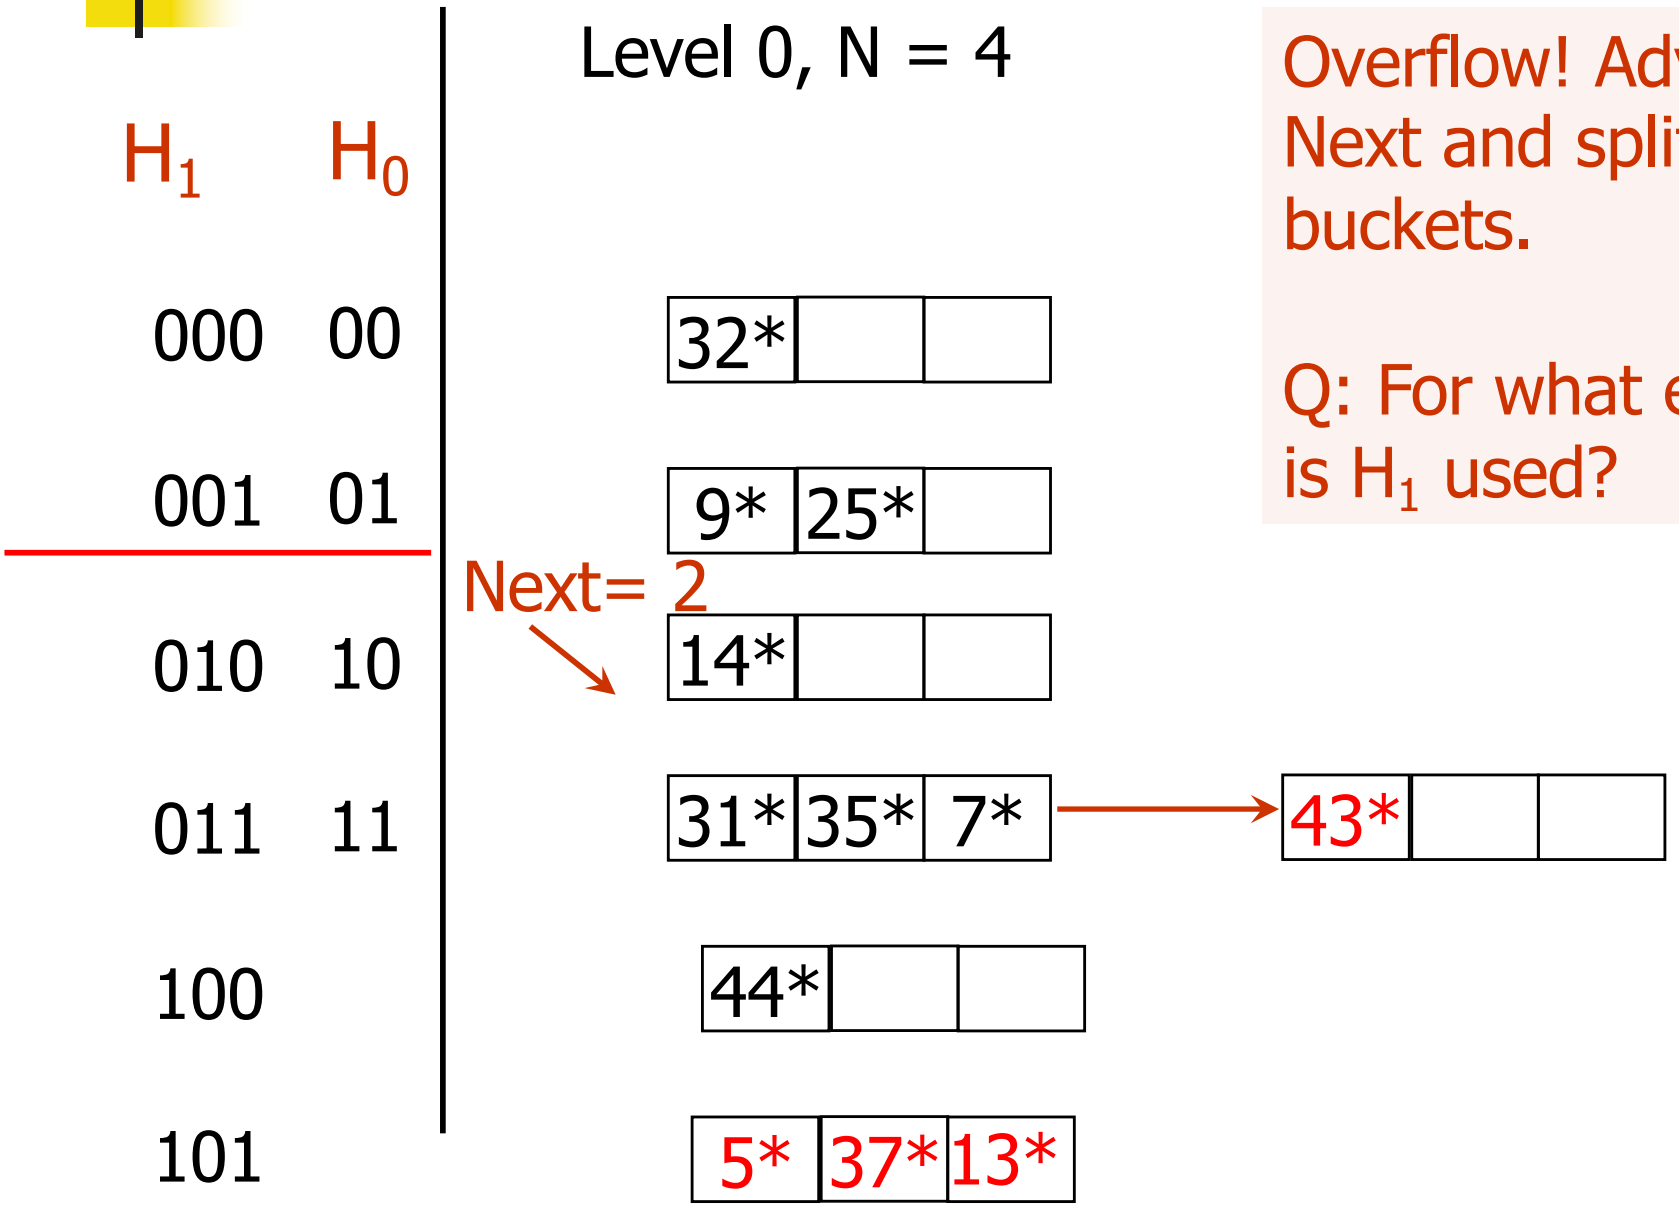

Overflow! Advance Next and split the buckets.

Q: For what entries is H_1 used?

Linear Hashing Example

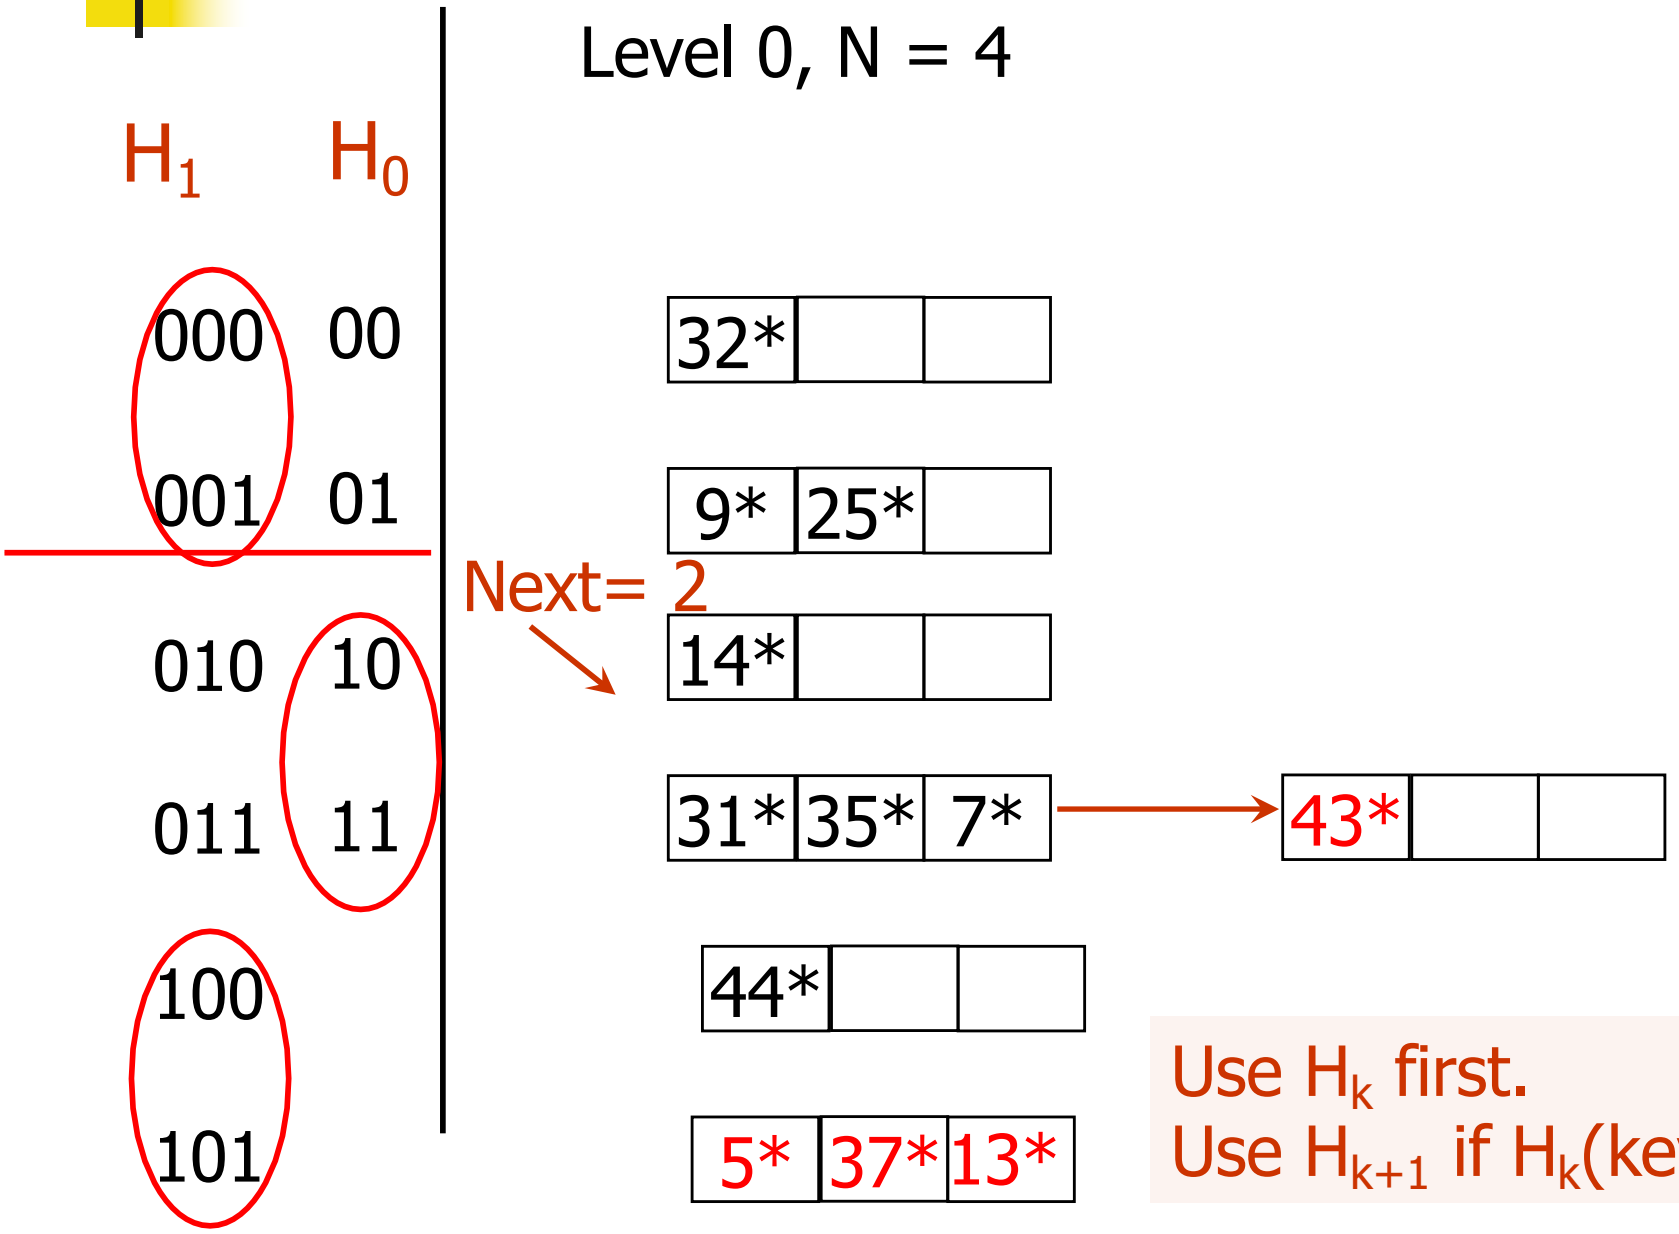

Use H_k first.
Use H_{k+1} if H_k(key) < Next.

Linear Hashing Example

Level 0, N = 4

H_1	H_0
000	00
001	01
010	10
011	11
100	
101	
110	

Next= 3

What happens on the next split?
(Imagine the buckets)

Linear Hashing Example

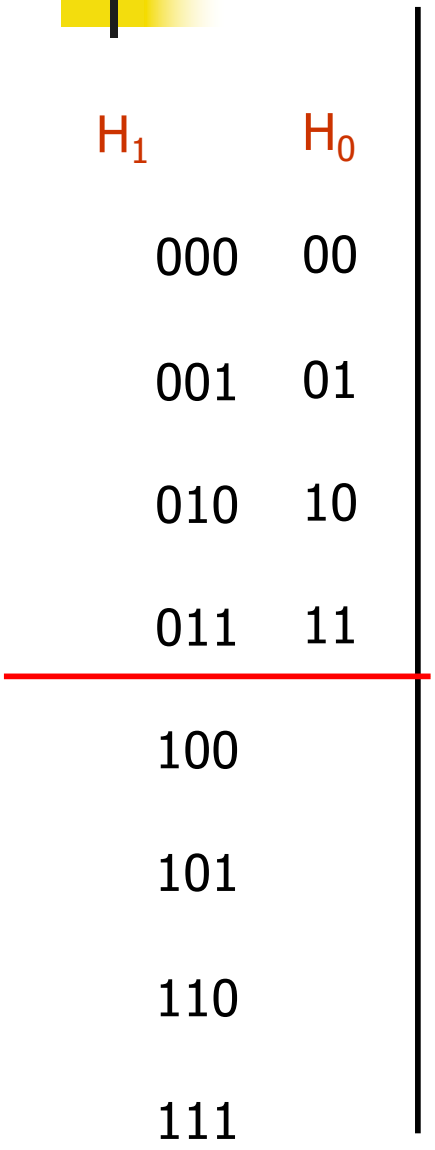

Level 0, N = 4

Next= 4?

No. At this point, H0 is useless – will never be used. Reset Next to and advance Level.

Linear Hashing Example

Level 1, N = 4

H_1

Next= 0

000

001

010

011

100

101

110

111

Level advanced. H_0 no longer needed.
H_2 will be used when Next advances further.

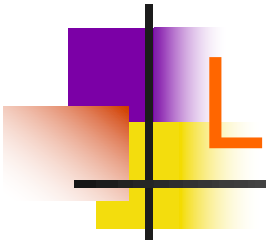

Linear Hashing (Generalization)

- *Can choose any criterion to trigger split.*
  - Split on a overflow (as we assumed)
  - Space utilization on the page > 90%
- Since buckets are split round-robin, long overflow chains don't develop!
- **Deletes**: see textbook

Quiz: Add 13* and 29*

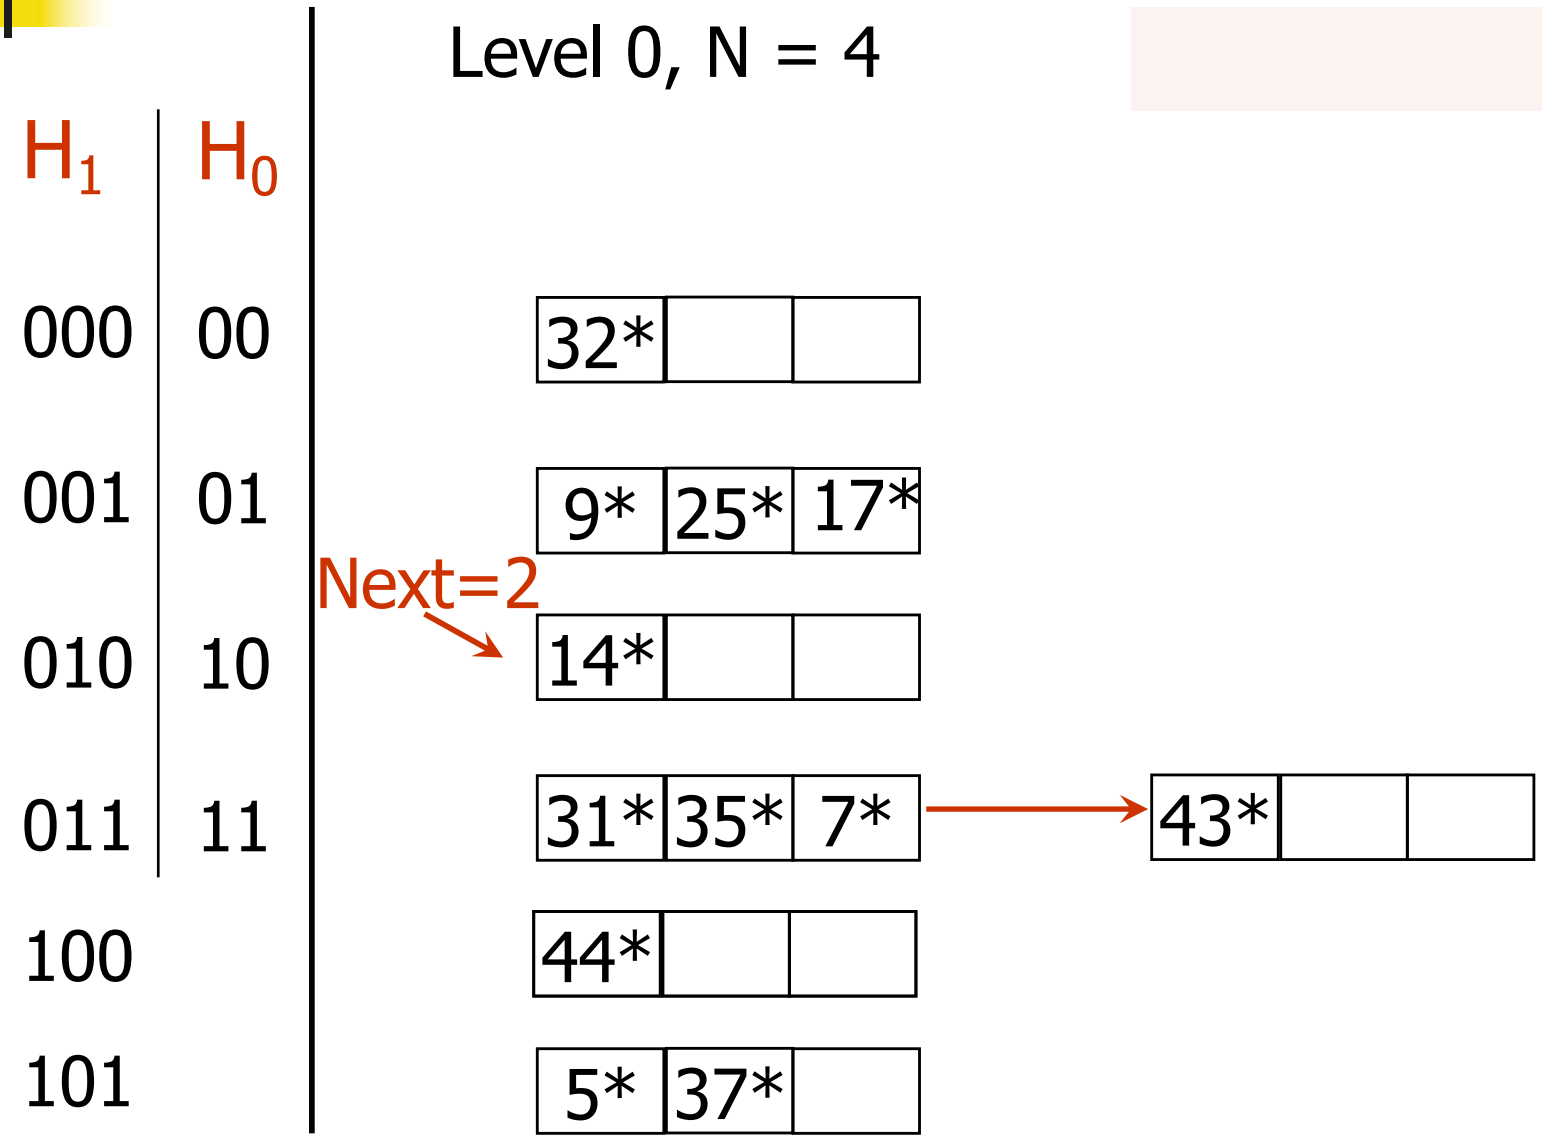

Quiz: After adding 13

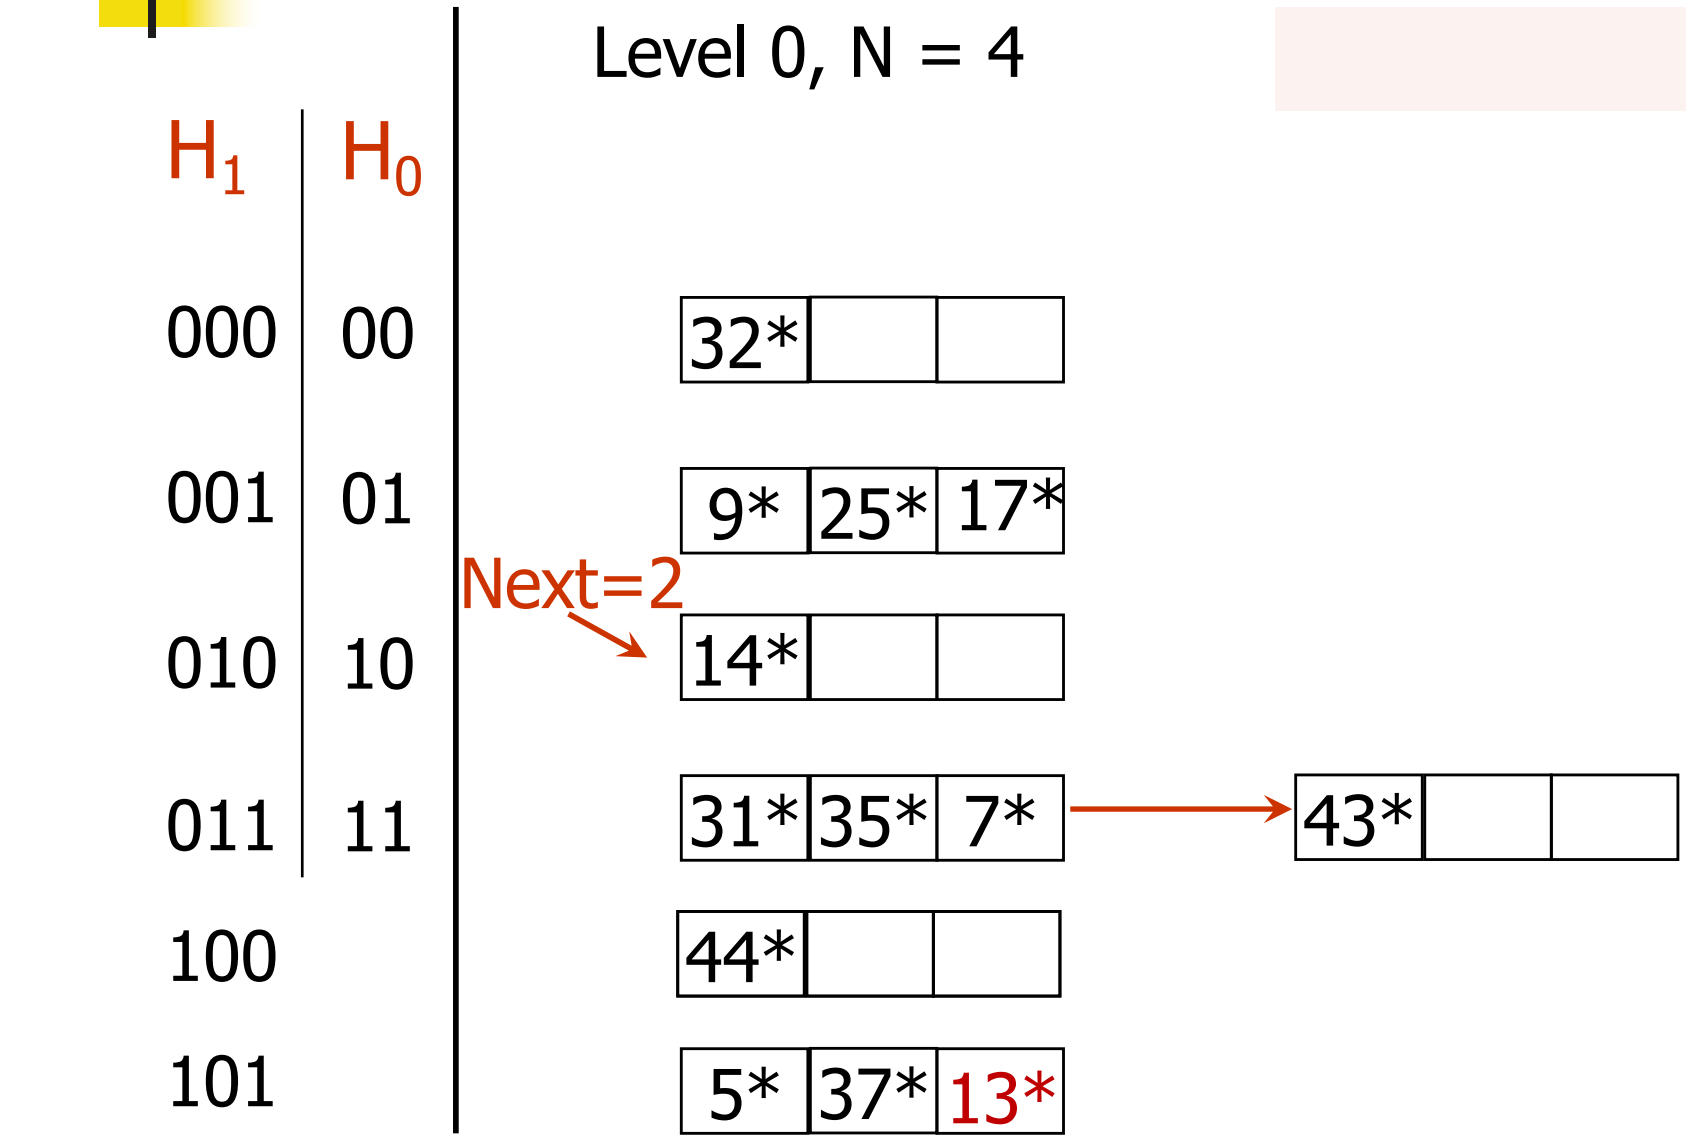

Quiz: After adding 29* (Step 1)

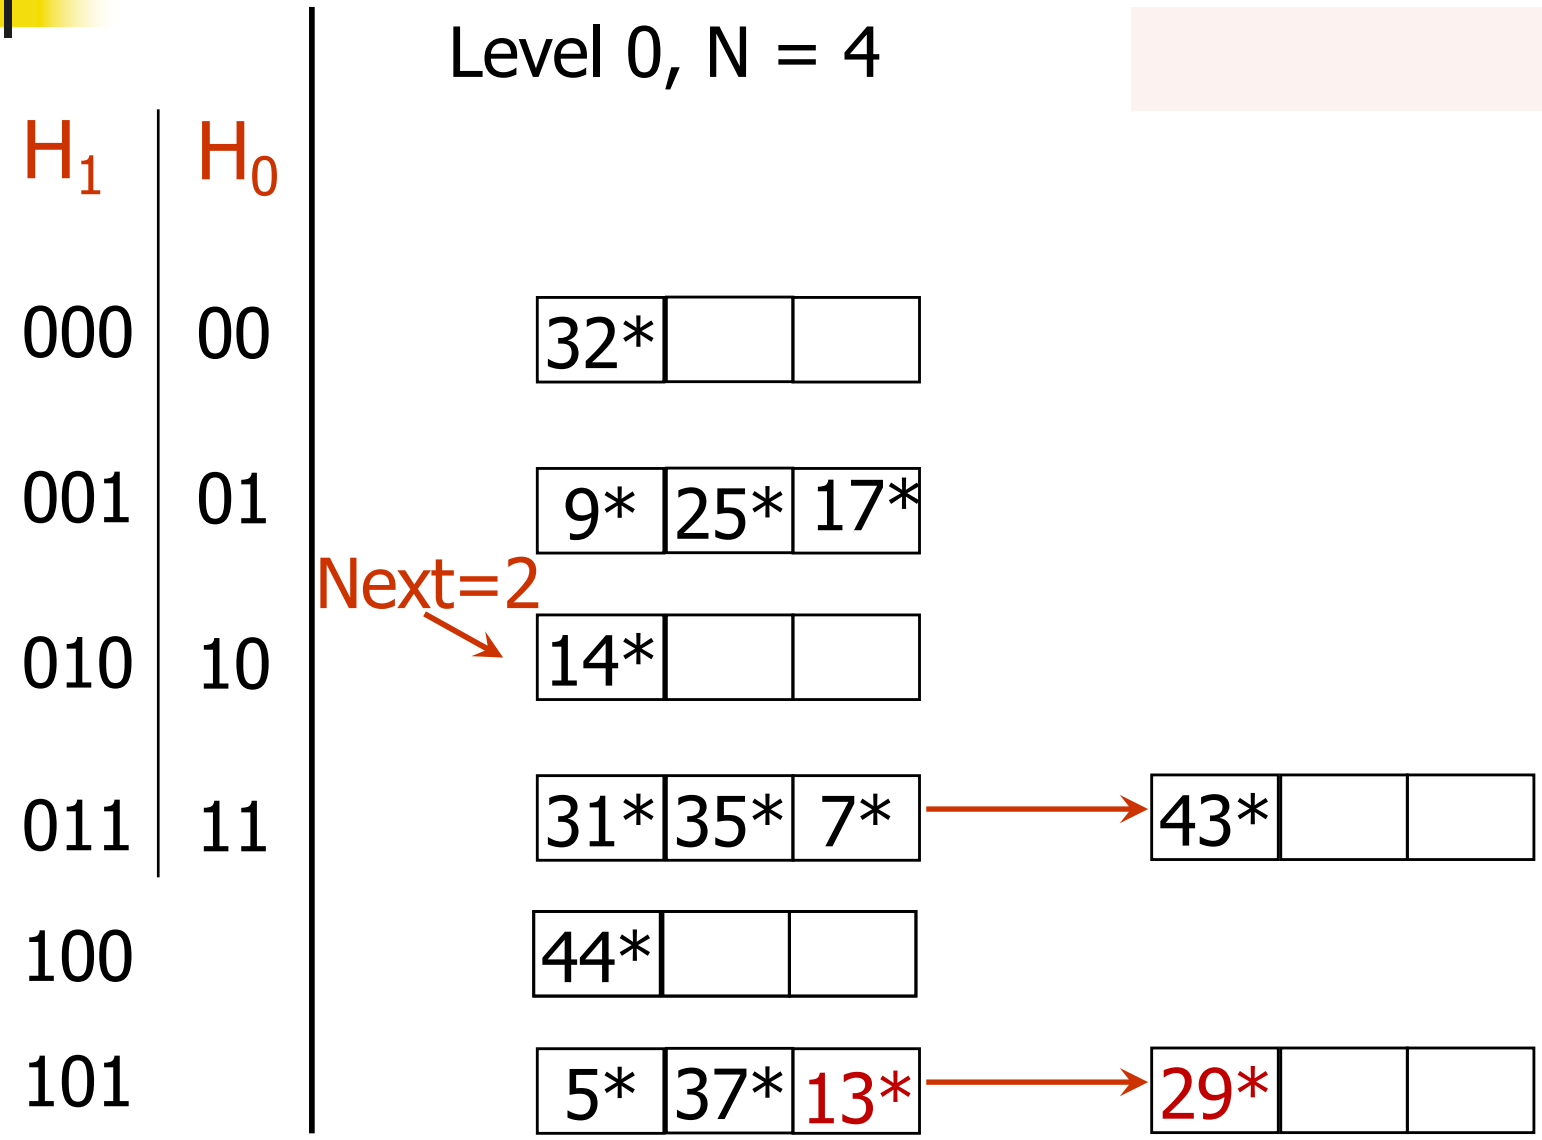

Quiz: After adding 29* (done)

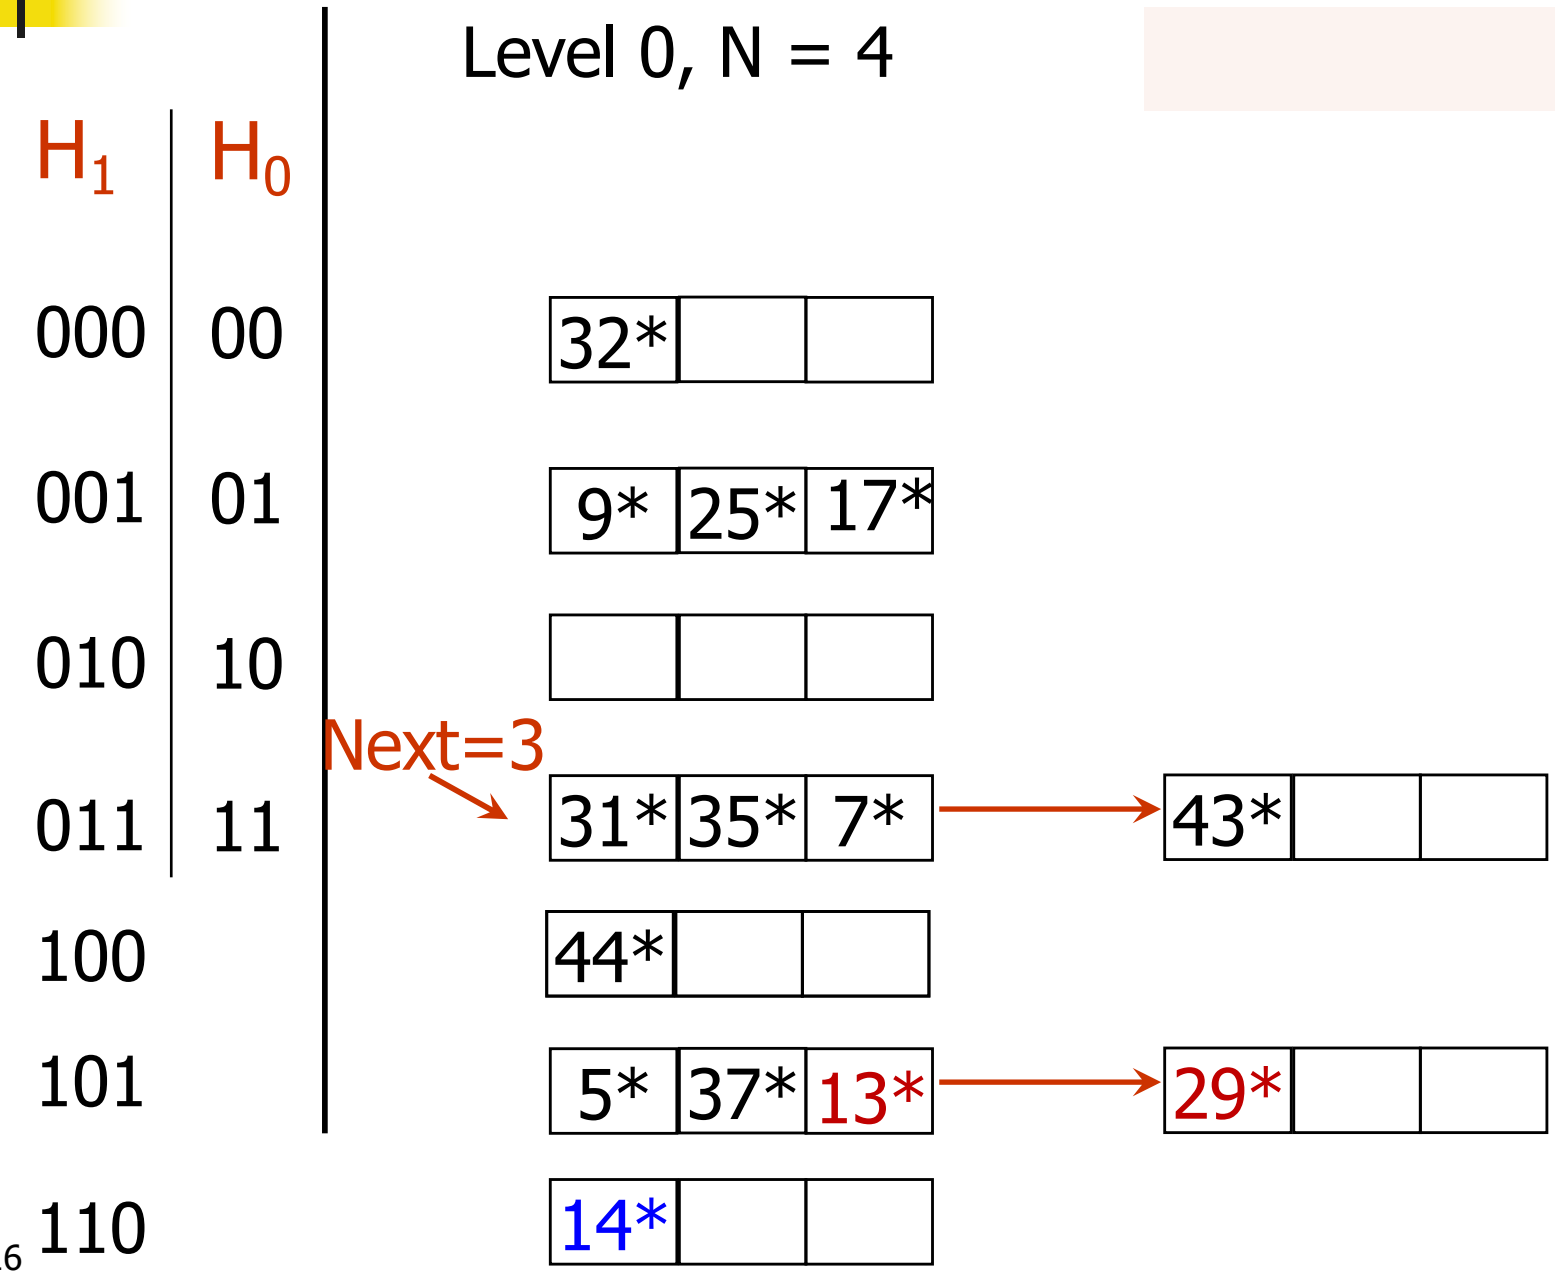

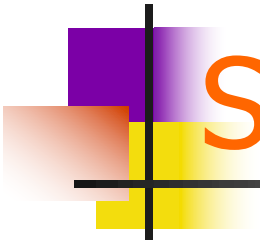

Summary

- Discussed 3 kinds of hash-based indexes
- Static Hashing can lead to long overflow chains.
- Extensible Hashing
  - Directory to keep track of buckets, doubles periodically.
  - Always splits the “right” bucket.
- Linear Hashing
  - Split buckets round-robin, and use overflow pages.
  - Space utilization could be lower than EH.

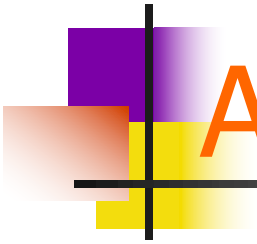

Announcements

- Review problems for hash indexes
  - 11.1, 11.3, 11.7, 11.9