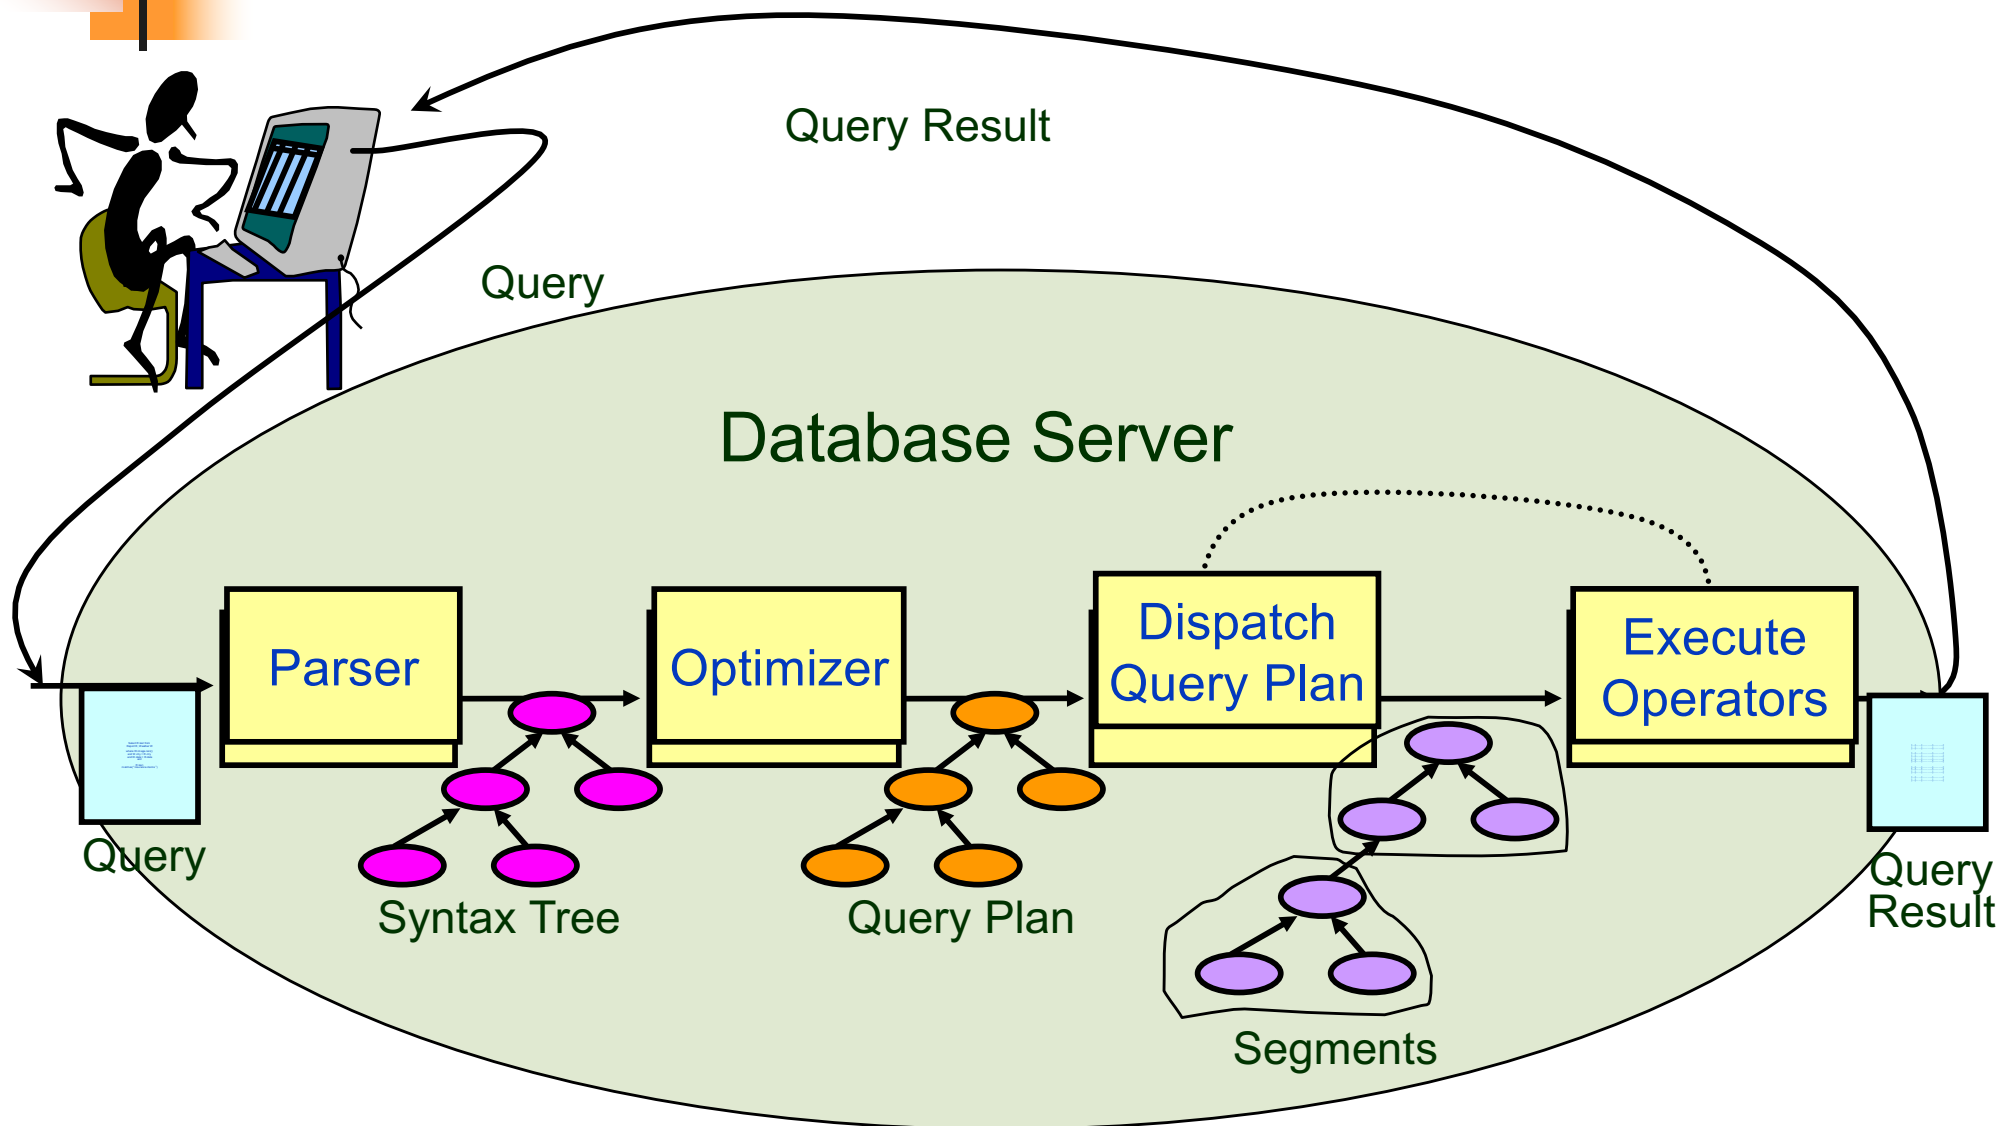




Evaluation of Relational Operations

Chapter 12 and 14

Query Execution Life-Cycle





Operator Evaluation


- How to implement common operators?
 - Selection
 - Projection (optional DISTINCT)
 - Join
 - Set Difference
 - Union
 - Aggregate operators (SUM, MIN, MAX, AVG)
 - GROUP BY
- Next week – How to choose a plan?
- Catalogs consulted to parse, optimize, execute plan

Query Evaluation Plan

EMP (ssn, ename, addr, sal, did)


DEPT (did, dname, floor, mgr)

```
SELECT E.ename
FROM Emp E, Dept D
WHERE D.dname = 'Toy'
AND D.did = E.did
```





Query Optimizer selects the evaluation plan

Query Execution Life-Cycle



Query Optimization

- May modify the query plan.



Also, different algorithms for SELECT, JOIN, and PROJECT based on whether an index is available, table sizes, etc.



Selection (on one table)

- An access path defines the strategy to use to do a *selection* on a table, possibly utilizing an index
- Example of a selection condition
 - a predicate: $\text{gpa} > 3.0$ and $\text{age} = 21$
- Examples of access paths
 - File scan
 - Index that *matches* a selection in the query. Examples:
 - B+tree index on the $\langle \text{gpa}, \text{age} \rangle$ attributes
 - B+tree index on gpa
 - B+tree index on age
 - Hash index on age



Selection

- Where R.a *op* value

- Options:

- Heap file

Cost: $O(N)$

- Sorted File

Cost: $O(\log_2 N) + \dots$

- Index

- Hash

Cost: $O(1) + \dots$

- B+Tree: Clustered/Unclustered

Cost: $O(\log_F N) + \dots$



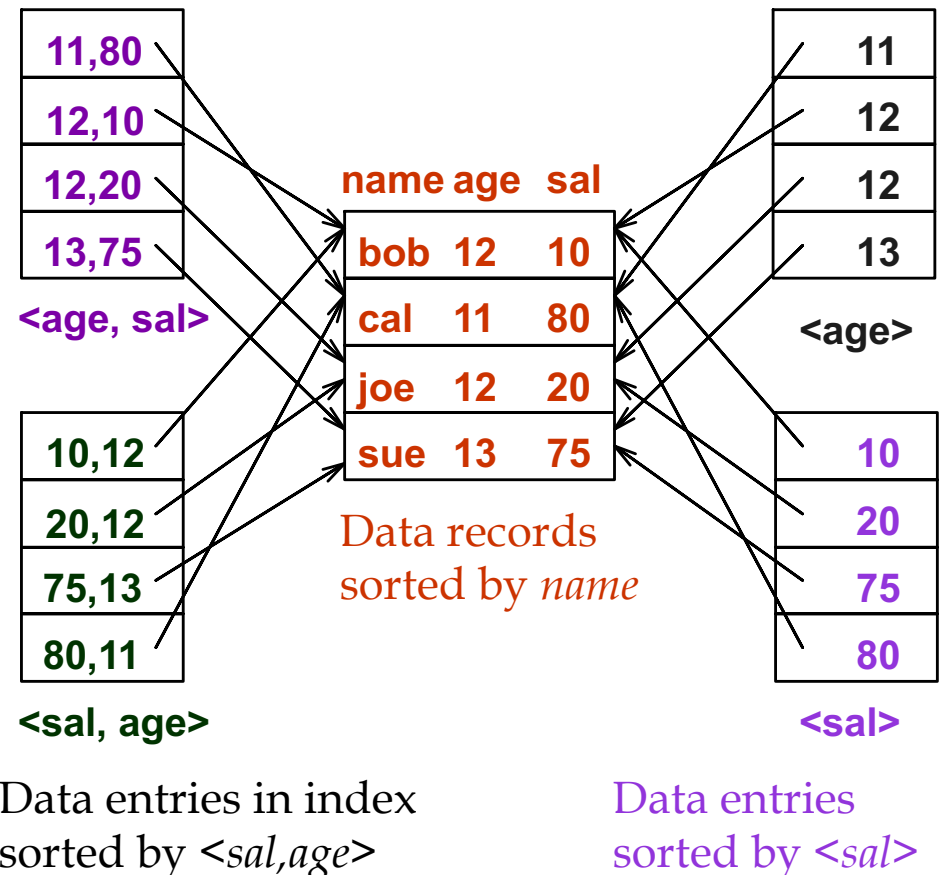
Composite Search keys

- Index on $\langle \text{age} \rangle$: search key is a single-attribute age
- Index on $\langle \text{name}, \text{age} \rangle$: search key is *composite* (name, age) pair.
 - For B+-tree, name is the primary comparison attribute and age matters only when names are equal
 - For hash-index, $h((\text{name}, \text{age}))$ used – both name and age are needed to hash.

Indexes with Composite Search Keys

- **Composite Search Keys:** Search on a combination of fields.
 - **Equality query:** Every field value is equal to a constant value. e.g. wrt $\langle \text{sal}, \text{age} \rangle$ index:
 - age=12 and sal =75
 - **Range query:** Some field value is not a constant. e.g.:
 - age =12; or age=12 and sal > 10
- Data entries in index sorted by search key to support range queries.

Examples of composite key indexes using lexicographic order





Index Matching

- When can we use an index to evaluate a selection **predicate**?
- An index *matches* a predicate if index can be used to evaluate the predicate

Exercise: Index Matching

- Index on $\langle a, b, c \rangle$
 - $a=5$ and $b=3$?
 - $a > 5$ and $b < 3$
 - $b=3$
 - $a=7$ and $b=5$ and $c=4$ and **$d > 4$**
 - $a=7$ and $c=5$

Tree Idx

- yes
- yes
- no!
- yes
- yes

Hash Idx

- no!
- no!
- no!
- yes
- no!

- Index matches (part of) a predicate if
 - Conjunction of terms involving only attributes (**no disjunctions**)
 - Hash: only equality operation, predicate has all index attributes.
 - Tree: Attributes are a prefix of the search key, any ops.



Index Selectivity

- To retrieve Emp records with *age*=30 AND *sal*=4000, an index on $\langle \textit{age}, \textit{sal} \rangle$ would be better than an index on *age* or an index on *sal*.
 - $\langle \textit{age}, \textit{sal} \rangle$ is more *selective* than just $\langle \textit{age} \rangle$ or just $\langle \textit{sal} \rangle$
 - It identifies fewer spurious records that will later be rejected
- If condition is: *age*=80 AND $3000 < \textit{sal} < 20,000$:
 - $\langle \textit{age} \rangle$ index much better than $\langle \textit{sal} \rangle$ index!
- *Ideally, the more selective index preferred*



Index Matching

- Predicate could match more than 1 index
- Hash index on $\langle a, b \rangle$ and B+tree index on $\langle a, c \rangle$

Predicate: $a=7$ and $b=5$ and $c=4$. Which index?

- Option 0: Neither. Simply use file scan
- Option 1: More selective one. Then, scan among the selected records.
- Option2: Use both!
- Algorithm: Intersect rid sets.
 - Sort rids, retrieve rids in both sets.



Quiz: Selection

- Hash index on $\langle a \rangle$ and Hash index on $\langle b \rangle$
 - $a=7$ **or** $b>5$ Which index?
 - Neither! File scan required for $b>5$
- Hash index on $\langle a \rangle$ and B+-tree on $\langle b \rangle$
 - $a=7$ **or** $b>5$ Which index?
 - Option 1: Neither
 - Option 2: Use both! Fetch rids and union
 - Note: Option 1 could be better sometimes. (When?)
- Hash index on $\langle a \rangle$ and B+-tree on $\langle b \rangle$
 - $(a=7$ **or** $c>5)$ and $b > 5$ Which index?
 - B+-tree (high selectivity) or File Scan (poor selectivity)

When to use a B+tree index

- Consider
 - A relation with 1M tuples
 - 100 tuples on a page
 - 500 (key, rid) pairs on a page

data pages
 $= 1\text{M}/100 = 10\text{K pages}$

leaf idx pgs
 $= 1\text{M} / (500 * 0.67)$
 $\sim 3\text{K pages}$

| | 1% Selection | 10% Selection |
|----------------|---------------------------|----------------------------|
| Clustered | $\sim 30 + 100$ | $\sim 300 + 1000$ |
| Non-Clustered | $\sim 30 + 10,000$ | $\sim 300 + 100,000$ |
| NC + Sort Rids | $\sim 30 + (\sim 10,000)$ | $\sim 300 + (\sim 10,000)$ |

- ⇒ Choice of Index access plan, consider:
- 1. Index Selectivity** **2. Clustering**
- ⇒ Similar consideration for hash-based indices



System Catalogs

- To help optimize queries, the system keeps information on each relation
 - name, file name, file structure (e.g., Heap file)
 - attribute name and type, for each attribute
 - index name, for each index
 - integrity constraints
- For each index:
 - structure (e.g., B+ tree) and search key fields
- For each view:
 - view name and definition
- Plus statistics, authorization, buffer pool size, etc.

Example catalog: Attribute Catalog

| attrName | relName | type | position |
|-----------|---------------|---------|----------|
| attr_name | Attribute_Cat | string | 1 |
| rel_name | Attribute_Cat | string | 2 |
| type | Attribute_Cat | string | 3 |
| position | Attribute_Cat | integer | 4 |
| sid | Students | string | 1 |
| name | Students | string | 2 |
| login | Students | string | 3 |
| age | Students | integer | 4 |
| gpa | Students | real | 5 |
| fid | Faculty | string | 1 |
| fname | Faculty | string | 2 |
| sal | Faculty | real | 3 |

Catalogs are themselves stored as relations



Basic Join Strategies



Join Operator

```
SELECT *  
FROM   Reserves R, Sailors S  
WHERE  R.sid = S.sid
```


- Commercial systems spend a lot of effort optimizing equality joins
 - Why is this important?
 - What is the major source of performance cost when joining two (large) relations?
- Cost Metric: # of I/Os
 - (We'll ignore final output cost)

Join Operator

- Many different ways of evaluating joins

Sailors

| sid | name |
|-----|-------|
| 1 | Lucky |
| 2 | Rusty |
| 3 | Bob |
| 4 | Fred |



Reserves

| sid | bid |
|-----|-----|
| 1 | 100 |
| 1 | 200 |
| 3 | 300 |
| 4 | 200 |

How would you evaluate this join in memory?

Simple Nested Loops Join

```
foreach tuple r in R do
    foreach tuple s in S do
        if r.sid == s.sid then add <r, s> to result
```

- Cost Model

- $||R||$ = # tuples in R
- $|R|$ = # pages in R

Slightly different
notation from
textbook!

- How many I/Os ?

$$|R| + ||R|| * |S|$$



Page-Oriented Nested Loops

- Page-oriented Nested Loops join:
 - For each page of R, get each page of S, and join
 - How many I/Os?
$$|R| + |R| * |S|$$
 - If S is the outer, then $|S| + |R| * |S|$

How many buffer pages does this use?

Block Nested Loops

- Can we exploit available memory?
- Suppose we have B buffer pages available.
- Use B-2 pages to hold a block of **outer** R.

```
foreach block of B-2 pages of R do
  foreach page of S do
    foreach r in the B-2 R pages do
      foreach tuple s in the S page do
        if r.sid == s.sid then add <r, s> to result
```

■ Cost: $|R| \bowtie |S| * \left\lceil \frac{|R|}{B-2} \right\rceil$

What is the cost
if the smaller relation
fits entirely in memory?

PNL vs BNL

- $|R| = 128$ $|S| = 64$ $B = 8$
tuples/page for both S and R = 10
- Page NL
 - Scan outer: 64
 - Join: $64 * 128 = 8192$
 - TOTAL: 8256

In PNL, which rel. should be the outer?

- Block NL
 - Scan outer: 64
 - Join: $\lceil 64/6 \rceil * 128 = 1408$
 - TOTAL: 1472

What about simple nested loops?



Announcements

- Optional Exercises:
 - 12.1 (1-4), 12.3, 12.5