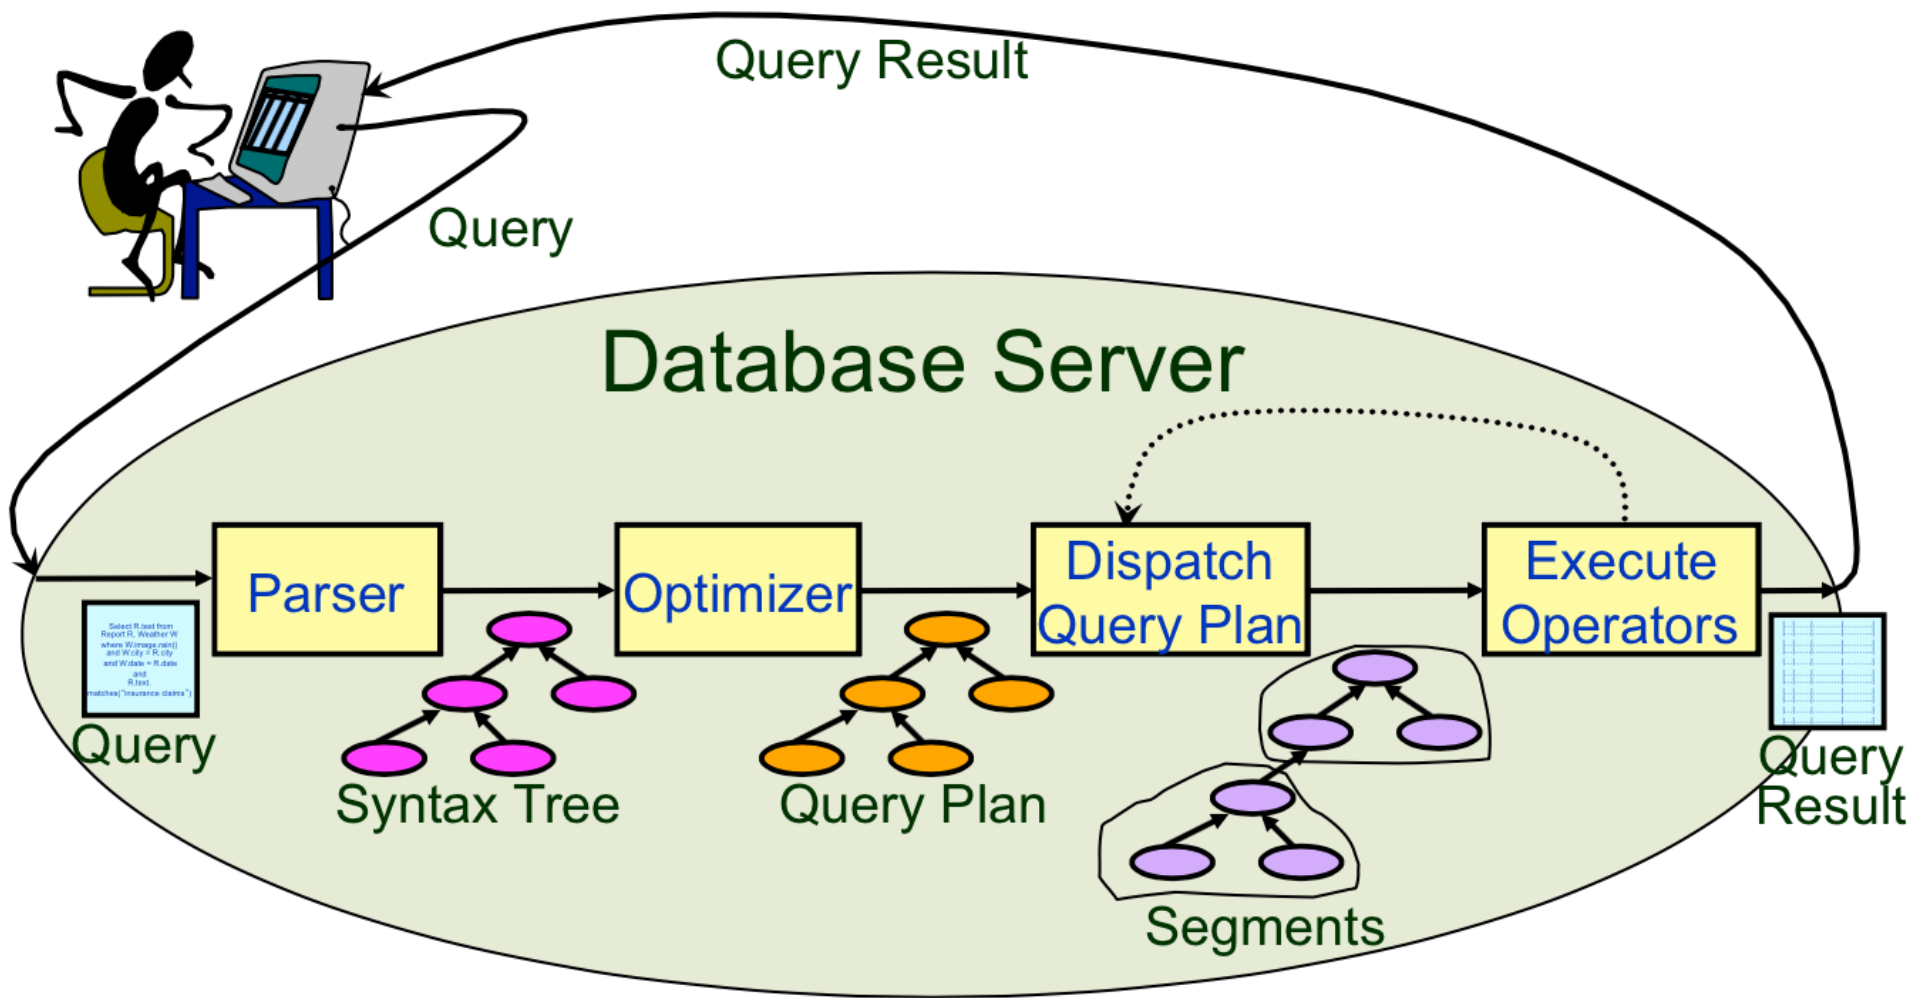




Projections

Chapter 12 and 14

Query Execution Life-Cycle



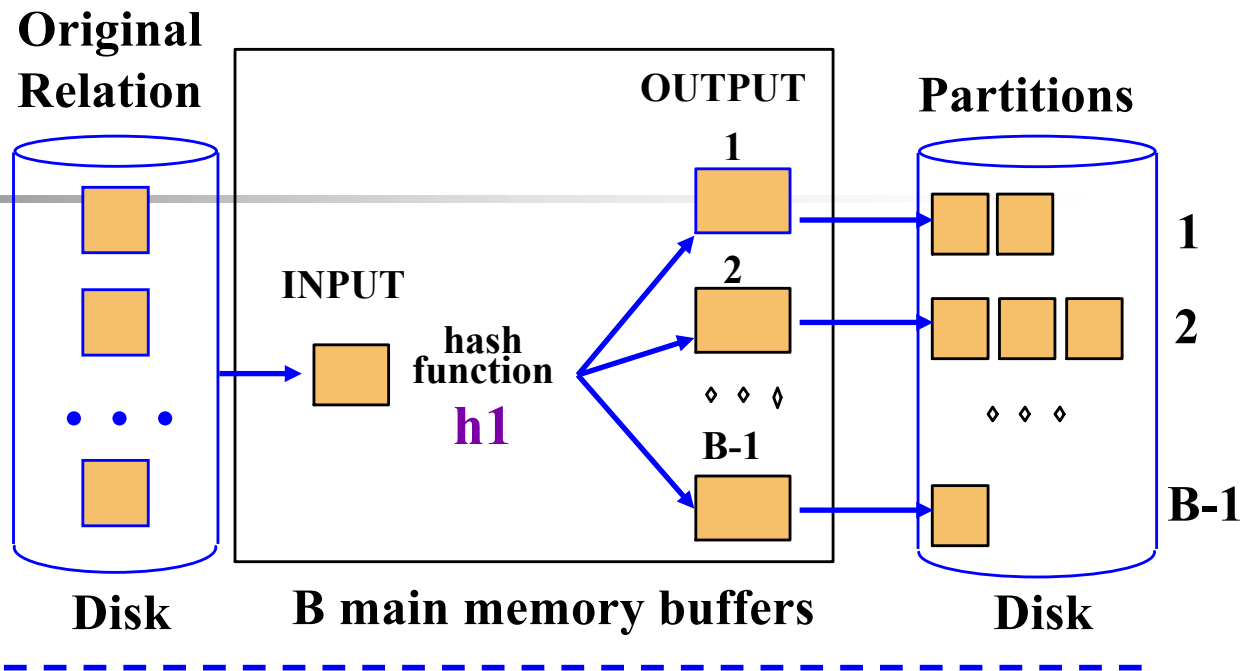


Projection

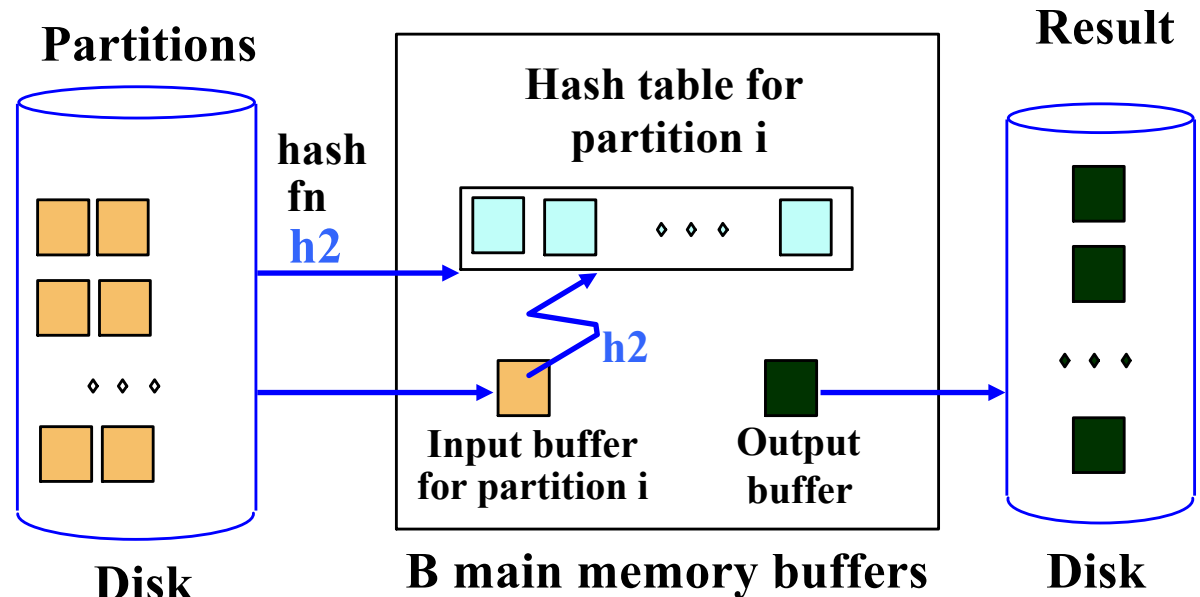
- Select R.a, R.d
 - Straightforward implementation!
- Select DISTINCT R.a, R.d
 - Remove attributes
 - Eliminate duplicates
- Algorithms for Projection DISTINCT:
 - Sorting: Sort on all the projected attributes
 - Pass 0: eliminate unwanted fields. Tuples in the sorted-runs may be smaller
 - Eliminate duplicates in the merge pass & sort
 - Hashing: Two phases
 - Partitioning
 - Duplicate elimination

Hashing

**Hash on projected
Attributes to
Eliminate duplicates**



Recursively apply
hash-based
projection
technique to
handle partition
overflow problem





Projection

- Sort-based approach for SELECT DISTINCT
 - better handling of skew
 - result is sorted
 - Thus, more commonly used than hash-based approach



Index-only Scans

- Index-only scan can be much more efficient if
 - Projection attributes subset of index attributes
- Apply projection techniques to data entries (much smaller!)
- For handling SELECT DISTINCT, an additional optimization possible if the ordered (i.e., tree) index contains all projection attributes as prefix of search key:
 - Retrieve index data entries in order (no sorting necessary)
 - Discard unwanted fields
 - Compare adjacent entries to eliminate duplicates (if required)



Set Operations

- \cap and \bowtie special cases of join
- \cup and $-$ similar; we'll do \cup
 - Both require duplicate elimination
- Duplicate elimination algorithms for \cup :
 1. Sorting:
 - Sort both relations (on all attributes).
 - Merge sorted relations eliminating duplicates.
 2. Hashing:
 - Partition R and S
 - Build hash table for R_i .
 - Probe with tuples in S_i , add to table if not a duplicate



Aggregates

- Sorting Approach
 - Sort on GROUP BY attributes (if any)
 - Scan sorted tuples, computing running aggregate
 - Min, Max
 - Count
 - Sum
 - Average: compute from sum and count
 - During scan, when the group by attribute changes (e.g., 2, 2, 2, **3**), output aggregate result



Aggregates

- Hashing Approach
 - Hash on GROUP BY attributes (if any)
 - Hash entry: grouped attributes + running aggregate
 - Scan tuples, probe hash table, update hash entry
 - Scan hash table, and output each hash entry
- Cost: Scan relation!



Using an Index for Aggregation

- Usually, index is not useful for aggregation operators.
- But sometimes it is:
 - When the search key contains all the relevant attributes: index-only scan may be feasible, rather than fetching all the data records.
 - If the GROUP BY attribute list is a prefix of the search key, we can retrieve data records in the order required for the grouping operation and thereby avoid the sorting step.



Announcements

- Optional Exercises:
 - 12.1 (1-4), 12.3, 12.5
 - 13.1, 13.3
 - 14.1 (2, 3, 4, 6, 7, 8, 9, 10), 14